



Modular QoS Configuration Guide for Cisco NCS 5500 Series Routers, IOS XR Release 7.8.x

First Published: 2022-10-21

Last Modified: 2023-01-26

Americas Headquarters

Cisco Systems, Inc.
170 West Tasman Drive
San Jose, CA 95134-1706
USA
<http://www.cisco.com>
Tel: 408 526-4000
800 553-NETS (6387)
Fax: 408 527-0883

THE SPECIFICATIONS AND INFORMATION REGARDING THE PRODUCTS IN THIS MANUAL ARE SUBJECT TO CHANGE WITHOUT NOTICE. ALL STATEMENTS, INFORMATION, AND RECOMMENDATIONS IN THIS MANUAL ARE BELIEVED TO BE ACCURATE BUT ARE PRESENTED WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED. USERS MUST TAKE FULL RESPONSIBILITY FOR THEIR APPLICATION OF ANY PRODUCTS.

THE SOFTWARE LICENSE AND LIMITED WARRANTY FOR THE ACCOMPANYING PRODUCT ARE SET FORTH IN THE INFORMATION PACKET THAT SHIPPED WITH THE PRODUCT AND ARE INCORPORATED HEREIN BY THIS REFERENCE. IF YOU ARE UNABLE TO LOCATE THE SOFTWARE LICENSE OR LIMITED WARRANTY, CONTACT YOUR CISCO REPRESENTATIVE FOR A COPY.

The Cisco implementation of TCP header compression is an adaptation of a program developed by the University of California, Berkeley (UCB) as part of UCB's public domain version of the UNIX operating system. All rights reserved. Copyright © 1981, Regents of the University of California.

NOTWITHSTANDING ANY OTHER WARRANTY HEREIN, ALL DOCUMENT FILES AND SOFTWARE OF THESE SUPPLIERS ARE PROVIDED "AS IS" WITH ALL FAULTS. CISCO AND THE ABOVE-NAMED SUPPLIERS DISCLAIM ALL WARRANTIES, EXPRESSED OR IMPLIED, INCLUDING, WITHOUT LIMITATION, THOSE OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NON-INFRINGEMENT OR ARISING FROM A COURSE OF DEALING, USAGE, OR TRADE PRACTICE.

IN NO EVENT SHALL CISCO OR ITS SUPPLIERS BE LIABLE FOR ANY INDIRECT, SPECIAL, CONSEQUENTIAL, OR INCIDENTAL DAMAGES, INCLUDING, WITHOUT LIMITATION, LOST PROFITS OR LOSS OR DAMAGE TO DATA ARISING OUT OF THE USE OR INABILITY TO USE THIS MANUAL, EVEN IF CISCO OR ITS SUPPLIERS HAVE BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

Any Internet Protocol (IP) addresses and phone numbers used in this document are not intended to be actual addresses and phone numbers. Any examples, command display output, network topology diagrams, and other figures included in the document are shown for illustrative purposes only. Any use of actual IP addresses or phone numbers in illustrative content is unintentional and coincidental.

All printed copies and duplicate soft copies of this document are considered uncontrolled. See the current online version for the latest version.

Cisco has more than 200 offices worldwide. Addresses and phone numbers are listed on the Cisco website at www.cisco.com/go/offices.

The documentation set for this product strives to use bias-free language. For purposes of this documentation set, bias-free is defined as language that does not imply discrimination based on age, disability, gender, racial identity, ethnic identity, sexual orientation, socioeconomic status, and intersectionality. Exceptions may be present in the documentation due to language that is hardcoded in the user interfaces of the product software, language used based on standards documentation, or language that is used by a referenced third-party product.

Cisco and the Cisco logo are trademarks or registered trademarks of Cisco and/or its affiliates in the U.S. and other countries. To view a list of Cisco trademarks, go to this URL: <https://www.cisco.com/c/en/us/about/legal/trademarks.html>. Third-party trademarks mentioned are the property of their respective owners. The use of the word partner does not imply a partnership relationship between Cisco and any other company. (1721R)

© 2022 Cisco Systems, Inc. All rights reserved.



CONTENTS

PREFACE

Preface ix

Changes to this Document ix

Communications, Services, and Additional Information ix

CHAPTER 1

New and Changed QoS Features 1

New and Changed QoS Features 1

CHAPTER 2

YANG Data Models for QoS Features 3

Using YANG Data Models 3

CHAPTER 3

Configuring Modular QoS Service Packet Classification 5

Packet Classification Overview 6

Traffic Class Elements 8

Default Traffic Class 9

Create a Traffic Class 9

QoS Classification Based on Packet Length 11

Increase in QoS Policer Allocation on Cisco NCS 5700 Series Line Cards 14

Traffic Policy Elements 15

Create a Traffic Policy 16

Scaling of Unique Ingress Policy Maps 17

Limitations and Restrictions 18

Attach a Traffic Policy to an Interface 18

Packet Marking 20

Handling QoS for Locally Originated Packets 22

Protecting Locally Originated BFD Packets 24

Prioritization of IS-IS and ARP Packets to Manage Transit Traffic 25

QoS Re-marking of IP Packets in Egress Direction	27
QoS Re-marking of Ethernet Packets in Egress Direction	28
QoS L2 Re-marking of Ethernet Packets in Egress Direction	28
QoS L2 Re-Marking of Ethernet Packets on L3 Flows in Egress Direction	29
QoS L2 Re-Marking of Ethernet Packets on L3 Flows in Egress Direction on L3 sub-interfaces	32
Layer 2 Ingress QoS Matching for IPv4 and IPv6 Destination Addresses	37
Overview	38
Guidelines and Limitations	38
Configure Layer 2 Ingress QoS Matching for IPv4 and IPv6 Destination Addresses	38
Bundle Traffic Policies	41
Shared Policy Instance	41
Restrictions and Guidelines	43
Attaching a Shared Policy Instance to Multiple Subinterfaces	43
Ingress Short-Pipe	45
Restrictions and Other Important Points	45
Configure Ingress Short-Pipe	46
Selective Egress Policy-Based Queue Mapping	47
Restrictions and Other Important Points	49
Configure Selective Egress Policy-Based Queue Mapping	49
Configuring QoS Groups with an ACL	52
Configuring an ACL with Fragment Match	55
Restrictions and Guidelines	56
Configuring an ACL with Fragment Match	56
Restrictions	57
In-Place Policy Modification	60
References for Modular QoS Service Packet Classification	61
Specification of the CoS for a Packet with IP Precedence	61
IP Precedence Bits Used to Classify Packets	61
IP Precedence Value Settings	61
IP Precedence Compared to IP DSCP Marking	62
Conditional Marking of MPLS Experimental bits for L3VPN Traffic	62
QoS DSCP Preservation	62
QoS IP DSCP Preservation for SR-TE	64
QoS DSCP Preservation for mLDP	67

Policy-map for conditional marking of incoming IPv4 and IPv6 traffic	69
Policy-map for conditional marking of outgoing MPLS traffic	70
Conditional Marking of MPLS Experimental bits for L2VPN Traffic	71
Policy-map for conditional marking of incoming traffic	71
Policy-map for conditional marking of outgoing MPLS traffic	72
Conditional Marking of MPLS Experimental Bits for EVPN-VPWS Single-Homing Services	73
QPPB	75
Configuration Workflow	77
Configuring QPPB on an Interface	79
Egress Interface Configuration	80

CHAPTER 4

Configuring Modular QoS Congestion Avoidance	81
Modular QoS Congestion Avoidance	81
Tail Drop and the FIFO Queue	82
Configure Tail Drop	82
Random Early Detection and TCP	84
Configure Random Early Detection	84
Weighted Random Early Detection	86
Average Queue Size for WRED	86
Configure Weighted Random Early Detection	87
Configure WRED Counters by Class	89
Explicit Congestion Notification (ECN)	93
Configure ECN Maximum Marking Probability	97

CHAPTER 5

Priority Flow Control Overview	101
Restrictions	101
Configuring Priority Flow Control Thresholds	103
Priority Flow Control Watchdog Overview	109
Configure a Priority Flow Control Watchdog Interval	109
Monitoring and Logging Packet Drops on Lossless PFC-Enabled Queues	112
A Brief Background	113
Highlights of Logging and Monitoring Packet Drops	113
Restrictions for Monitoring and Logging Packet Drops	114

CHAPTER 6**Configuring Modular QoS Congestion Management 115**

Congestion Management Overview	115
Ingress Traffic Management Model	116
Class-based Weighted Fair Queueing	118
Bandwidth Remaining	118
Configure Minimum Bandwidth and Bandwidth Remaining	118
Low-Latency Queuing with Strict Priority Queuing	120
Configuring Low Latency Queuing with Strict Priority queuing	120
Overhead Accounting	122
Traffic Shaping	125
Configure Traffic Shaping	125
4K Pseudowire on Bundle with QoS Enhancement	127
Restrictions and Guidelines	128
Configure Two-Level Shaper Policy on Main Interface	128
Traffic Policing	129
Committed Bursts and Excess Bursts	130
Single-Rate Policer	131
Configure Traffic Policing (Single-Rate Two-Color)	132
Configure Traffic Policing (Single-Rate Three-Color)	133
Two-Rate Policer	135
Configure Traffic Policing (Two-Rate Three-Color)	135
Packets-Per-Second-Based Policer	138
Per-thousand and Per-million Units	141
Shared Policer	141
Policer Bucket Shared	141
Policer Bucket Referred	142
Shared Policer Statistics	142
Restrictions and Guidelines	142
Configuring Shared Policer	143
References for Modular QoS Congestion Management	145
Committed Bursts	145
Excess Bursts	145
Two-Rate Policer Details	146

CHAPTER 7	Configuring Modular QoS on Link Bundles	147
	QoS on Link Bundles	147
	Load Balancing	148
	Configure QoS on Link Bundles	148

CHAPTER 8	Configuring Hierarchical Modular QoS	153
	Overview of Hierarchical Modular QoS	153
	Restrictions for Configuring H-QoS	154
	Configuring Hierarchical Queuing	155
	Conform Aware Hierarchical Policy Overview	161
	Conform Aware Hierarchical Policy Restrictions	161
	Configuring Conform Aware Hierarchy Policy	161

CHAPTER 9	Configure Egress Traffic Management	165
	Egress Traffic Management	166
	More About the Egress Traffic Management Model	168
	QoS Policy Configuration Rules	169
	Configure Egress Traffic Management	170
	Low Rate Shaper	186



Preface

This preface contains these sections:

- [Changes to this Document, on page ix](#)
- [Communications, Services, and Additional Information, on page ix](#)

Changes to this Document

This table lists the changes made to this document since it was first published.

Date	Summary
November 2022	Initial release of this document.

Communications, Services, and Additional Information

- To receive timely, relevant information from Cisco, sign up at [Cisco Profile Manager](#).
- To get the business impact you're looking for with the technologies that matter, visit [Cisco Services](#).
- To submit a service request, visit [Cisco Support](#).
- To discover and browse secure, validated enterprise-class apps, products, solutions and services, visit [Cisco Marketplace](#).
- To obtain general networking, training, and certification titles, visit [Cisco Press](#).
- To find warranty information for a specific product or product family, access [Cisco Warranty Finder](#).

Cisco Bug Search Tool

[Cisco Bug Search Tool](#) (BST) is a web-based tool that acts as a gateway to the Cisco bug tracking system that maintains a comprehensive list of defects and vulnerabilities in Cisco products and software. BST provides you with detailed defect information about your products and software.



CHAPTER 1

New and Changed QoS Features

- [New and Changed QoS Features, on page 1](#)

New and Changed QoS Features

Table 1: QoS Features Added or Modified in IOS XR Release 7.8.x

Feature	Description	Changed in Release	Where Documented
Configure WRED Counters by Class on Cisco NCS 5700 Series Base (Non-SE) Version Line Cards	You can now view WRED statistics per class on Cisco NCS 5700 Series base (non-SE) line cards operating in native mode.	Release 7.8.1	Configure WRED Counters by Class
Increase in QoS Policer Allocation on Cisco NCS 5700 Series Line Cards	You can now allocate up to a maximum of 64000 policers for QoS.	Release 7.8.1	Increase in QoS Policer Allocation on Cisco NCS 5700 Series Line Cards



CHAPTER 2

YANG Data Models for QoS Features

This chapter provides information about the YANG data models for QoS features.

- [Using YANG Data Models, on page 3](#)

Using YANG Data Models

Cisco IOS XR supports a programmatic way of configuring and collecting operational data of a network device using YANG data models. Although configurations using CLIs are easier and human-readable, automating the configuration using model-driven programmability results in scalability.

The data models are available in the release image, and are also published in the [Github](#) repository. Navigate to the release folder of interest to view the list of supported data models and their definitions. Each data model defines a complete and cohesive model, or augments an existing data model with additional XPaths. To view a comprehensive list of the data models supported in a release, navigate to the **Available-Content.md** file in the repository. The unsupported sensor paths are documented as deviations. For example, `openconfig-acl.yang` provides details about the supported sensor paths, whereas `cisco-xr-openconfig-acl-deviations.yang` shows the unsupported sensor paths for `openconfig-acl.yang` model on Cisco IOS XR routers.

You can also view the data model definitions using the [YANG Data Models Navigator](#) tool. This GUI-based and easy-to-use tool helps you explore the nuances of the data model and view the dependencies between various containers in the model. You can view the list of models supported across Cisco IOS XR releases and platforms, locate a specific model, view the containers and their respective lists, leaves, and leaf lists presented visually in a tree structure. This visual tree form helps you get insights into nodes that can help you automate your network.

To get started with using the data models, see the *Programmability Configuration Guide for Cisco NCS 5500 Series Routers*.



CHAPTER 3

Configuring Modular QoS Service Packet Classification

This chapter covers these topics:

- [Packet Classification Overview, on page 6](#)
- [Traffic Class Elements, on page 8](#)
- [Traffic Policy Elements, on page 15](#)
- [Ingress Short-Pipe, on page 45](#)
- [Selective Egress Policy-Based Queue Mapping, on page 47](#)
- [Configuring QoS Groups with an ACL, on page 52](#)
- [Configuring an ACL with Fragment Match, on page 55](#)
- [Restrictions , on page 57](#)
- [In-Place Policy Modification, on page 60](#)
- [References for Modular QoS Service Packet Classification, on page 61](#)
- [QPPB, on page 75](#)

Packet Classification Overview

Table 2: Feature History Table

Feature Name	Release Information	Feature Description
Cisco NC57 Compatibility Mode: QoS Enablement on Layer 2 MPLS/BGP	Release 7.3.1	<p>This feature is now supported on routers that have the Cisco NC57 line cards installed and operate in the compatibility mode.</p> <p>The following Layer 2 services are supported:</p> <ul style="list-style-type: none"> • Local switching [xconnect or bridging] • Layer 2 VPN – Virtual Private Wire Service (VPWS) <p>Apart from packet classification, this feature is available for the following QoS operations:</p> <ul style="list-style-type: none"> • Modular QoS Congestion Avoidance , on page 81 • Configuring Modular QoS Congestion Management, on page 115 • QoS on Link Bundles, on page 147 • Configuring Hierarchical Modular QoS, on page 153

Feature Name	Release Information	Feature Description
Cisco NC57 Native Mode: QoS Enablement on Layer 2 Services	Release 7.4.1	<p>This feature is now supported on routers that have the Cisco NC57 line cards installed and operate in the native mode.</p> <p>The following Layer 2 services are supported:</p> <ul style="list-style-type: none"> • Local switching [xconnect or bridging] • Layer 2 VPN – Virtual Private Wire Service (VPWS) <p>Apart from packet classification, this feature is available for the following QoS operations:</p> <ul style="list-style-type: none"> • Modular QoS Congestion Avoidance , on page 81 • Configuring Modular QoS Congestion Management, on page 115 • QoS on Link Bundles, on page 147 • Configuring Hierarchical Modular QoS, on page 153

Packet classification involves categorizing a packet within a specific group (or class) and assigning it a traffic descriptor to make it accessible for QoS handling on the network. The traffic descriptor contains information about the forwarding treatment (quality of service) that the packet should receive. Using packet classification, you can partition network traffic into multiple priority levels or classes of service. The source agrees to adhere to the contracted terms and the network promises a quality of service. Traffic policers and traffic shapers use the traffic descriptor of a packet to ensure adherence to the contract.

Traffic policers and traffic shapers rely on packet classification features, such as IP precedence, to select packets (or traffic flows) traversing a router or interface for different types of QoS service. After you classify packets, you can use other QoS features to assign the appropriate traffic handling policies including congestion management, bandwidth allocation, and delay bounds for each traffic class.

The Modular Quality of Service (QoS) CLI (MQC) defines the traffic flows that must be classified, where each traffic flow is called a class of service, or class. Later, a traffic policy is created and applied to a class. All traffic not identified by defined classes fall into the category of a default class.

You can classify packets at the ingress on L3 subinterfaces for (CoS, DEI) for IPv4, IPv6, and MPLS flows. IPv6 packets are forwarded by paths that are different from those for IPv4. To enable classification of IPv6 packets based on (CoS, DEI) on L3 subinterfaces, run the `hw-module profile qos ipv6 short-l2qos-enable` command and reboot the line card for the command to take effect.

Starting with Cisco IOS XR Release 7.4.1 systems with Cisco NC57 line cards running in native mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Traffic Class Elements

The purpose of a traffic class is to classify traffic on your router. Use the **class-map** command to define a traffic class.

A traffic class contains three major elements:

- A name
- A series of **match** commands - to specify various criteria for classifying packets.
- An instruction on how to evaluate these **match** commands (if more than one **match** command exists in the traffic class)

Packets are checked to determine whether they match the criteria that are specified in the **match** commands. If a packet matches the specified criteria, that packet is considered a member of the class and is forwarded according to the QoS specifications set in the traffic policy. Packets that fail to meet any of the matching criteria are classified as members of the default traffic class.

This table shows the details of match types that are supported on the router.

Match Type Supported	Min, Max	Max Entries	Support for Match NOT	Support for Ranges	Direction Supported on Interfaces
IPv4 DSCP IPv6 DSCP DSCP	(0,63)	64	Yes	Yes	Ingress
IPv4 Precedence IPv6 Precedence Precedence	(0,7)	8	Yes	No	Ingress
MPLS Experimental Topmost	(0,7)	8	Yes	No	Ingress
Access-group	Not applicable	8	No	Not applicable	Ingress
QoS-group	(1,7) (1,511) for peering profile	7	No	No	<ul style="list-style-type: none"> • Egress • Ingress for QoS Policy Propagation Using Border Gateway Protocol (QPPB) • Ingress for peering profile
Traffic-class	(1,7)	7	No	No	<ul style="list-style-type: none"> • Egress

Match Type Supported	Min, Max	Max Entries	Support for Match NOT	Support for Ranges	Direction Supported on Interfaces
Protocol	(0,255)	1	Yes	Not applicable	Ingress



Note Egress queue statistics are displayed only for those classes which have a corresponding match criteria in the egress. Therefore, if you have a **set traffic-class** *x* configured in the ingress, you must have a corresponding **match traffic-class** *x* in the egress, in order to see the statistics in the egress side.



Note A maximum value of up to 64 unique queues is supported.

Default Traffic Class

Unclassified traffic (traffic that does not meet the match criteria specified in the traffic classes) is treated as belonging to the default traffic class.

If the user does not configure a default class, packets are still treated as members of the default class. However, by default, the default class has no enabled features. Therefore, packets belonging to a default class with no configured features have no QoS functionality. These packets are then placed into a first in, first out (FIFO) queue and forwarded at a rate determined by the available underlying link bandwidth. This FIFO queue is managed by a congestion avoidance technique called tail drop.

For egress classification, match on **traffic-class** (1-7) is supported. Match **traffic-class 0** cannot be configured. The class-default in the egress policy maps to **traffic-class 0**.

This example shows how to configure a traffic policy for the default class:

```
configure
policy-map ingress_policy1
class class-default
  police rate percent 30
!
```

Create a Traffic Class

To create a traffic class containing match criteria, use the **class-map** command to specify the traffic class name, and then use the **match** commands in class-map configuration mode, as needed.

Guidelines

- Users can provide multiple values for a match type in a single line of configuration; that is, if the first value does not meet the match criteria, then the next value indicated in the match statement is considered for classification.
- Use the **not** keyword with the **match** command to perform a match based on the values of a field that are not specified.

- All **match** commands specified in this configuration task are considered optional, but you must configure at least one match criterion for a class.
- If you specify **match-any**, one of the match criteria must be met for traffic entering the traffic class to be classified as part of the traffic class. This is the default. If you specify **match-all**, the traffic must match all the match criteria.
- From Release 7.7.1 onwards, for the **match access-group** command, QoS classification based on the packet length field in the IPv4 and IPv6 headers is supported. Prior to this, support was not available for packet length and TTL (time to live) fields.
- Deny action in the ACL or an empty ACL (contains no rules, only remarks) is not supported for the **match access-group command** when an ACL list is used within a **class-map**.
The configuration is rejected when you attach such a **policy-map** to an interface.
- The **traffic-class** and **discard-class** are supported only in egress direction, and these are the only match criteria supported in egress direction.
- The egress default class implicitly matches **qos-group 0** for marking policy and **traffic-class 0** for queuing policy.
- Multicast takes a system path that is different than unicast on router, and they meet later on the egress in a multicast-to-unicast ratio of 20:80 on a per interface basis. This ratio is maintained on the same priority level as that of the traffic.
- Egress QoS for multicast traffic treats traffic classes 0-5 as low-priority and traffic classes 6-7 as high priority. Currently, this is not user-configurable.
- Egress shaping does not take effect for multicast traffic in the high priority (HP) traffic classes. It only applies to unicast traffic.
- If you set a traffic class at the ingress policy and do not have a matching class at egress for the corresponding traffic class value, then the traffic at ingress with this class will not be accounted for in the default class at the egress policy map.
- Only traffic class 0 falls in the default class. A non-zero traffic class assigned on ingress but with no assigned egress queue, falls neither in the default class nor any other class.

Configuration Example

You have to accomplish the following to complete the traffic class configuration:

1. Creating a class map
2. Specifying the match criteria for classifying the packet as a member of that particular class
(For a list of supported match types, see [Traffic Class Elements, on page 8](#).)

```
Router# configure
Router(config)# class-map match-any qos-1
Router(config-cmap)# match qos-group 1
Router(config-cmap)# end-class-map
Router(config-cmap)# commit
```

Use this command to verify the class-map configuration:

```
Router#show class-map qos-1
1) ClassMap: qos-1      Type: qos
   Referenced by 2 Polymaps
```

Also see, [Running Configuration, on page 19](#).

Also see, [Verification, on page 19](#).

Related Topics

- [Traffic Class Elements, on page 8](#)
- [Traffic Policy Elements, on page 15](#)

Associated Commands

- [class-map](#)
- [match access-group](#)
- [match dscp](#)
- [match mpls experimental topmost](#)
- [match precedence](#)
- [match qos-group](#)

QoS Classification Based on Packet Length

Table 3: Feature History Table

Feature Name	Release Information	Feature Description
QoS Classification Based on Packet Length	Release 7.7.1	<p>You can add an ingress QoS policy on an ACL that filters IPv6 and IPv4 packets based on the packet length. The ACL specifies packet length criteria such as equal to, lesser than, greater than, and so on, to filter packets and prevent distributed-denial-of-service (DDoS) attacks. The QoS policy specifies the forwarding decision.</p> <p>The ACL provides an additional QoS match criteria for IP traffic, thereby providing the combined benefit of network security and traffic classification.</p> <p>This feature is supported on the NCS 5500 Series routers and NCS 5700 Series routers on NC57-18DD-SE and NC57-24DD PIDs.</p>

An ACL consists of one or more access control entries that collectively define the network traffic profile and describe the access rights related to a particular security identifier or user. This profile can then be referenced by Cisco IOS XR software features such as traffic filtering, route filtering, QoS classification, and access control. For more information on packet classification, see [Packet Classification Overview](#). For more

information on creating IP ACLs, see the [Understanding Access Lists](#) section in the [IP Addresses and Services Configuration Guide](#).

Guidelines and Limitations for QoS Classification Based on Packet Length

- Packet-length configurations for IPv6 packets with optional extension headers will be classified under **class-default**.
- Now, the **match access-group** command can define the QoS classification based on the packet length field in the IPv6 and IPv4 headers. The **show policy-map interface** command displays the corresponding configuration information.
- To enable IPv6 traffic classification based on the packet length, configure the **hw-module profile qos ipv6 short-l2qos-enable** command in configuration mode.
- When the packet-length of the IPv6 and IPv4 traffic is considered, the L2 header length is included.

Configure QoS Classification Based on Packet Length

Follow these steps to classify ingress traffic based on the packet length:

- Create an ACL that allows IPv4 traffic with the specified packet length.
- Create QoS class-map and policy-map for ingress IPv4 traffic classification based on its packet length.
- Apply the policy-map on the designated interface.

Configuration Example for QoS Classification Based on Packet Length



Note This example enables traffic filtering for a specific packet length value by using the **packet-length eq** command. Instead of **eq**, you can also use one of these options:

- **gt** to filter packets with a greater value.
- **lt** to filter packets with a lesser value.
- **neq** to filter packets that do not contain the specified value.
- **range** to filter packets in the specified range of values.

```
/* Create an ACL that allows IPv4 traffic with the specified packet length */
```

```
Router# configure
Router(config)# ipv4 access-list aclv4_pktlen
Router(config-ipv4-acl)# 10 permit ipv4 any any packet-length eq 100
Router(config-ipv4-acl)# commit
```

Based on ACL **aclv4_pktlen**, create a QoS policy and apply it on the ingress IPv4 traffic.

```
/* Create QoS class-map and policy-map for ingress IPv4 traffic classification based on its
packet length */
```

```
Router# configure
Router(config)# class-map match-any c_pktlen
Router(config-cmap)# match access-group ipv4 aclv4_pktlen
Router(config-cmap)# end-class-map
```

```

Router(config)# policy-map p_pktlen
Router(config-pmap)# class c_pktlen
Router(config-pmap-c)# set traffic-class 2
Router(config-pmap-c)# exit
Router(config-pmap)# class class-default
Router(config-pmap-c)# end-policy-map
Router(config-pmap-c)# commit
Router(config-pmap-c)# root

```

The policy-map **p_pktlen** is associated with interface HundredGigabitEthernet 0/0/0/0.

```
/* Apply the policy-map on the designated Ethernet interface */
```

```

Router(config)# interface hundredGigE 0/0/0/0
Router(config-if)# service-policy input p_pktlen
Router(config-if)# commit

```

Running Configuration

```
/* Running configuration for QoS Match Packet Length */
```

```

Router# show running-configuration

..
..

ipv4 access-list aclv4_pktlen
 10 permit ipv4 any packet-length eq 100

class-map match-any c_pktlen
 match access-group ipv4 aclv4_pktlen
end-class-map
!
policy-map p_pktlen
 class c_pktlen
   set traffic-class 2
!
class class-default
!
end-policy-map
!
interface HundredGigE0/0/0/0
 service-policy input p_pktlen
!

```

Verification

```
/* This displays the policy information applied in the input direction: */
```

```

Router# show policy-map interface hundredGigE 0/0/0/0 input

HundredGigE0/0/0/0 input: p_pktlen

Class c_pktlen
  Classification statistics      (packets/bytes)      (rate - kbps)
  Matched                      :      53186/54090162      3769
  Transmitted                  :      53186/54090162      3769
  Total Dropped                :                0/0                0

Class class-default
  Classification statistics      (packets/bytes)      (rate - kbps)
  Matched                      :                0/0                0
  Transmitted                  :                0/0                0

```

Total Dropped : 0/0 0

Policy Bag Stats time: 1657119549887 [Local Time: 07/06/22 14:59:09.887]

Increase in QoS Policer Allocation on Cisco NCS 5700 Series Line Cards

Table 4: Feature History Table

Feature Name	Release Information	Feature Description
Increase in QoS Policer Allocation on Cisco NCS 5700 Series Line Cards	Release 7.8.1	<p>NCS 5700 Series routers can now handle a maximum of 64000 QoS policers by reallocating unused BGP Flowspec policer resources to QoS.</p> <p>Before this release, the QoS policer count was limited to 32000. With more QoS policers, you can apply policies on more sub-interfaces.</p> <p>This feature is supported on NC57-18DD-SE and NC57-36H-SE line cards.</p> <p>This feature introduces the following command:</p> <ul style="list-style-type: none"> • hw-module profile qos policer-scale

Traffic policing involves checking the packet rate for a specific classification of flows to determine if it conforms to specified rate settings. The policer is like a gate that allows a certain number of packets (or bytes) through a network, through physical and logical interfaces. Policers are allocated from the SE line card's Forwarding ASIC for different technologies. 32000 policers are allocated for QoS and BGP Flowspec, each.

To accommodate QoS policing on more sub-interfaces, you can increase the QoS policer count to 64000. However, when you allocate more QoS policers, the BGP Flowspec policer count is reduced. For example, when you allocate 48000 policers for QoS, the default BGP Flowspec pool of 32000 is reduced to 16000. If you do not need or have limited use of Flowspec policers, you can enable this feature to increase QoS policers.

Some examples of increase in QoS policer allocation:

- If you allocate 48000 policers with 16 class-maps, then 2999 policers are allocated per core and 5998 policers per NPU.
- If you allocate 64000 policers with 16 class-maps, then 3999 policers are allocated per core and 7998 policers per NPU.

Guidelines and Limitations

- This feature is supported only on Cisco NCS 5700 series SE line cards.
- Enhanced QoS mode does not support more than 32000 policers. So, you cannot use this feature with Enhanced QoS mode configuration, which is enabled using the **hw-module profile stats qos-enhanced** command.
- The QoS policer scale numbers are only applicable when the **hw-module profile mdb l3max-se** or **hw-module profile mdb l2max-se** command is enabled. For more information on hardware MDB profiles, see [NCS5700 MDB Profile](#) in the *NCS 5500 System Setup and Software Installation Guide*.

Configure QoS Policer Pool Size

The following example shows how to allocate 48000 QoS policers:

```
RP/0/RP0/CPU0:router# config
RP/0/RP0/CPU0:router(config)# hw-module profile qos policer-scale 48000
RP/0/RP0/CPU0:router(config)# commit
```



Note After you configure, reload the router to apply the profile. Use the **admin hw-module reload location all** command to reload the router.

Traffic Policy Elements

A traffic policy contains three elements:

- Name
- Traffic class
- QoS policies

After choosing the traffic class that is used to classify traffic to the traffic policy, the user can enter the QoS features to be applied to the classified traffic.

The MQC does not necessarily require that the users associate only one traffic class to one traffic policy.

The order in which classes are configured in a policy map is important. The match rules of the classes are programmed into the TCAM in the order in which the classes are specified in a policy map. Therefore, if a packet can possibly match multiple classes, only the first matching class is returned and the corresponding policy is applied.

The router supports 32 classes per policy-map in the ingress direction and 8 classes per policy-map in the egress direction.

This table shows the supported class-actions on the router.

Supported Action Types	Direction supported on Interfaces
minimum-bandwidth	egress
bandwidth-remaining	egress
mark	(See Packet Marking , on page 20)
police	ingress
priority	egress (level 1 to level 7)
queue-limit	egress
shape	egress
wred	egress

WRED supports **default** and **discard-class** options; the only values to be passed to the discard-class being 0 and 1.

Create a Traffic Policy

The purpose of a traffic policy is to configure the QoS features that should be associated with the traffic that has been classified in a user-specified traffic class or classes.

To configure a traffic class, see [Create a Traffic Class, on page 9](#).

After you define a traffic policy with the **policy-map** command, you can attach it to one or more interfaces to specify the traffic policy for those interfaces by using the **service-policy** command in interface configuration mode. With dual policy support, you can have two traffic policies, one marking and one queuing attached at the output. See, [Attach a Traffic Policy to an Interface, on page 18](#).

Configuration Example

You have to accomplish the following to complete the traffic policy configuration:

1. Creating a policy map that can be attached to one or more interfaces to specify a service policy
2. Associating the traffic class with the traffic policy
3. Specifying the class-action(s) (see [Traffic Policy Elements, on page 15](#))

```
Router# configure
Router(config)# policy-map test-shape-1
Router(config-pmap)# class qos-1

/* Configure class-action ('shape' in this example).
Repeat as required, to specify other class-actions */
Router(config-pmap-c)# shape average percent 40
Router(config-pmap-c)# exit

/* Repeat class configuration as required, to specify other classes */

Router(config-pmap)# end-policy-map
Router(config)# commit
```

See, [Running Configuration, on page 19](#).

See, [Verification, on page 19](#).

Related Topics

- [Traffic Policy Elements, on page 15](#)
- [Traffic Class Elements, on page 8](#)

Associated Commands

- [bandwidth](#)
- [bandwidth remaining](#)
- [class](#)

- `police`
- `policy-map`
- `priority`
- `queue-limit`
- `service-policy`
- `set discard-class`
- `set dscp`
- `set mpls experimental`
- `set precedence`
- `set qos-group`
- `shape`

Scaling of Unique Ingress Policy Maps

Table 5: Feature History Table

Feature Name	Release Information	Feature Description
Scaling of Unique Ingress Policy Maps	Release 7.4.1	This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native mode.
Scaling of Unique Ingress Policy Maps	Release 7.3.1	With this feature, unique policy maps associated to the same template are shared in TCAM, thus enabling scaling of — or creating more number of — policy maps.

Traditionally, when unique policy maps were associated to the same template — that is, having the same match criteria and actions in the *same* order — each map was assigned a different TCAM entry. This resulted in inefficient TCAM entry management and also restricted the number of policy maps that could be created.

With this functionality, unique policy maps associated to the same template are shared in TCAM, thus enabling scaling of—in other words, creating more number of—policy maps. The other way to understand this functionality is that two policy maps with the same combination of criteria and actions use one template. This way, up to 250 templates are supported for association to policy map combinations.

As an example, consider the following policy maps (**`policy-map ncs_input1`** and **`policy-map ncs_input2`**) having the same class maps (**`class COS7_DEI0`** and **`class COS7_DEI1`**):

```
class-map match-all COS7_DEI0
  match cos 0
end-class-map
class-map match-all COS7_DEI1
  match cos 1
end-class-map
```

```

policy-map ncs_input1
  class COS7_DEI0
    set traffic class 1
    police rate 10 mbps
  !
  class COS7_DEI1
    set traffic class 2
    policer rate 20 mbps
  !

policy-map ncs_input2
  class COS7_DEI0
    set traffic class 1
    police rate 30 mbps
  !
  class COS7_DEI1
    set traffic class 2
    policer rate 40 mbps
  !

```

Earlier, when the policy maps were attached to interface, they used different TCAM entries, although the match criteria and actions were the same, except for the policer action.

With this functionality, both policy maps share the TCAM entry instead of selecting different entries, thus freeing up TCAM entries for more policy maps.

Limitations and Restrictions

- Policy Maps share TCAM entries only for the same match criteria and actions or template. However, the policer action can be different for the same template.
- For all unique policy maps the maximum number of templates supported is 250.

Attach a Traffic Policy to an Interface

After the traffic class and the traffic policy are created, you must attach the traffic policy to interface, and specify the direction in which the policy should be applied.



Note When a policy-map is applied to an interface, the transmission rate counter of each class is not accurate. This is because the transmission rate counter is calculated based on the exponential decay filter.

Configuration Example

You have to accomplish the following to attach a traffic policy to an interface:

1. Creating a traffic class and the associated rules that match packets to the class (see [#unique_27](#))
2. Creating a traffic policy that can be attached to one or more interfaces to specify a service policy (see [Create a Traffic Policy, on page 16](#))
3. Associating the traffic class with the traffic policy
4. Attaching the traffic policy to an interface, in the ingress or egress direction

```

Router# configure
Router(config)# interface HundredGigE 0/6/0/18
Router(config-int)# service-policy output test-shape-1
Router(config-int)# commit

```

Running Configuration

```

/* Class-map configuration */

class-map match-any traffic-class-1
  match traffic-class 1
end-class-map
!
- - -
- - -

/* Traffic policy configuration */
policy-map test-shape-1
  class traffic-class-1
    shape average percent 40
  !
  class class-default
  !
end-policy-map
!
- - -
- - -

/* Attaching traffic policy to an interface in egress direction */
interface HundredGigE0/6/0/18
  service-policy output test-shape-1
!

```

Verification

```

Router# show qos interface hundredGigE 0/6/0/18 output

```

NOTE:- Configured values are displayed within parentheses Interface HundredGigE0/6/0/18 ifh
 0x30001f8 -- output policy

```

NPU Id: 3
Total number of classes: 2
Interface Bandwidth: 100000000 kbps
VOQ Base: 11112
VOQ Stats Handle: 0x88430698
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)

```

```

-----
Level1 Class = qos-1
Egressq Queue ID = 11113 (LP queue)
Queue Max. BW. = 40329846 kbps (40 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 40000000 kbps
TailDrop Threshold = 50069504 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class = class-default
Egressq Queue ID = 11112 (Default LP queue)
Queue Max. BW. = 101803495 kbps (default)
Queue Min. BW. = 0 kbps (default)

```

Inverse Weight / Weight	=	1 / (BWR not configured)
Guaranteed service rate	=	50000000 kbps
TailDrop Threshold	=	62652416 bytes / 10 ms (default)
WRED not configured for this class		

Related Topics

- [Traffic Policy Elements, on page 15](#)
- [Traffic Class Elements, on page 8](#)

Associated Commands

- [service-policy](#)

Packet Marking



Note L2 packet marking is not supported on NC55-24DD and NC55-18DD-SE line cards for Cisco IOS XR Release 7.0.2.

The packet marking feature provides users with a means to differentiate packets based on the designated markings. The router supports egress packet marking. match on **discard-class** on egress, if configured, can be used for a marking policy only.

The router also supports L2 ingress marking.

For ingress marking:

Ingress traffic— For the ingress pop operation, re-marking the customer VLAN tag (CoS, DEI) is not supported.

Egress traffic— The ingress ‘pop VLAN’ is translated to a ‘push VLAN’ for the egress traffic, and (CoS, DEI) marking is supported for newly pushed VLAN tags. If two VLAN tags are pushed to the packet header at the egress side, both inner and outer VLAN tags are marked. For example:

1. rewrite ingress tag pop 1 symmetric
2. rewrite ingress tag pop 2 symmetric
3. rewrite ingress tag translate 2-to-1 dot1q/dot1ad <> symmetric

In case of pop action, the outer VLAN tag (CoS, DEI) is retained on NC57 line cards.

Single tag— When symmetrical pop 1 action is performed, the outer tag (CoS, DEI) is retained as the original frame.

Double tag— When symmetrical pop 2 action is performed, the outer tag (CoS, DEI) is retained as the original frame and inner tag (CoS, DEI) is set to 0/0.

Limitations

- The statistics and counters for the egress marking policy cannot be viewed on the router.
- QoS EXP matching for egress doesn’t work for Layer 2 VPN and Layer 3 VPN traffic flowing from:

- Cisco NCS 5700 series line cards at ingress to Cisco NCS 5500 series line cards at the egress
- and
- from Cisco NCS 5500 series line cards at ingress to Cisco NCS 5700 series line cards at egress.
 - For QoS EXP-Egress marking applied on a Layer 3 interface, there is a limit of 2 unique policy-maps per NPU. You can apply these policies to as many interfaces as your system resources allow. However, if you apply more than two unique policies, you may encounter unexpected failure.
 - For QoS egress marking (CoS, DEI) applied on a Layer 2 interface, there is a limit of 13 unique policy-maps per NPU. If you exceed this number, you may encounter unexpected failure.

Supported Packet Marking Operations

This table shows the supported packet marking operations.

Supported Mark Types	Range	Support for Unconditional Marking	Support for Conditional Marking
set dscp	0-63	ingress	No
set QoS-group	0-7	ingress	No
set traffic-class	0-7	ingress	No

Class-based Unconditional Packet Marking

Table 6: Feature History Table

Feature Name	Release Information	Feature Description
Class-based Unconditional Packet Marking	Release 7.4.1	This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.

The packet marking feature allows you to partition your network into multiple priority levels or classes of service, as follows:

- Use QoS unconditional packet marking to set the IP precedence or IP DSCP values for packets entering the network. Routers within your network can then use the newly marked IP precedence values to determine how the traffic should be treated.

On ingress direction, after matching the traffic based on either the IP Precedence or DSCP value, you can set it to a particular discard-class. Weighted random early detection (WRED), a congestion avoidance technique, thereby uses discard-class values to determine the probability that a packet is dropped.

If however, you set a discard-class of 3, the packet is dropped at ingress itself.

- Use QoS unconditional packet marking to assign MPLS packets to a QoS group. The router uses the QoS group to determine how to prioritize packets for transmission. To set the traffic class identifier on MPLS packets, use the **set traffic-class** command in policy map class configuration mode.



Note Setting the traffic class identifier does not automatically prioritize the packets for transmission. You must first configure an egress policy that uses the traffic class.



- Note**
- Unless otherwise indicated, the class-based unconditional packet marking for Layer 3 physical interfaces applies to bundle interfaces.
 - From IOS XR Release 7.2.1 onwards with NC57 line cards, propagation of PREC->COS marking happens by default on egress Layer 3 subinterfaces. This applies to single and double-tag L3 subinterfaces, and to NC57 line cards in Native mode.

Handling QoS for Locally Originated Packets

Packets that are generated and transmitted by a router are called Locally Originated Packets (LOPs). These are different from packets that pass through the router. Each device uses a default precedence value as determined by the device. The default value, used by Locally Originated Control Protocols (LOCPs) such as BGP, OSPF, CCM(CSM), and RSVP, is a precedence of 6 or Differentiated Services Codepoint (DSCP) of 48. Locally Originated Management Protocols (LOMPs) such as Telnet and SSH use a precedence value of 2 or DSCP of 16. SNMP uses a precedence value of 0. Some protocols such as BGP, RSVP, CFM, and LDP and the management protocols are capable of setting a specific precedence or DSCP value.



Note Bidirectional Forwarding Detection (BFD) uses a DSCP value of 48 (CS6) for single- and multi-hops sessions on NCS55xx.

The following applies to Traffic Class (TC) alignment:

- For releases before Release 7.6.1, locally generated control plane packets, such as IS-IS and BGP, are generated using traffic-class 6.
- From Release 7.6.1 onwards, locally generated control plane packets, such as IS-IS and BGP, are generated using traffic-class 7.
- Locally generated BFD over Bundle (IETF) packets, which are generated on the Network Processing Unit (NPU), are generated using traffic-class 7.
- When the BFD packets are offloaded to the hardware and generated on the NPU, the egress QoS policies are applied. These packets are classified along with the regular data plane traffic.

On the router, datapath packets and injected packets aren't differentiated if both their traffic classes share the same Virtual Output Queues (VOQs). Therefore, in the case of a congested VOQ, the LOCP packets are dropped. To avoid the LOCP packets drop, Cisco recommends that you have a different traffic class for data path traffic. Alternatively, you can also specify a higher bandwidth for traffic-class 7 (if ingress traffic rate is predictable).

Classifying traffic helps the router to recognize traffic as a certain type and mark that traffic. By marking traffic early on its travel, you can prevent excessive reclassification later. You can mark traffic at the protocol level as shown in the following examples:

Ethernet

The following configuration shows that the outbound Control Hub packets are marked with a precedence value of 2 and EXP of 2, instead of a precedence and EXP value of 6. The SSH packets have a precedence value of 3 instead of 2.

```

ethernet cfm
  mep domain FOO service FOOBAR mep-id 1
  cos 2
ssh server dscp 24

```

BGP

```
neighbor x.x.x.x dscp
```

MPLS LDP

```
mpls ldp signalling dscp
```

Telnet

```
telnet ipv4 dscp
```

SNMP

```
snmp-server ipv4 precedence/dscp
```

Syslog

```

logging ipv4 precedence/dscp
netflow
flow exporter-map TEST dscp

```

NTP

```

ntp ipv4 precedence/dscp
ssh client dscp 56
ssh server dscp 56

```



Note By default, the router marks the Precision Time Protocol (PTP) traffic as high priority. Therefore, the need to prioritize PTP traffic in the QoS configuration is not required.

All LOCPs originating on the RP or LC CPU have the discard priority set in the appended Buffer Header (BHDR). The discard priority ensures that the LOCPs are not dropped internally (under normal circumstances). Such LOCPs include non-IP (IS-IS and ARP) based control packets. The discard priority is not set for LOMPs. Therefore, such packets are treated as normal traffic, both in terms of classification and re-marking, and may be dropped under congestion conditions. Therefore, you must ensure that you do not inadvertently re-mark and drop such traffic.



Note By default, all LOCPs are assigned to traffic-class 7. Considering that LOCPs and LOMPs are generated by the RP, an Ingress QoS policy cannot be applied. Therefore, you must ensure that the egress QoS policy includes a class-map which matches traffic-class 7. By definition, the egress QoS policy matches all implicitly marked packets.

LOCPs are not subject to traffic policing, Weighted Random Early Detection (WRED), or Tail-drop queue-limit operation. The LOCP packets are not subject to WRED, even if the **max_th** value is being met. The tail-drop queue-limit must be hit before the LOCP packets are dropped.

All LOCPs with the discard priority set are by default put into an implicitly allocated high priority queue of each physical egress interface.

By default, all LOCPs that have the discard priority set are put into an implicitly allocated high priority queue of each physical egress interface.

When configuring QoS policies, you may attach a policy to the physical interface, which then references the sub-interfaces. Or, alternatively, you may attach QoS policies to the sub-interfaces directly. If you attach QoS policies to the sub-interfaces directly, the operator is prevented from attaching a QoS policy to the physical interface. LOCPs, including those being transmitted on a sub-interface, are always sent out on the default high-priority queue of the physical interface. The operator is therefore prevented from assigning any bandwidth to the physical interface, which could be reserved for use by LOCPs. During over-subscription, it may lead to a LOCPs drop and as a result, sessions may be terminated.

To prevent session termination, a minimum bandwidth of MIN (1% of interface BW, 10 mbps) is reserved for the default high-priority queue associated with the physical interface that has no QoS policy applied. If a QoS policy is applied to the physical interface, the minimum bandwidth for the default HP queue is controlled by the configured policy.

- Any QoS classification does not affect the queue-selection for LOCP.
- Irrespective of the QoS policy configured, non-IP LOP control packets are always sent to the high-priority queue. For example, ISIS and ARP

LOCPs can be mapped to a corresponding QoS group. The following example illustrates how this can be achieved:

```
control-plane
! local control-packets
copy precedence qos-group
```

The precedence value of the control packet is mapped to the respective QoS group number.

Protecting Locally Originated BFD Packets

For releases before Release 7.6.1, BFD packets are injected into traffic-class 6, with drop priority 0 (equivalent of discard-class 0). If transit traffic is also classified into traffic-class 7 and the associated egress queue is congested, BFD packets may be dropped.

From Release 7.6.1, BFD packets are injected into traffic-class 7, with drop priority 0 (equivalent of discard-class 0). If transit traffic is also classified into traffic-class 7 and the associated egress queue is congested, BFD packets may be dropped.

The recommendation is to configure transit traffic-class 7 in the ingress QoS policy with discard-class 1 or 2. You must then configure WRED in the egress QoS policy that drops these packets before dropping discard-class 0.



Note The default queue length is 16 ms.

Example

```

class-map match-any NC
  match traffic-class 6 (for releases before Release 7.6.1)
  match traffic-class 7 (from Release 7.6.1)
!
policy-map CORE-OUT
  class NC
    random-detect discard-class 1 7 ms 8 ms

```

Hardware Programming

RP/0/RP0/CPU0:R26-Spine2-5508#sh qos interface hu0/0/0/35 output

NOTE:- Configured values are displayed within parentheses

Interface HundredGigE0/0/0/35 ifh 0x1408 -- output policy

```

NPU Id: 5
Total number of classes: 2
Interface Bandwidth: 100000000 kbps
Policy Name: 7 wred-out
SPI Id: 0x0
VOQ Base: 1592
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----

```

```

Level1 Class = NC
Egressq Queue ID = 1598 (LP queue)
Queue Max. BW. = 20480000 kbps (20 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 20000000 kbps
Peak burst = 33600 bytes (default)
TailDrop Threshold = 40108032 bytes / 16 ms (default)

```

```

WRED profile for Discard_Class 1
WRED Min. Threshold = 17563648 bytes (7 ms)
WRED Max. Threshold = 20054016 bytes (8 ms)

```

```

Default RED profile
WRED Min. Threshold = 0 bytes (0 ms)
WRED Max. Threshold = 0 bytes (0 ms)

```

```

WRED ECN = Disabled

```

Prioritization of IS-IS and ARP Packets to Manage Transit Traffic

Table 7: Feature History Table

Feature Name	Release Information	Feature Description
Prioritization of IS-IS and ARP Packets to Manage Transit Traffic	Release 7.5.1	<p>This feature gives you the option to assign the highest priority to IS-IS and Address Resolution Protocol (ARP) packets in transit. This feature is disabled by default.</p> <p>The feature provides more flexibility in transit traffic management on a per-hop basis and also fine-tunes the traffic profile management for transit traffic.</p> <p>This feature introduces the hw-module profile qos arp-isis-priority-enable command.</p>

Overview of IS-IS and ARP Traffic Prioritization

Transit traffic refers to all traffic that enters an ingress interface, is compared against the forwarding table entries, and forwarded out an egress interface toward its destination. While the exact path of the transit path may not be of interest to the sender or receiver, you may still want some of the Integrated Intermediate System-to-Intermediate System (IS-IS) and Address Resolution Protocol (ARP) transit traffic to be managed and routed efficiently between specific source and destination addresses. You can now achieve higher levels of flexibility and fine-tune the traffic profile management for transit traffic by enabling the ability to assign the highest priority level to IS-IS and ARP traffic on Layer 2 networks.

This feature is useful if you manage environments such as data centers where you have complete end-to-end control over your network, and you want to avoid any drops in IS-IS and ARP traffic during congestion.

Cisco IOS XR Release 7.5.1 introduces the **hw-module profile qos arp-isis-priority-enable** command to enable prioritization of IS-IS and ARP traffic in transit on Layer 2 networks. Configuring this command assigns a priority level of TC 7 to transit traffic.



Note Assigning highest priority levels to IS-IS and ARP traffic lead to higher volumes of ARP traffic and flood high-priority queues, which may cause other traffic to drop. It's best to assign priority levels based on your network configuration and traffic volume.

Guidelines

- This feature is disabled by default. Configure the **hw-module profile qos arp-isis-priority-enable** command to enable prioritization of IS-IS and ARP traffic in transit on Layer 2 networks.
- This feature is supported on routers that have the Cisco NC57 line cards installed and operate in the native and compatibility mode.
- Configuring **hw-module profile qos arp-isis-priority-enable** assigns a priority level of TC 7 for ISIS and ARP traffic. When this feature is in its default state of disabled, the default priority level is TC 0.
- This feature doesn't allow you to assign TC values.
- Reload the line card for the **hw-module** command to be functional.

Enabling IS-IS and ARP Traffic Prioritization

To enable IS-IS and ARP traffic prioritization, configure the **hw-module profile qos arp-isis-priority-enable** command.

```
Router#config
Router(config)#hw-module profile qos arp-isis-priority-enable
Router(config)#commit
Router(config)#exit
Router# reload location <lc location>
```

Prioritization is based on IS-IS destination MAC address (01:80:c2:00:00:14 and 01:80:c2:00:00:15) and ARP ether type 0x080. When you configure the **hw-module profile qos arp-isis-priority-enable** command, priority level for IS-IS and ARP and traffic is set as TC 7.

Verification

The following example shows the verification command for NC 57 line cards. The assigned priority level is **TC 07**.

```
RP/0/RP0/CPU0:ios#show controllers fia diag 0 "pp vis pkttm" location 0/0/CPU0
```

```
Node ID: 0/0/CPU0
```

```
R/S/I: 0/0/0
```

Packet TM Information Core=0 Results			
Name	Value	Size	
		(bits)	
Fwd_Action_Destination	0c353e	21	
type	PORT_ID	6	
value	353e	15	
Tc	07	3	
Dp	00	2	
Int_Stat_Meter_Dp_Cmd	00	4	
Lag_Lb_Key	1772	16	
St_Vsq_Ptr	00	8	
Visibility	01	1	
System-Headers_Size	28	7	
NWK_Header_Truncate_Size	00	8	
NWK_Header_Append_Size_Ptr	02	8	
Counter_ID_0	N/A	0	
Counter_ID_1	N/A	0	
Counter_ID_2	N/A	0	
Counter_ID_3	N/A	0	
Counter_ID_4	N/A	0	
Counter_ID_5	N/A	0	
Counter_ID_6	N/A	0	
Counter_ID_7	00a000	20	
Counter_ID_8	N/A	0	
Counter_ID_9	N/A	0	
Meter_ID_0	N/A	0	
Meter_ID_1	N/A	0	
Meter_ID_2	N/A	0	
Ethernet_Meter_ID	02	3	
snif0_cmd	00	5	
snif0_code	0009	9	
snif0_qualifier	00	8	
snif1_cmd	00	5	
snif1_code	01e0	9	
snif1_qualifier	00	8	
snif2_cmd	00	5	
snif2_code	0000	9	
snif2_qualifier	00	8	

QoS Re-marking of IP Packets in Egress Direction

The router support the marking of IP DSCP bits of all IP packets to zero, in the egress direction. This feature helps to re-mark the priority of IP packets, which is mostly used in scenarios like IP over Ethernet over MPLS over GRE. This functionality is achieved using the ingress policy-map with **set dscp 0** option configured in class-default.

Configuration Example

```
Router# configure
Router(config)# policy-map ingress-set-dscp-zero-policy
Router(config-pmap)# class class-default
```

```
Router(config-pmap-c)# set dscp 0
Router(config-pmap-c)# end-policy-map
Router(config-pmap)# commit
```

Running Configuration

```
policy-map ingress-set-dscp-zero-policy
class class-default
  set dscp 0
!
end-policy-map
!
```

QoS Re-marking of Ethernet Packets in Egress Direction

The router supports Layer 2 marking of Ethernet packets in the egress direction.

QoS L2 Re-marking of Ethernet Packets in Egress Direction

The router supports Layer 2 marking of Ethernet packets in the egress direction.

To enable this feature, you must:

- Configure the policy maps for queuing and marking at the egress interface.
- Set traffic-class in the ingress and use **match traffic-class** in the egress for queuing.
- Ensure that the **set qos-group** command is configured in ingress policy and the corresponding **match qos-group** command is configured in the egress marking policy. If there is no corresponding QoS group, you will experience traffic failure.

The ingress ‘push VLAN’ is translated to ‘pop VLAN’ for the egress traffic. In this case, (CoS, DEI) re-marking is not supported for the VLAN tag. For example:

1. rewrite ingress tag push dot1q/dot1ad <> symmetric
2. rewrite ingress tag push dot1q/dot1ad <> second-dot1q <> symmetric
3. rewrite ingress tag translate 1-to-2 dot1q/dot1ad <> second-dot1q <> symmetric

Running Configuration

```
policy-map egress-marking
class qos1
  set cos 1
!
class qos2
  set cos 2
  set dei 1
!
class qos3
  set cos 3
!
class class-default
  set cos 7
!
```

```
end-policy-map
!
```

QoS L2 Re-Marking of Ethernet Packets on L3 Flows in Egress Direction

The router supports Layer 2 marking of Ethernet packets on Layer 3 flows in the egress direction.

To enable this feature, you must:

- Configure the policy maps for marking at the egress interface.
- Ensure that the **set qos-group** command is configured in ingress policy and the corresponding **match qos-group** command is configured in the egress marking policy. If there is no corresponding QoS group, you will experience traffic failure.

Restrictions

The following restrictions apply while configuring the Layer 2 marking of Ethernet packets on Layer 3 flows in the egress direction.

- **set discard-class** is not supported in ingress policy with peering mode.
- Egress marking statistics are not available.
- Layer 2 (802.1p) Egress marking is supported on Layer 3 flows for these types of traffic: IP-to-IP, IP-to-MPLS, and MPLS-to-IP traffic.
- Layer 2 marking of Ethernet packets on Layer 3 flows in the egress direction is supported only in the peering mode.

Running Configuration

Ingress Policy:

You must first set up the qos-group at ingress.

```
class-map match-any Class0
match mpls experimental topmost 0
match precedence routine
match dscp 0-7
end-class-map
class-map match-any Class1
match mpls experimental topmost 1
match precedence priority
match dscp 8-15
end-class-map
class-map match-any Class2
match mpls experimental topmost 2
match precedence immediate
match dscp 16-23
end-class-map
class-map match-any Class3
match mpls experimental topmost 3
match precedence flash
match dscp 24-31
end-class-map
class-map match-any Class4
match mpls experimental topmost 4
match precedence flash-override
match dscp 32-39
```

```

    end-class-map
class-map match-any Class5
  match mpls experimental topmost 5
  match precedence critical
  match dscp 40-47
end-class-map
class-map match-any Class6
  match mpls experimental topmost 6
  match precedence internet
  match dscp 48-55
end-class-map
class-map match-any Class7
  match mpls experimental topmost 7
  match precedence network
  match dscp 56-63
end-class-map
!

policy-map ncs_input
  class Class7
    set traffic-class 7
    set mpls experimental imposition 7
    set qos-group 7
  !
  class Class6
    set traffic-class 6
    set mpls experimental imposition 6
    set qos-group 6
  !
  class Class5
    set traffic-class 5
    set mpls experimental imposition 5
    set qos-group 5
  !
  class Class4
    set traffic-class 4
    set mpls experimental imposition 4
    set qos-group 4
  !
  class Class3
    set traffic-class 4
    set mpls experimental imposition 3
    set qos-group 3
  !
  class Class2
    set traffic-class 2
    set mpls experimental imposition 2
    set qos-group 2
  !
  class Class1
    set traffic-class 2
    set mpls experimental imposition 1
    set qos-group 1
  !
  class Class0
    set traffic-class 0
    set mpls experimental imposition 0
    set qos-group 0
  !
end-policy-map
!
```

Egress Policy:

At the egress, run these commands to mark the packets.


```
class-map match-any qos7
match qos-group 7
end-class-map
!
class-map match-any qos6
match qos-group 6
end-class-map
!
class-map match-any qos5
match qos-group 5
end-class-map
!
class-map match-any qos4
match qos-group 4
end-class-map
!
class-map match-any qos3
match qos-group 3
end-class-map
!
class-map match-any qos2
match qos-group 2
end-class-map
!
class-map match-any qos1
match qos-group 1
end-class-map
!

policy-map ncs_output
class qos7
set cos 7
set dei 1
!
class qos6
set cos 6
set dei 1
!
class qos5
set cos 5
set dei 1
!
class qos4
set cos 4
set dei 1
!
class qos3
set cos 3
set dei 1
!
class qos2
set cos 2
set dei 1
!
class qos1
set cos 1
set dei 1
!
end-policy-map
!
```

QoS L2 Re-Marking of Ethernet Packets on L3 Flows in Egress Direction on L3 sub-interfaces

Table 8: Feature History Table

Feature Name	Release Information	Feature Description
Cisco NC57 Native Mode: 802.1p marking	Release 7.4.1	This feature is now supported on routers that have the Cisco NC57 line cards installed and operate in the native mode.

The router supports Layer 2 marking of Ethernet packets on Layer 3 flows in the egress direction on L3 subinterfaces.

To enable this feature, you must:

- Configure the policy maps for marking at the egress interface.
- Ensure that the **set qos-group** command is configured in ingress policy and the corresponding **match qos-group** command is configured in the egress marking policy. If there is no corresponding QoS group, you experience traffic failure.

Restrictions

The following restrictions apply while configuring the Layer 2 marking of Ethernet packets on Layer 3 flows in the egress direction.

- **set discard-class** is not supported in ingress policy with peering mode.
- Egress marking statistics are not available.
- Layer 2 (CoS, DEI) Egress marking is supported on Layer 3 flows on L3 subinterfaces for these types of traffic: IP-to-IP, IP-to-MPLS, and MPLS-to-IP traffic.

Running Configuration

Ingress Policy:

You must first set up the qos-group at ingress. This is applicable only when you want to mark packets at the egress.

```
class-map match-all COS0_DEI0
  match cos 0
  match dei 0
end-class-map
class-map match-all COS0_DEI1
  match cos 0
  match dei 1
end-class-map
class-map match-all COS1_DEI0
  match cos 1
  match dei 0
end-class-map
class-map match-all COS1_DEI1
  match cos 1
  match dei 1
end-class-map
class-map match-all COS2_DEI0
```

```
match cos 2
match dei 0
end-class-map
class-map match-all COS2_DEI1
match cos 2
match dei 1
end-class-map
class-map match-all COS3_DEI0
match cos 3
match dei 0
end-class-map
class-map match-all COS3_DEI1
match cos 3
match dei 1
end-class-map
class-map match-all COS4_DEI0
match cos 4
match dei 0
end-class-map
class-map match-all COS4_DEI1
match cos 4
match dei 1
end-class-map
class-map match-all COS5_DEI0
match cos 5
match dei 0
end-class-map
class-map match-all COS5_DEI1
match cos 5
match dei 1
end-class-map
class-map match-all COS6_DEI0
match cos 6
match dei 0
end-class-map
class-map match-all COS6_DEI1
match cos 6
match dei 1
end-class-map
class-map match-all COS7_DEI0
match cos 7
match dei 0
end-class-map
class-map match-all COS7_DEI1
match cos 7
match dei 1
end-class-map

policy-map ncs_input
class COS7_DEI0
set qos-group 7
set discard-class 0
!
class COS7_DEI1
set qos-group 7
set discard-class 1
!
class COS6_DEI0
set qos-group 6
set discard-class 0
!
class COS6_DEI1
set qos-group 6
set discard-class 1
```

```

!
class COS5_DEI0
  set qos-group 5
  set discard-class 0
!
class COS5_DEI1
  set qos-group 5
  set discard-class 1
!
class COS4_DEI0
  set qos-group 4
  set discard-class 0
!
class COS4_DEI1
  set qos-group 4
  set discard-class 1
!
class COS3_DEI0
  set qos-group 3
  set discard-class 0
!
class COS3_DEI1
  set qos-group 3
  set discard-class 1
!
class COS2_DEI0
  set qos-group 2
  set discard-class 0
!
class COS2_DEI1
  set qos-group 2
  set discard-class 1
!
class COS1_DEI0
  set qos-group 1
  set discard-class 0
!
class COS1_DEI1
  set qos-group 1
  set discard-class 1
!
class COS0_DEI0
  set qos-group 0
  set discard-class 0
!
class COS0_DEI1
  set qos-group 0
  set discard-class 1
!

```

Egress Policy:

At the egress, run these commands to mark the packets.

```

class-map match-all qos7_dc0
match qos-group 7
match discard-class 0
end-class-map
!
class-map match-all qos7_dc1
match qos-group 7
match discard-class 1
end-class-map
!

```

```
class-map match-all qos6_dc0
match qos-group 6
match discard-class 0
end-class-map
!
class-map match-all qos6_dc1
match qos-group 6
match discard-class 1
end-class-map
!
class-map match-all qos5_dc0
match qos-group 5
match discard-class 0
end-class-map
!
class-map match-all qos5_dc1
match qos-group 5
match discard-class 1
end-class-map
!
class-map match-all qos4_dc0
match qos-group 4
match discard-class 0
end-class-map
!
class-map match-all qos4_dc1
match qos-group 4
match discard-class 1
end-class-map
!
class-map match-all qos3_dc0
match qos-group 3
match discard-class 0
end-class-map
!
class-map match-all qos3_dc1
match qos-group 3
match discard-class 1
end-class-map
!
class-map match-all qos2_dc0
match qos-group 2
match discard-class 0
end-class-map
!
class-map match-all qos2_dc1
match qos-group 2
match discard-class 1
end-class-map
!
class-map match-all qos1_dc0
match qos-group 1
match discard-class 0
end-class-map
!
class-map match-all qos1_dc1
match qos-group 1
match discard-class 1
end-class-map
!
class-map match-all qos0_dc0
match qos-group 0
match discard-class 0
end-class-map
```

```

!
class-map match-all qos0_dc1
match qos-group 0
match discard-class 1
end-class-map
!

policy-map ncs_output
class qos7_dc0
set cos 7
set dei 0
set mpls experimental imposition 7
!
class qos7_dc1
set cos 7
set dei 1
set mpls experimental imposition 7
!
class qos6_dc0
set cos 6
set dei 0
set mpls experimental imposition 6
!
class qos6_dc1
set cos 6
set dei 1
set mpls experimental imposition 6
!
class qos5_dc0
set cos 5
set dei 0
set mpls experimental imposition 5
!
class qos5_dc1
set cos 5
set dei 1
set mpls experimental imposition 5
!
class qos4_dc0
set cos 4
set dei 0
set mpls experimental imposition 4
!
class qos4_dc1
set cos 4
set dei 1
set mpls experimental imposition 4
!
class qos3_dc0
set cos 3
set dei 0
set mpls experimental imposition 3
!
class qos3_dc1
set cos 3
set dei 1
set mpls experimental imposition 3
!
class qos2_dc0
set cos 2
set dei 0
set mpls experimental imposition 2
!

```

```

class qos2_dc1
  set cos 2
  set dei 1
  set mpls experimental imposition 2
!
class qos1_dc0
  set cos 1
  set dei 0
  set mpls experimental imposition 1
!
class qos1_dc1
  set cos 1
  set dei 1
  set mpls experimental imposition 1
!
class qos0_dc0
  set cos 0
  set dei 0
  set mpls experimental imposition 0
!
class qos0_dc1
  set cos 0
  set dei 1
  set mpls experimental imposition 0
!
end-policy-map
!

```

Layer 2 Ingress QoS Matching for IPv4 and IPv6 Destination Addresses

Table 9: Feature History Table

Feature Name	Release Information	Feature Description
Layer 2 Ingress QoS Matching for IPv4 and IPv6 Destination Addresses	Release 7.5.1	<p>Using this feature, you can match class maps to IPv4 and IPv6 destination addresses on Layer 2 networks. The Layer 2 interface service policy has the relevant class maps, actioning them for ingress QoS operations.</p> <p>This feature provides you with an additional level of classification for aggregated customer traffic at your ingress, thus giving you granular control on traffic flows.</p> <p>This feature introduces the following commands:</p> <ul style="list-style-type: none"> • match destination-address • hw-module profile qos l2-match-dest-addr-v4v6

Overview

As a service provider, you provide Layer 2 connectivity for different classes of customer traffic across your network. With aggregated customer traffic arriving at your ingress, you need to provide differential treatment depending on specific destination addresses for the traffic. Such ability gives you granular control over traffic, allowing you to classify specific traffic flows depending on the type of services for which your customers have signed up.

You can match class maps to IPv4 and IPv6 destination addresses on Layer 2 networks to ensure such granular control. The interface service policy has the relevant class maps, actioning them for ingress QoS marking.

Guidelines and Limitations

- You can match up to 4 IPv4 and IPv6 addresses each in a class.
- For match on IPv6, only up to 64-bit prefix match is supported.
- The L2VPN traffic can be Virtual Private Wire Service (VPWS) or Virtual Private LAN Service (VPLS).
- Redundant and non-redundant pseudowires are supported.
- This feature isn't supported with egress ACL enabled.
- This feature isn't supported on NC57 line cards.
- Traffic classification for VLAN tags is supported as shown in the following table.

VLAN Tag Condition	IPv4 Addresses	IPv6 Addresses	Combination of IPv4 and IPv6 Addresses
With no VLAN tags	☑	☑	☑
With a single VLAN tag	☑	☑	☑
With a double VLAN tag	☑	☑	☑

Configure Layer 2 Ingress QoS Matching for IPv4 and IPv6 Destination Addresses

Perform the following steps to configure Layer 2 ingress QoS matching for IPv4 and IPv6 destination addresses. This example covers:

- **match-all** criteria for an IPv4 address and a Layer 2 classification (**match dscp**) in the same class map.



Note You can use the **match-all** criteria only when you want to match one specific destination address with other Layer 2 classification options such (CoS, DEI) or DSCP.

- **match-any** criteria for IPv4 and IPv6 addresses in the same class map.

1. Enable the ability to match class maps to IPv4 and IPv6 destination addresses on Layer 2 networks. Reload the router for the **hw-module** command to be functional.
2. Create a class map and specify **match-all** criteria for an IPv4 address and DSCP.

3. Create a class map and specify **match-any** criteria for IPv4 and IPv6 addresses.
4. Create a policy map and associate the class maps you created with the traffic policy and specify **class-action**.
5. Attach the policy map to the interface.

Configuration

```
/*Enable the ability to match class maps to IPv4 and IPv6 destination addresses on Layer 2
 networks*/
Router(config)#hw-module profile qos l2-match-dest-addr-v4v6
Router(config)#commit
Router#reload

/*Create a class map and specify match-all criteria for an IPv4 address and DSCP*/
Router(config)#class-map match-all ipv4_dst_cs1
Router(config-cmap)#match destination-address ipv4 192.168.1.4 255.255.255.255
Router(config-cmap)#match dscp cs1

/*Create a class map and specify match-any criteria for IPv4 and IPv6 addresses*/
Router(config-cmap)#class-map match-any V4_V6_MATCH
Router(config-cmap)#match destination-address ipv4 10.0.0.0 255.0.0.0
Router(config-cmap)#match destination-address ipv4 20.1.0.0 255.255.0.0
Router(config-cmap)#match destination-address ipv4 20.1.1.1 255.255.255.255
Router(config-cmap)#match destination-address ipv4 30.1.0.1 255.255.255.0
Router(config-cmap)#match destination-address ipv6 101:1:12::1/64
Router(config-cmap)#match destination-address ipv6 201:1:1::1/32
Router(config-cmap)#match destination-address ipv6 201:1:3::2/64
Router(config-cmap)#match destination-address ipv6 301:1:3::2/64
Router(config-cmap)#commit

/*Create a policy map, associate the class maps with the traffic policy; specify class-action:
 police rate, in this example*/
Router(config-cmap)#policy-map PMAP_L2_V4_V6_MATCH
Router(config-pmap)#class ipv4_dst_cs1
Router(config-pmap-c)#police rate 10 mbps
Router(config-pmap-c-police)#class V4_V6_MATCH
Router(config-pmap-c)#police rate 10 mbps
Router(config-pmap-c-police)#commit

/*Attach the policy map with class-actions that you set in the class maps*/
Router(config-pmap-c-police)#int Bundle-Ether100.2
Router(config-if)#service-policy input PMAP_L2_V4_V6_MATCH
Router(config-if)#commit
```

You have successfully configured Layer 2 ingress QoS matching for IPv4 and IPv6 destination addresses.

Running Configuration

```
qos l2-match-dest-addr-v4v6
!

class-map match-all ipv4_dst_cs1
 match destination-address ipv4 192.168.1.4 255.255.255.255
 match dscp cs1
!

class-map match-any V4_V6_MATCH
 match destination-address ipv4 10.0.0.0 255.0.0.0
 match destination-address ipv4 20.1.0.0 255.255.0.0
 match destination-address ipv4 20.1.1.1 255.255.255.255
```

```

match destination-address ipv4 30.1.0.1 255.255.255.0
match destination-address ipv6 101:1:12::1/64
match destination-address ipv6 201:1:1::1/32
match destination-address ipv6 201:1:3::2/64
match destination-address ipv6 301:1:3::2/64

!
!

policy-map PMAP_L2_V4_V6_MATCH
  class ipv4_dst_cs1
    police rate 10 mbps
    class V4_V6_MATCH
    police rate 10 mbps
!
!
```

Verification

To verify that the configuration was successful, run the **sh policy-map pmap-name** command for the policy map you created with all class maps associated. The output displays all the **match-any** and **match-all** configurations for IPv4 and IPv6 addresses.

```

Router#sh policy-map pmap-name PMAP_L2_V4_V6_MATCH detail
class-map match-all ipv4_dst_cs1
  match destination-address ipv4 192.168.1.4 255.255.255.255
  match dscp cs1
end-class-map

!
class-map match-any V4_V6_MATCH
  match destination-address ipv4 10.0.0.0 255.0.0.0
  match destination-address ipv4 20.1.0.0 255.255.0.0
  match destination-address ipv4 20.1.1.1 255.255.255.255
  match destination-address ipv4 30.1.0.1 255.255.255.0
  match destination-address ipv6 101:1:12::1/64
  match destination-address ipv6 201:1:1::1/32
  match destination-address ipv6 201:1:3::2/64
  match destination-address ipv6 301:1:3::2/64
end-class-map
!
policy-map PMAP_L2_V4_V6_MATCH
  class ipv4_dst_cs1
    police rate 10 mbps
  !
  !
  class V4_V6_MATCH
    police rate 10 mbps
  !
  !
  class class-default
  !
end-policy-map
!
!
```

```

Router#sh run interface bundle-ether 100
interface Bundle-Ether100
  service-policy input ipv4_dst_cs1
  ipv4 address 192.168.1.4 255.255.255.255
```

```
dscp cs1
service-policy input PMAP_L2_V4_V6_MATCH
ipv4 address 10.1.0.1 255.255.255.0
ipv6 address 10:1::1/96
!
```

Bundle Traffic Policies

A policy can be bound to bundles. When a policy is bound to a bundle, the same policy is programmed on every bundle member (port). For example, if there is a policer or shaper rate, the same rate is configured on every port. Traffic is scheduled to bundle members based on the load balancing algorithm.

Both ingress and egress traffic is supported. Percentage-based policies, absolute rate-based policies, and time-based policies are supported.



Note Egress marking is not supported on BVI interfaces.

For details, see [Configure QoS on Link Bundles, on page 148](#).

Shared Policy Instance

Table 10: Feature History Table

Feature Name	Release Information	Feature Description
Shared Policy Instance	Release 7.3.1	This feature allows you to share a single instance of QoS policy across multiple subinterfaces, allowing for aggregate shaping of the subinterfaces to one rate. The ability to facilitate queue consumption in this manner offers the advantage of saving on QoS and hardware resources, while ensuring that the specified rate is not exceeded.
Shared Policy Instance	Release 7.4.1	This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.

Traditionally, when services required by your end-customers mapped one-on-one to an interface, attaching the QoS policy-map directly to the interface was the way to meet customer SLAs. However, with increasing demand for triple play configurations—requiring the management of voice and video queues in addition to data queues—you may have several forwarding constructs. This scenario calls for the need to apply an aggregate QoS policy across interfaces to provide the necessary traffic.

After you create the traffic class and traffic policy, you can optionally use a shared policy instance to allocate a single set of QoS resources and share them across a group of subinterfaces.

With shared policy instance, you can share a single instance of a QoS policy across multiple subinterfaces, allowing for aggregate shaping, policing, and marking of the subinterfaces to one rate. All the subinterfaces that share the instance of a QoS policy must belong to the same main interface. The number of subinterfaces that share the QoS policy instance can range from 2 to the maximum number of subinterfaces on the main interface.

When a shared policy instance of a policy map is shared by several subinterfaces, QoS operations such as aggregate shaping, policing, and marking are applied for traffic on all the interfaces that use the same shared policy instance.

Traditionally, policies were bound to interfaces. However, different types of interfaces, such as Layer 2 and Layer 3, can use a single shared-policy-instance, which allows flexibility in the "attachment point" that binds the policy map.

As an example, consider the following policy configuration:

```
policy-map hqos_gold
  class class-default
    service-policy child_hqos_gold
    shape average 20 mbps
  !
end-policy-map
!
policy-map child_hqos_gold
  class voice
    priority level 1
    shape average 64 kbps
  !
  class video
    priority level 1
    shape average 4 mbps
  !
  class data
    bandwidth 5 mbps
  !
  class class-default
  !
end-policy-map
!

interface TenGigE 0/1/0/10.300 l2transport
  service-policy output hqos_gold shared-policy-instance hqos_gold_customer1
!
interface TenGigE 0/1/0/10.400 l2transport
  service-policy output hqos_gold shared-policy-instance hqos_gold_customer1
!
```

The keyword **shared-policy-instance** and the instance name **hqos_gold_customer1** identify the subinterfaces that share an aggregate SLA. These are shared on a physical main interface or a bundle member. In other words, in a mix of Layer 2 and Layer 3 subinterfaces in the same shared policy instance, both layers support classification criteria and action.

In the case of bundles, sharing is applicable within a bundle member and not the entire bundle. Depending on the traffic hashing, shared policy instance may or may not take effect across the subinterface under the bundle main interface.

All subinterfaces that share the same shared policy instance share resources as well. Hence, the **show policy-map** statistics values and **show qos** values for all the subinterfaces are the same.

Restrictions and Guidelines

The following restrictions and guidelines apply while configuring shared policy instance for a policy map.

- Subinterfaces that are part of the same shared policy must belong to the same main interface. In other words, subinterfaces of different main interfaces cannot be part of the same shared policy.
- There is no restriction on the number of unique shared policies across a system. However, the limit of maximum number of subinterfaces with QoS policies applies.
- There is no restriction on the number of unique shared policies per main interface, port, core, NPU, or line card.
- You cannot use the same shared policy name on the ingress and egress of the same subinterface.
- Shared policy instance is not supported with multi-policies. For example, on the egress, you cannot apply a marking policy and a queueing policy under a shared policy.
- A shared policy can include a combination of Layer 2 and Layer 3 subinterfaces.

Attaching a Shared Policy Instance to Multiple Subinterfaces

To attach a shared policy instance to multiple subinterfaces:

1. Enter interface configuration mode and configure a subinterface.
2. Attach a policy map to an input or output subinterface for it to be the service policy for that subinterface.

```
RP/0/RP0/CPU0:router(config)#interface HundredGigE0/3/0/0.1
RP/0/RP0/CPU0:router(config-subif)#service-policy output pm-out shared-policy-instance spi1
```

Running Configuration

```
interface HundredGigE0/3/0/0.1
service-policy output pm-out shared-policy-instance spi1
ipv4 address 20.0.0.1 255.255.255.0
encapsulation dot1q 1
!
```

Verification

The `show policy-map shared-policy-instance` command includes an option to display counters for the shared policy instance.



Note

- For bundle subinterfaces, use `RP` as the location keyword.
- For physical subinterfaces, use `LC` as the location keyword.

For example, for a physical interface:

```
RP/0/RP0/CPU0:ios#show policy-map shared-policy-instance spi1 output location 0/3/CPU0

Shared Policy Instance spi1 output: pm-out

Class cm-tc-1
```

```

Classification statistics          (packets/bytes)          (rate - kbps)
  Matched                        :      772637560/1143503679080      9622860
  Transmitted                    :      731260312/1082265352040      5052880
  Total Dropped                  :      41377248/61238327040       4569980
Queueing statistics
  Queue ID                      :      1433
  Taildropped(packets/bytes)     :      41377248/61238327040
Class class-default
  Classification statistics      (packets/bytes)          (rate - kbps)
    Matched                     :              0/0              0
    Transmitted                  :              0/0              0
    Total Dropped                :              0/0              0
  Queueing statistics
    Queue ID                    :      1432
    Taildropped(packets/bytes)   :      0/0
Policy Bag Stats time: 1604675533816 [Local Time: 11/06/20 15:12:13.816]

```

Use the **clear qos counters shared-policy-instance** command to clear counters for the shared policy instance.

**Note**

- For bundle subinterfaces, use RP as the location keyword.
- For physical subinterfaces, use LC as the location keyword.

For example, for a physical interface:

```
RP/0/RP0/CPU0:ios#clear qos counters shared-policy-instance spil output location 0/3/CPU0
```

The **show qos shared-policy-instance** command allows you to display the QoS hardware programming values.

**Note**

- For bundle subinterfaces, use RP as the location keyword.
- For physical subinterfaces, use LC as the location keyword.

For example, for a physical interface:

```

RP/0/RP0/CPU0:ios#show qos shared-policy-instance spil output location 0/3/CPU0
Fri Nov 6 15:21:44.200 UTC
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/3/0/0.1 ifh 0x60040c8 -- output policy
NPU Id:                                0
Total number of classes:                2
Interface Bandwidth:                    100000000 kbps
Policy Name:                            pm-out
SPI Id:                                0x3000001
VOQ Base:                               1432
Accounting Type:                        Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class                            = cm-tc-1
Egressq Queue ID                        = 1433 (LP queue)
Queue Max. BW.                          = 5118857 kbps (5 %)
Queue Min. BW.                          = 0 kbps (default)
Inverse Weight / Weight                  = 1 / (BWR not configured)
Guaranteed service rate                  = 5000000 kbps
Peak burst                              = 33600 bytes (default)
TailDrop Threshold                      = 6258688 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class                = class-default
Egressq Queue ID            = 1432 (Default LP queue)
Queue Max. BW.              = no max (default)
Queue Min. BW.              = 0 kbps (default)
Inverse Weight / Weight     = 1 / (BWR not configured)
Guaranteed service rate     = 50000000 kbps
Peak burst                  = 33600 bytes (default)
TailDrop Threshold          = 62652416 bytes / 10 ms (default)
WRED not configured for this class

```

Ingress Short-Pipe

When QoS traffic leaves an MPLS network, the MPLS label stack is removed on the penultimate ingress Label Switch Router (LSR), leaving an IPv4 or IPv6 packet to be forwarded. MPLS experimental bits (or EXP or pipe mode) carries out this disposition process and the packet is marked with a Differentiated Services Code Point (DSCP) or precedence value (also called DSCP or Precedence-based classification).

Usually, QoS traffic supports DSCP and precedence-based classifications only when there is no MPLS label in the packet. Using the ingress short-pipe feature, however, you can classify a packet that contains one MPLS label using the type-of-service (ToS) field of the IPv4 or IPv6 header. This classification method is called ingress short-pipe. To classify an IP packet this way, you must:

1. Create a child class map.
2. Specify a ToS value in the child class map.
3. Attach the child class map to a parent class map.
4. Create a policy map containing the parent class map.
5. Set any ingress action such as traffic class or QoS group. From Release 7.1.1 onwards, you can also set ingress action DSCP (or precedence value).

With the ingress short-pipe feature, you get an increased visibility into traffic packets. Plus, the feature also removes the limitation of classifying MPLS packets that come into IPv4 or IPv6 networks.

Restrictions and Other Important Points

Ensure that you read these points before you configure the ingress short-pipe feature.

- This feature isn't supported on:
 - NC57-24DD
 - NC57-18DD-SE
 - NC57-36H-SE
 - NC57-36H6D-S
 - NC57-MOD-S
 - NCS-57B1-6D24-SYS
 - NCS-57B1-5DSE-SYS

- NCS-57C3-MOD-SYS
- NCS-57D2-18DD-SYS
- This feature works only when there is one MPLS header in the traffic packet. If there are two or more MPLS headers, the ingress-short pipe feature fails. For example, in case of Explicit Null where there are two labels at the disposition, this feature will not work.
- You can carry out ingress classification using either the MPLS experimental bits (or EXP or pipe mode) classification OR the DSCP/precedence (or short-pipe) classification. Ensure that you do not mix the classification methods, else it may result in an unknown behavior, and the classification may not work at all.
- This feature is supported only on L3VPN, and not supported on L2VPN.
- This feature works for regular IPv4/IPv6 traffic, but will not work for IPv6 VPN Provider Edge over MPLS (6VPE).
- You can add only one child class map to a parent class map.
- This feature supports the invocation of short-pipe and legacy DSCP classification for the same parent class map.
- The child class map can contain only match precedence and match dscp commands.
- This feature is not supported in peering mode.

Configure Ingress Short-Pipe

This section details a sample configuration for the ingress short-pipe feature and another sample to configure classification for labeled and non-labeled packets under the same parent class.

Sample configuration to classify a packet that contains one MPLS label using the type-of-service (ToS) field of the IPv4 or IPv6 header (or the ingress short-pipe method):

```
class-map match-any in_pipe
  match mpls disposition class-map child_pipe
end-class-map
!
class-map match-any child_pipe
  match precedence 1
  match dscp ipv4 af11
end-class-map
!
class-map match-any ingress-business-high
  match dscp af21 af22
end-class-map

class-map match-any ingress-business-low
  match dscp af11 af12
end-class-map

policy-map ingress-classifier
  class in_pipe
    set traffic-class 5
    set dscp af31
  class ingress-business-high
    set traffic-class 4
  class ingress-business-low
```



```

set traffic-class 2
class class-default
set traffic-class 0
!

```



Note The **set dscp** option is available from Release 7.1.1 onwards.

You can configure classification for both labeled and non-labeled packets under the same parent class as in the following sample configuration. In this example, for MPLS labeled packets, DSCP configured under the child class is classified, while for non-labeled packets, DSCP/ToS configured in the **match dscp <value>** statement is classified.

DSCP value range is from 0 through 63. The range option is not supported. Up to 8 items per class are supported. Up to 64 **match dscp** values in total.

```

class-map match-any in_pipe
match mpls disposition class-map child_pipe (labeled case)
match dscp af11 (non-labeled case)
end-class-map
!
class-map match-any child_pipe
match precedence 1
match dscp ipv4 af11
end-class-map
!
class-map match-any ingress-business-high
match dscp af21 af22
end-class-map

class-map match-any ingress-business-low
match dscp af11 af12
end-class-map

policy-map ingress-classifier
class in_pipe
set traffic-class 5

class ingress-business-high
set traffic-class 4
class ingress-business-low
set traffic-class 2
class class-default
set traffic-class 0
!

```

Associated Commands

- match mpls disposition class-map

Selective Egress Policy-Based Queue Mapping

With selective egress policy-based queue mapping, you can combine traffic class (TC) maps in various permutations at the egress.



Note Modular chassis do not support this feature.

The primary aim of introducing the egress TC (traffic class) mapping is to classify the traffic in the ingress using a single policy and place the classified traffic into queues, by assigning the traffic classes. At the egress, you can support different grouping of TCs.

Based on different Service Level Agreements (SLAs) that each customer has signed up for, you can group some TCs into priority queues for real time (RT) traffic, other TCs into guaranteed bandwidth (BW) traffic, and the rest into best effort (BE) traffic delivery.

Let us consider an example where three customers have purchased these services, based on their requirements:

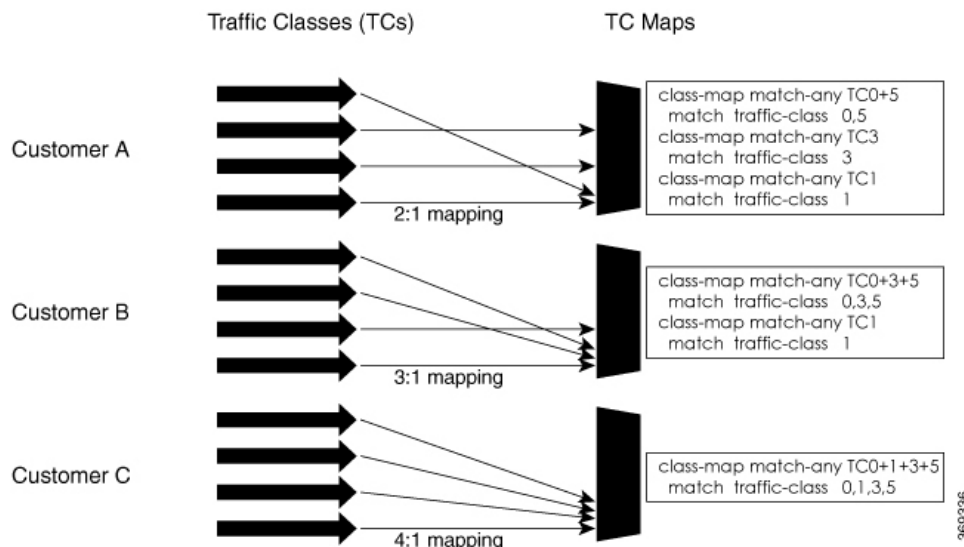
- Customer A - Requires RT traffic, reserved BW traffic and BE traffic delivery.
- Customer B – Requires reserved BW traffic and BE traffic delivery.
- Customer C – Needs only BE traffic delivery.

Using the selective egress policy-based queue mapping, you can create three profiles this way:

- Customer A – Priority queue RT traffic (TC1), Guaranteed BW traffic (TC3), Best effort traffic (TC0, TC5)
- Customer B – Guaranteed BW traffic (TC1), Best effort traffic (TC0, TC3, TC5)
- Customer C - Best effort traffic (TC0, TC1, TC3, TC5)

Using the egress TC-mapping, you can create three different profiles that you can use for each customer based on their SLAs with the provider.

Figure 1: Selective Egress Policy-Based Queue Mapping Helps Create Customer Profiles Based on Their SLAs



Restrictions and Other Important Points

Ensure that you read these points before you configure the selective egress policy-based queue-mapping feature.

- There can be only one TC (Traffic Class) mapped class to a PM (Policy Map).
- You cannot use a TC that you used in a mapped class, in a non-mapped class under the same PM.
- You can have a maximum of three unique TC mapped PMs or profiles per platform.
- Every TC mapped class must include **traffic-class 0** in the range values.
- The TC-mapping range is from 0 through 5.
- When a TC-mapped class is present in a PM, the class default becomes a dummy class. This means that the class default statistics and QoS values are not applicable.
- All the class default limitations apply to the TC-mapped class; for example, you cannot configure **priority** command under the TC mapped class.



Note A TC-mapped PM or profile is a PM that contains a TC-mapped class.

Example of a TC-mapped class:

```
match traffic-class 0 1 2 3
```

Example of a TC non-mapped class:

```
match traffic-class 1
```

Configure Selective Egress Policy-Based Queue Mapping

This section details a sample configuration for the selective egress policy-based queue-mapping feature and a use case to show how this feature works.

Sample configuration

```
class-map match-any <name>
  match traffic-class <value>
commit

policy-map tc_pmap
  class tc035
    shape average percent 1
  !
  class class-default
  !
end-policy-map
!
class-map match-any tc035
match traffic-class 0 3 5
end-class-map
!
```

Verification

Run the **show qos interface** and **show policy-map interface** commands.

When TC mapping class is present in a policy map, the class default does not have any values calculated.

show qos interface bundle-Ether 44 output sample

```
NOTE:- Configured values are displayed within parentheses
NPU Id: 0
Total number of classes: 3
Interface Bandwidth: 100000000 kbps
Policy Name: tc_pmap
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = tc1

Level1 Class = tc035

Level1 Class = class-default

Interface HundredGigE0/0/0/30 Ifh 0xf000208 (Member) -- output policy
NPU Id: 0
Total number of classes: 3
Interface Bandwidth: 100000000 kbps
Policy Name: tc_pmap
VOQ Base: 1264
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = tc1
Egressq Queue ID = 1265 (LP queue)
Queue Max. BW. = 10063882 kbps (10 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 10000000 kbps
TailDrop Threshold = 12517376 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class = tc035
Egressq Queue ID = 1264 (LP queue)
Queue Max. BW. = 1011732 kbps (1 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 1000000 kbps
TailDrop Threshold = 1253376 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class = class-default
Queue Max. BW. = no max (default)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 0 / (BWR not configured)
```

show policy-map interface bundle-Ether 44 output sample

Bundle-Ether44 output: tc_pmap

```
Class tc1
  Classification statistics (packets/bytes) (rate - kbps)
    Matched : 429444/53823648 0
    Transmitted : 429444/53823648 0
    Total Dropped : 0/0 0
  Queueing statistics
    Queue ID : None (Bundle)
    Taildropped(packets/bytes) : 0/0
Class tc035
```

```

Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                        :      1288331/161470820      0
  Transmitted                    :      1288331/161470820      0
  Total Dropped                  :              0/0              0
Queueing statistics
  Queue ID                      : None (Bundle)
  Taildropped(packets/bytes)     : 0/0
Class class-default
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                     :              0/0              0
    Transmitted                  :              0/0              0
    Total Dropped               :              0/0              0
  Queueing statistics
    Queue ID                    : None (Bundle)
    Taildropped(packets/bytes)   : 0/0
Policy Bag Stats time: 1557216940000 [Local Time: 05/07/19 08:15:40.000]
RP/0/RP0/CPU0:BB1#

```

Use Case

With the ingress traffic matching the same match criteria, you can group the egress traffic up to three unique TC mapped profiles. Using this feature, you can provide differentiated services to customers based on the SLAs they have signed up for.

In the example that follows, the ingress policy-map sets the ingress match criteria for the traffic class from 0 through 5. Based on the SLAs, you can group the TC values at the egress PM to deliver differentiated services.

After you group the TC values, you can apply specific egress actions under that class.

Ingress match:

```

class EXP1
  set traffic-class 1
!
class EXP2
  set traffic-class 2
!
class EXP3
  set traffic-class 3
!
class EXP4
  set traffic-class 4
!
class EXP5
  set traffic-class 5
!
class class-default
!
end-policy-map
!

```

Egress match:

Sample TC mapped class for policy-map PM1

```

class-map match-any TC2:1
match traffic-class 0 1
end-class-map

```

Sample TC mapped class for policy-map PM2

```

class-map match-any TC3:1
match traffic-class 0 1 2
end-class-map

```

Sample TC mapped class for policy-map PM3

```
class-map match-any TC6:1
match traffic-class 0 1 2 3 4 5
end-class-map
```

Configuring QoS Groups with an ACL

You can create QoS groups and configure ACLs to classify traffic into the groups based on a specified match condition. In this example, we match by the QoS group value (0-511).

Supported ACL Types

Your router supports the following ACL types.



Note If you configure QoS group with an unsupported ACL type, the system doesn't display any error message.

ACL Type	Example
IPv4 DSCP	permit ipv4 any any dscp af43
UDP DSCP	permit udp any any dscp af43
UDP Fragments IPv4	udp any any fragments - IPv4
UDP Fragments IPv6	udp any any fragments - IPv6
TCP Fragments IPv4	tcp any any fragments - IPv4
TCP Fragments IPv6	tcp any any fragments - IPv6
IPv4 DSCP Fragments	permit ipv4 any any dscp af43 fragments
UDP DSCP Fragments	permit udp any any dscp af43 fragments
UDP Host Fragments	permit udp host <sip> host <dip> fragments
TCP Host Fragments	permit tcp host <sip> host <dip> dscp af43 fragments
TCP DSCP	permit tcp <source network > <destination network> dscp af43
TCP Port based	permit tcp any any eq <port>
UDP Port based	permit udp any any eq <port>
TCP Flags	permit tcp host <sip> host <dip> established fin psh syn urg

Restrictions

- ACLs with fragment match are supported on systems with only NC57 line cards, also referred to as native mode.
- IPv6 ACLs with fragment match are supported only in **short** and **short-l2-qos** mode.

Prerequisites

Before you can configure QoS groups with an ACL, the QoS peering profile must be enabled on the router or the line card. After enabling QoS peering, the router or line card must be reloaded, as shown in the following configuration.

Enabling QoS Peering Profile on the Router

Enter the global configuration mode and enable the QoS peering profile for the router as shown:

```
RP/0/RP0/CPU0:router(config)# hw-module profile qos ingress-model peering
RP/0/RP0/CPU0:router(config)# exit
RP/0/RP0/CPU0:router# reload
```

Enabling QoS Peering Profile on the Line Card

Enter the global configuration mode and enable the QoS peering profile for the line card as shown:

```
RP/0/RP0/CPU0:router(config)# hw-module profile qos ingress-model peering location 0/0/CPU0
RP/0/RP0/CPU0:router(config)# exit
RP/0/RP0/CPU0:router# reload location 0/0/CPU0
```

Configuration

Use the following set of configuration statements to configure an ACL with QoS groups.

```
/*
  Enter the global configuration mode, and configure an ACL with the required QoS groups.
*/
RP/0/RP0/CPU0:router# configure
RP/0/RP0/CPU0:router(config)# ipv4 access-list qos-acl
RP/0/RP0/CPU0:router(config-ipv4-acl)# 10 permit ipv4 host 5.0.0.1 any set qos-group 1
RP/0/RP0/CPU0:router(config-ipv4-acl)# 11 permit ipv4 host 6.0.0.1 any set qos-group 2
RP/0/RP0/CPU0:router(config-ipv4-acl)# 12 permit ipv4 host 7.0.0.1 any set qos-group 3
RP/0/RP0/CPU0:router(config-ipv4-acl)# 13 deny ipv4 any any

/* Create a policy map with the required classes.
In this example, we also create a default class for traffic that does not belong to any of
the specified
classes. */
RP/0/RP0/CPU0:router(config)# policy-map qos-acl-map
RP/0/RP0/CPU0:router(config-pmap)# class qos1
RP/0/RP0/CPU0:router(config-pmap-c)# set dscp af43
RP/0/RP0/CPU0:router(config-pmap-c)# set traffic-class 2
RP/0/RP0/CPU0:router(config-pmap-c)# exit

RP/0/RP0/CPU0:router(config-pmap)# class qos2
RP/0/RP0/CPU0:router(config-pmap-c)# set precedence critical
RP/0/RP0/CPU0:router(config-pmap-c)# set traffic-class 7
RP/0/RP0/CPU0:router(config-pmap-c)# exit

RP/0/RP0/CPU0:router(config-pmap)# class qos3
RP/0/RP0/CPU0:router(config-pmap-c)# set precedence 2
```

```

RP/0/RP0/CPU0:router(config-pmap-c)# set traffic-class 2
RP/0/RP0/CPU0:router(config-pmap-c)# exit

RP/0/RP0/CPU0:router(config-pmap)# class qos4
RP/0/RP0/CPU0:router(config-pmap-c)# set traffic-class 4
RP/0/RP0/CPU0:router(config-pmap-c)# set dscp cs4
RP/0/RP0/CPU0:router(config-pmap-c)# exit

RP/0/RP0/CPU0:router(config-pmap)# class class-default
RP/0/RP0/CPU0:router(config-pmap-c)# police rate percent 20
RP/0/RP0/CPU0:router(config-pmap-c-police)# exit

/* Create the class maps for specifying the match conditions. */
RP/0/RP0/CPU0:router(config)# class-map match-any qos1
RP/0/RP0/CPU0:router(config-cmap)# match qos-group 1
RP/0/RP0/CPU0:router(config-cmap)# end-class-map

RP/0/RP0/CPU0:router(config)# class-map match-any qos2
RP/0/RP0/CPU0:router(config-cmap)# match qos-group 2
RP/0/RP0/CPU0:router(config-cmap)# end-class-map

RP/0/RP0/CPU0:router(config)# class-map match-any qos3
RP/0/RP0/CPU0:router(config-cmap)# match qos-group 3
RP/0/RP0/CPU0:router(config-cmap)# end-class-map

RP/0/RP0/CPU0:router(config)# class-map match-any qos4
RP/0/RP0/CPU0:router(config-cmap)# match qos-group 4
RP/0/RP0/CPU0:router(config-cmap)# end-class-map

/* Apply the access list and the QoS map to the Gigabit interface, and commit your
configuration. */
RP/0/RP0/CPU0:router(config)# interface TenGigE0/0/0/1
RP/0/RP0/CPU0:router(config-if)# ipv4 address 12.0.0.1/24
RP/0/RP0/CPU0:router(config-if)# no shut
RP/0/RP0/CPU0:router(config-if)# service-policy input qos-acl-map
RP/0/RP0/CPU0:router(config-if)# ipv4 access-group qos-acl ingress compress level 3

RP/0/RP0/CPU0:router(config-if)# commit
Tue Mar 28 10:23:34.106 IST

RP/0/0/CPU0:Mar 28 10:37:48.570 : ifmgr[397]: %PKT_INFRA-LINK-3-UPDOWN : Interface
TenGigE0/0/0/1, changed state to Down
RP/0/0/CPU0:Mar 28 10:37:48.608 : ifmgr[397]: %PKT_INFRA-LINK-3-UPDOWN : Interface
TenGigE0/0/0/1, changed state to Up

RP/0/RP0/CPU0:router(config-if)# exit

```

Running Configuration

Confirm your configuration.

```

RP/0/RP0/CPU0:router(config)# show run
Tue Mar 28 10:37:55.737 IST

Building configuration...
!! IOS XR Configuration 0.0.0

ipv4 access-list qos-acl
10 permit ipv4 host 5.0.1.1 any set qos-group 1
11 permit ipv4 host 6.0.1.1 any set qos-group 2

```



```

12 permit ipv4 host 7.0.1.1 any set qos-group 3
13 deny ipv4 any any

class-map match-any qos1
match qos-group 1
end-class-map
!
class-map match-any qos2
match qos-group 2
end-class-map
!
class-map match-any qos3
match qos-group 3
end-class-map
!
class-map match-any qos4
match qos-group 4
end-class-map
!

policy-map qos-acl-map
class qos1
    set dscp af43
    set traffic-class 2
!
class qos2
    set precedence critical
    set traffic-class 7
!
class qos3
    set precedence 2
    set traffic-class 2
!
class qos4
    set traffic-class 4
    set dscp cs4
!
class class-default
    police rate percent 20
!
!
end-policy-map
!

interface TenGigE0/0/0/1
service-policy input qos-acl-map
ipv4 address 12.0.0.1 255.255.255.0
ipv4 access-group qos-acl ingress compress level 3

!

```

You have successfully configured an ACL with QoS groups.

Configuring an ACL with Fragment Match

Usually, IP ACLs process non-fragmented packets and the first fragments of a packet using permit and deny actions. These packets may have Layer 3 and 4 information that the ACLs match for a permit or deny action. By default, however, ACLs permit noninitial fragments. This is because noninitial fragments do not contain Layer 4 information, and the ACL logic extrapolates that if the Layer 3 information matches, then the Layer

4 information, had it been available, would also match. This could lead to potential security issues with users with malicious intent using the noninitial fragments to launch denial of service (DoS) attacks.

With this feature, you now have more granular control over noninitial IP fragments of a packet. The fragment has all the information the system requires to filter, so the ACL entry applies to the fragments of a packet. Before configuring this functionality to set QoS policies for noninitial fragment packets, ensure that you configure the fragment offset value to a non-zero value. To know more about fragment offset, see the *IP Addresses and Services Configuration Guide for Cisco NCS 5500 Series Routers* and [fragment-offset](#). Also, see *User Datagram Protocol (UDP)* under [deny \(IPv4\)](#) and [deny \(IPv6\)](#).

Restrictions and Guidelines

The following restrictions and guidelines apply while configuring an ACL with fragment match.

- Until Cisco IOS XR Release 7.8.1, this feature is supported on systems with NC57 line cards running in native mode. You can configure IPv4 and IPv6 ACLs with fragment match on these systems.
- To enable IPv6 fragment classification support, configure **hw-module profile qos ipv6 short-l2qos-enable** or **hw-module profile qos ipv6 short**.
- IPv6 fragmentation is supported with only one Extension Header (EH).

Configuring an ACL with Fragment Match

For IPv4 and IPv6 ACLs

To configure an ACL with fragment match, you must:

1. Create IPv4 and IPv6 ACLs with fragment match.
2. Create two class maps, one for IPv4 and IPv6, and attach the respective ACLs to the class maps.
3. Create a policy map with these two class maps and set action.



Note You can configure IPv4 and IPv6 ACLs with fragment match only on systems with NC57 line cards running in native mode.

```
RP/0/RP0/CPU0:router(config)#ipv4 access-list v4_ace
RP/0/RP0/CPU0:routerios(config-ipv4-acl)# permit ipv4 any any fragments
RP/0/RP0/CPU0:router(config)#ipv6 access-list v6_ace
RP/0/RP0/CPU0:routerios(config-ipv6-acl)# permit ipv6 any any fragments
RP/0/RP0/CPU0:router(config)#class-map match-any v4_class
RP/0/RP0/CPU0:router(config-cmap)#match access-group ipv4 v4_ace
RP/0/RP0/CPU0:router(config)#class-map match-any v6_class
RP/0/RP0/CPU0:router(config-cmap)#match access-group ipv6 v6_ace
RP/0/RP0/CPU0:router(config)#policy-map frag_policy
RP/0/RP0/CPU0:router(config-pmap)#class v4_class
RP/0/RP0/CPU0:router(config-pmap-c)#set traffic-class 3
RP/0/RP0/CPU0:router(config-pmap)#class v6_class
RP/0/RP0/CPU0:router(config-pmap-c)#set traffic-class 4
RP/0/RP0/CPU0:router(config-pmap-c)#commit
RP/0/RP0/CPU0:router(config)# exit
```

Running Configuration

```

ipv4 access-list v4_ace
permit ipv4 any any fragments
exit
!
ipv6 access-list v6_ace
permit ipv6 any any fragments
exit
!
class-map match-any v4_class
match access-group ipv4 v4_ace
exit
!
class-map match-any v6_class
match access-group ipv6 v6_ace
exit
!

policy-map frag_policy
class v4_class
set traffic-class 3
exit
!
class v6_class
set traffic-class 4
commit

```

Verification

```

RP/0/RP0/CPU0:ios#show policy-map pmap-name v4_policy detail
ipv4 access-list v4_ace
10 permit udp any any fragments
20 permit ipv4 any any fragments

ipv6 access-list v6_ace
10 permit ipv6 any any fragments
20 permit udp any any fragments

class-map match-any v4_class
match access-group ipv4 v4_ace
end-class-map
!
class-map match-any v6_class
match access-group ipv6 v6_ace
end-class-map
!
policy-map v4_policy
class v4_class
set dscp af41
!
class v6_class
set traffic-class 3
!
class class-default
!
end-policy-map

```

Restrictions

Refer to the following table for Ingress QoS Scale limitation.

Table 11: Ingress QoS Scale Limitation

QoS Mode	Class-Map Size	Maximum number of Interfaces with Ingress QoS Applied	
		Per Core	Per NPU
Normal	4	1023	2046
Normal	8	511	1022
Normal	16	255	510
Normal	32	127	254
Enhanced	4	871	1742
Enhanced	8	435	870
Enhanced	16	217	434
Enhanced	32	108	216



Note If you apply an ingress policy map to a bundle that has bundle members only from a single core of an NPU, the QoS resources are consumed on both cores of that NPU.

Example: For Default Configuration, which is Normal (2 counter mode) QoS Mode & 32 Class Map-Size, you can configure 127 interfaces with Ingress Policy per core.

Other restrictions to follow:

- If you have a **set traffic class** statement explicitly configured in ingress service policy, it is mandatory to have a corresponding **match traffic class** on egress for the traffic to be correctly matched and the stats to be accounted in **show policy-map interface <> output** command. To match the ingress traffic to egress class-default, traffic class should be set to 0 on ingress.
- If you have a **set traffic class that is configured** in ingress service policy, and no corresponding **match traffic class** on egress, the traffic will not go to class default and the stats for this traffic flow will not be seen in **show policy-map interface <> output** command.
- If you do not have any **set traffic class** statement in ingress, then traffic will hit the default-class on egress.
- If you have a **set discard-class** statement configured in ingress service policy, it is mandatory to have a corresponding **match discard-class** on egress for the traffic to be correctly matched and the stats to be accounted in **show policy-map interface <> output** command.
- If you have a **set discard-class** statement configured in ingress service policy and do not have a corresponding **match discard-class** on egress, the traffic will not hit the class-default and the stats for this flow will not be accounted in **show policy-map interface <> output** command.
- The system does not support class-map size on peering mode.

- Even if you have an egress policy that has the drop action configured, the transmitted counter stats still shows an increment.
- Depending on the packet size, the traffic shaped value for low shaper rates, such as 10mbps, have greater deviation than 5% of tolerance from the shaper value. For higher shaper rates, the deviation is within the limit of 5% of tolerance from the shaper value for all packet sizes.

Restrictions for Peering QoS Profile

- **explicit set discard-class** statement is not supported.
- This feature is supported only on L3 interfaces and is limited to 1000 L3 interfaces per system.
- **set mpls exp topmost** statement is not supported within QoS in peering mode.
- **access group** statement is not supported.
- (Only in Release 6.2.x and Release 6.3.x) **set mpls exp imposition** statement is not supported on ingress interface.
- 2-Level ingress policer is not supported.
- (From Release 6.5.x) Egress H-QOS with peering profile support is enabled, but ingress H-QOS with peering profile is not supported.
- Depending on the packet size, the traffic shaped value for low shaper rates, such as 10mbps, have greater deviation than 5% of tolerance from the shaper value. For higher shaper rates, the deviation is within the limit of 5% of tolerance from the shaper value for all packet sizes.
- (From Release 7.2.1) On NC57 line cards, QoS ingress peering profile is not supported. This restriction is applicable for systems operating in compatibility mode and native mode.

Restrictions for QoS on BVI

- The system does not support the egress policy on Bridge-Group Virtual Interface (BVI), but BVI (CoS, DEI) marking is supported by applying the policy to its corresponding Layer 2 interface, which is part of the same bridge domain.
- If you apply L3 ingress QoS policy on L2 interface, which is a part of the same bridge-domain as BVI, the classification might not work if packets are destined to the BVI MAC address.
- If a QoS policy is attached to BVI, the policy is inherited by the L2 interfaces, which are part of the same bridge-domain. Hence, any other policy cannot be applied on the L2 interfaces. Similarly, if a QoS policy is attached to any of the L2 interfaces, any QoS policy cannot be applied on the BVI, which is part of the same bridge-domain.

Restrictions for Egress Drop Action

- A maximum of 8 interfaces can have the drop action configured and a maximum of 8 classes in any single policy can have the drop action.
- A drop action in any particular class cannot be combined with other actions.
- Drop action in a policy applied on the main interface is not inherited onto sub-interfaces.

- Match condition for drop action PM can only be based on qos-group, discard class based match is not supported.

In-Place Policy Modification

The In-Place policy modification feature allows you to modify a QoS policy even when the QoS policy is attached to one or more interfaces. A modified policy is subjected to the same checks that a new policy is subject to when it is bound to an interface. If the policy-modification is successful, the modified policy takes effect on all the interfaces to which the policy is attached. However, if the policy modification fails on any one of the interfaces, an automatic rollback is initiated to ensure that the pre-modification policy is in effect on all the interfaces.

You can also modify any class map used in the policy map. The changes made to the class map take effect on all the interfaces to which the policy is attached.



Note

- The QoS statistics for the policy that is attached to an interface are lost (reset to 0) when the policy is modified.
- When a QoS policy attached to an interface is modified, there might not be any policy in effect on the interfaces in which the modified policy is used for a short period of time.
- The system does not support the show policy-map statistics for marking policies.
- An in-place modification of an ACL does not reset the policy-map statistics counter.



Note

- For QOS EXP-Egress marking applied on a Layer 3 interface, there is a limit of 2 unique policy-maps per NPU. When the maximum limit for policy-maps is reached and you try to modify a policy-map which is shared between different interfaces, you may get an error.
- For QOS egress marking (CoS, DEI) applied on a Layer 2 interface, there is a limit of 13 unique policy-maps per NPU. When the maximum limit for policy-maps is reached and you try to modify a policy-map which is shared between different interfaces, you may get an error.

Verification


If unrecoverable errors occur during in-place policy modification, the policy is put into an inconsistent state on target interfaces. No new configuration is possible until the configuration session is unblocked. It is recommended to remove the policy from the interface, check the modified policy and then re-apply accordingly.

References for Modular QoS Service Packet Classification

Specification of the CoS for a Packet with IP Precedence

Use of IP precedence allows you to specify the CoS for a packet. You can create differentiated service by setting precedence levels on incoming traffic and using them in combination with the QoS queuing features. So that, each subsequent network element can provide service based on the determined policy. IP precedence is usually deployed as close to the edge of the network or administrative domain as possible. This allows the rest of the core or backbone to implement QoS based on precedence.

Figure 2: IPv4 Packet Type of Service Field



You can use the three precedence bits in the type-of-service (ToS) field of the IPv4 header for this purpose. Using the ToS bits, you can define up to eight classes of service. Other features configured throughout the network can then use these bits to determine how to treat the packet in regard to the ToS to grant it. These other QoS features can assign appropriate traffic-handling policies, including congestion management strategy and bandwidth allocation. For example, queuing features such as LLQ can use the IP precedence setting of the packet to prioritize traffic.

IP Precedence Bits Used to Classify Packets

Use the three IP precedence bits in the ToS field of the IP header to specify the CoS assignment for each packet. You can partition traffic into a maximum of eight classes and then use policy maps to define network policies in terms of congestion handling and bandwidth allocation for each class.

Each precedence corresponds to a name. IP precedence bit settings 6 and 7 are reserved for network control information, such as routing updates. These names are defined in RFC 791.

IP Precedence Value Settings

By default, the routers leave the IP precedence value untouched. This preserves the precedence value set in the header and allows all internal network devices to provide service based on the IP precedence setting. This policy follows the standard approach stipulating that network traffic should be sorted into various types of service at the edge of the network and that those types of service should be implemented in the core of the network. Routers in the core of the network can then use the precedence bits to determine the order of transmission, the likelihood of packet drop, and so on.

Because traffic coming into your network can have the precedence set by outside devices, we recommend that you reset the precedence for all traffic entering your network. By controlling IP precedence settings, you prohibit users that have already set the IP precedence from acquiring better service for their traffic simply by setting a high precedence for all of their packets.

The class-based unconditional packet marking and LLQ features can use the IP precedence bits.

IP Precedence Compared to IP DSCP Marking

If you need to mark packets in your network and all your devices support IP DSCP marking, use the IP DSCP marking to mark your packets because the IP DSCP markings provide more unconditional packet marking options. If marking by IP DSCP is undesirable, however, or if you are unsure if the devices in your network support IP DSCP values, use the IP precedence value to mark your packets. The IP precedence value is likely to be supported by all devices in the network.

You can set up to 8 different IP precedence markings and 64 different IP DSCP markings.

Conditional Marking of MPLS Experimental bits for L3VPN Traffic

The conditional marking of MPLS experimental bits is achieved for Layer 3 Virtual Private Network (L3VPN) traffic by applying a combination of ingress and egress policy-maps on the Provider Edge (PE) router. In the ingress policy-map, the qos-group or discard-class is set either based on the result of the policing action or implicitly. The egress policy-map matches on qos-group or discard-class and sets the mpls experiment bits to the corresponding value.

This feature is supported on both IPv4 and IPv6 traffic in the L3VPN network. Conditional marking can be used to mark the MPLS experimental bits differently for in-contract and out-of-contract packets. In-contract packets are the confirmed packets with the color green and discard-class set to 0. Out-of-contract packets are the packets which have exceeded the limit and have the color yellow and discard-class set to 1.

Conditional marking of MPLS experimental bits for L3VPN traffic is supported on both physical and bundle main interfaces as well as sub-interfaces.

Restrictions for Conditional Marking of MPLS Experimental bits on L3VPN

1. In the case of two PE routers connected back-to-back and the only label that the traffic between the routers have is the BGP label, then the explicit null label should be configured.
2. A maximum of three policy-maps which perform conditional marking of MPLS experimental bits can be configured per Network Processor Unit (NPU) of the Cisco NCS 5500 Series Routers.
3. In the ingress policy-map if qos-group is being set for the incoming traffic packets, then setting of dscp and mpls experimental bits will not work.
4. Both the ingress and egress policy-maps must be applied in order to attain the expected behaviour. If either one of them is not applied then it may lead to undefined behaviour.
5. If the egress policy-map does not match on qos-group or discard-class and set the mpls experiment bits to the required value, then the mpls experimental bits will be set to a value of zero, by default.

QoS DSCP Preservation

Table 12: Feature History Table

Feature Name	Release Information	Feature Description
QoS DSCP Preservation	Release 7.2.1	This feature is now also supported on routers that have the Cisco NC57 line cards installed and operate in the native mode.

The DSCP value of the packet is preserved only for L3VPN networks when the packets are destined to the directly connected routes on the PE routers. To preserve the DSCP value for packets to destinations beyond the egress PE routers for L3VPN, you should use the **label mode per-vrf** command under the VRF at the PE routers. Similarly for MPLS networks, the default behaviour is uniform mode where the penultimate hop copies the MPLS EXP bit to IP DSCP and IP DSCP value is not preserved beyond the egress PE router. To preserve the IP DSCP value, you should use the **mpls ip-ttl-propagate disable** command at the penultimate hop.

DSCP value preservation is not supported on NC55-24DD and NC55-18DD-SE line cards for Cisco IOS XR Release 7.0.2.

QoS IP DSCP Preservation for SR-TE

Table 13: Feature History Table

Feature Name	Release Information	Feature Description
QoS IP DSCP Preservation for SR-TE	Release 7.7.1	<p>In terms of preserving IP DSCP markings, this release covers two scenarios for SR-TE traffic:</p> <ul style="list-style-type: none"> • For two or less than two topmost or imposition labels: when you set the MPLS experimental bits (EXP) values (also called Traffic Class values), the IP DSCP markings are now preserved by default in the ingress policies when the MPLS labels are pushed into the packet. • For more than three imposition labels: you must enable this functionality to preserve IP DSCP markings. <p>With preservation, traffic with IP packets with DSCP marking for priority, flows as intended and there's no drop in traffic because of incorrect or missing labels.</p> <p>In previous releases, irrespective of the number of MPLS labels, when the EXP values were copied into the packet header during imposition, even the IP DSCP markings were modified. This modification resulted in traffic drops at the next-hop routers in SR-TE tunnels.</p> <p>This feature introduces the hw-module profile mpls-ext-dscp-preserve v4uc-enable command.</p>

A typical ingress scenario for SR-TE traffic involves:

1. Specifying DSCP values as match criteria for a class map.
2. Identifying specific three-bit experimental bits (EXP) field values (also called Traffic Class or TC values, as per [RFC 5462](#)) in the topmost MPLS label as match criteria for the same class map.

When the MPLS labels are pushed (also called imposition) into the packet, the EXP values are copied into the EXP field in the packet header. At this stage, for SR-TE traffic, even the DSCP values that you set initially are modified, which could cause traffic drops at the next-hop router.

It is this behavior that the QoS DSCP preservation for SR-TE functionality mitigates. To enable this functionality, configure the **hw-module profile mpls-ext-dscp-preserve v4uc-enable** command. This configuration ensures that the IP DSCP values are preserved for SR-TE traffic.



Note Ensure that you restart the router for the configuration to take effect.

Here's a quick look at what the QoS DSCP preservation for SR-TE functionality achieves.

Table 14: QoS DSCP Preservation for SR-TE and TI-LFA Scenarios

Scenario	What was happening in previous releases?	What's happening now?
SR-TE cases	Irrespective of the number of MPLS labels (1, 2, ... 12) for imposition, the push action modified IP DSCP values.	<ul style="list-style-type: none"> IP DSCP values are preserved by default for two or less than two MPLS labels. IP DSCP values are preserved for three or more labels only when you configure the hw-module profile mpls-ext-dscp-preserve v4uc-enable command.
SR Topology-Independent Loop-Free Alternate (TI-LFA) cases	For three MPLS labels for imposition, the push action modified IP DSCP values.	Irrespective of the number of MPLS labels, IP DSCP values are preserved when you configure the hw-module profile mpls-ext-dscp-preserve v4uc-enable command.

Guidelines and Limitations for QoS IP DSCP Preservation for SR-TE

- Disabling the propagation of IP Time-To-Live (TTL) to and from the MPLS header (using the **mpls ip-ttl-propagate** command) doesn't help preserve IP DSCP values for SR-TE traffic. This is because the IP TTL acts on the egress, while the DSCP preservation for SR-TE is for the ingress traffic.
- When you use the **on-demand steering** command to disable labeled services for SR-TE policies, your router supports 12 MPLS transport labels by default. With the QoS DSCP preservation for SR-TE functionality enabled, this default value changes to ten MPLS transport labels.
- By default, SR-TE (with steering services enabled) supports 2 service labels and 9 transport labels. This support changes to 2 service labels and 8 transport labels when you configure **hw-module profile mpls-ext-dscp-preserve v4uc-enable**.
- This functionality doesn't work for SR-TE traffic associated with Binding Single-ID (BSID) policy paths. This is because packets with BSID labels get terminated, and new SR-TE labels get pushed, modifying the IP DSCP values. Thus, for BSID traffic (labeled traffic), disabling the propagation of IP TTL (using

[mpls ip-ttl-propagate](#)) or configuring the **hw-module profile mpls-ext-dscp-preserve v4uc-enable** command to enable QoS DSCP preservation for SR-TE doesn't help in IP DSCP preservation.

- This functionality isn't applicable for routers that have Cisco NC57 line cards installed and operating in native and compatibility modes. For these routers, DSCP preservation is by default, irrespective of the number of labels for both SR TI-LFA and SR-TE cases.
- The [hw-module profile sr-policy v6-null-label-autopush](#) and **hw-module profile mpls-ext-dscp-preserve v4uc-enable** commands are mutually exclusive.

Configuration for QoS IP DSCP Preservation for SR-TE

Configuration Example

To configure QoS IP DSCP preservation for SR-TE:

1. Match the DSCP value for a class map. (See [match dscp](#) for command details.)

```
Router(config)#class-map
Router(config)#class-map match-any exp_policy
Router(config-cmap)#match dscp cs2
Router(config-cmap)#end-class-map
Router(config)#
```

2. Set the MPLS EXP value for the same map. (See [set mpls experimental](#) for command details.)

```
Router(config)#policy-map exp_policy
Router(config-pmap)#class exp_policy
Router(config-pmap-c)#set mpls experimental imposition 7
Router(config-pmap-c)#set traffic-class 7
Router(config-pmap-c)#commit
```

3. Apply the policy on the ingress interface.

```
Router(config)#interface TenGigE0/0/0/0
Router(config-if)#service-policy input exp_policy
Router(config-if)#ipv4 address 192.0.2.1 255.255.255.0
Router(config-if)#commit
```

4. Nexthop reachability for ingress traffic is through SR-TE Tunnels. To configure SR-TE policies, see *Configure SR-TE Policies* in the *Segment Routing Configuration Guide for Cisco NCS 5500 Series Routers*.
5. Configure the DSCP preservation for SR-TE traffic and reload the chassis for the configuration to take effect.

```
Router#config
Router(config)#hw profile mpls-ext-dscp-preserve v4uc-enable
Router(config)#commit
Router(config)#exit
Router#reload location all
```

Running Configuration

```
class-map
class-map match-any exp_policy
match dscp cs2
end-class-map
!
```

```

policy-map
policy-map exp_policy
  class exp_policy
    set mpls experimental imposition 7
    set traffic-class 7
!

interface TenGigE0/0/0/0
service-policy input exp_policy
ipv4 address 192.0.2.1 255.255.255.0
!

hw profile mpls-ext-dscp-preserve v4uc-enable
!

```

Verification

Run the command **show controller fia diagshell 0 "config" location 0/0/cpu0 | i dscp** and in the output, verify that **custom_feature_dscp_preserve_with_extend_label** equals 1. The value is 1 only when the **hw profile mpls-ext-dscp-preserve v4uc-enable** command takes effect after router reload.

```

Router#show controller fia diagshell 0 "config" location 0/0/cpu0 | i dscp
      custom_feature_dscp_preserve_with_extend_label=1

```

QoS DSCP Preservation for mLDP

Table 15: Feature History Table

Feature Name	Release Information	Feature Description
QoS DSCP Preservation for mLDP	Release 7.7.1	<p>In terms of preserving IP DSCP markings, when you set the MPLS experimental bits (EXP) values (also called Traffic Class values), the IP DSCP markings are now preserved by default in the ingress policies when the MPLS labels are pushed into the packet.</p> <p>Traffic with IP packets with DSCP marking for priority, flows as intended and there's no drop in traffic because of incorrect or missing labels.</p> <p>In previous releases, irrespective of the MPLS label, when the EXP values were copied into the packet header during disposition, even the IP DSCP markings were modified. This modification resulted in traffic drops at the last-hop routers.</p>

A typical ingress scenario for mLDP traffic involves:

1. Specifying DSCP values as match criteria for a class map.
2. Identifying specific three-bit experimental bits (EXP) field values (also called Traffic Class or TC values, as per [RFC 5462](#)) in the MPLS label as match criteria for the same class map.

For MPLS networks, the default behaviour is uniform mode, where the MPLS labels are popped (disposition) at the end of the mLDP tunnel, the EXP values in the popped label is copied to the IP DSCP field, which causes the DSCP values that you set initially get modified. This modification causes traffic drops at the last-hop router.

It is this behavior that the QoS DSCP preservation for mLDP functionality mitigates.

From Cisco IOS XR Release 7.7.1 onwards, the QoS DSCP preservation for mLDP functionality is enabled by default when you bring up an edge router with mLDP profile configuration. Earlier, to help preserve IP DSCP values for mLDP traffic, you could disable the propagation of IP Time-To-Live (TTL) to and from the MPLS header by using the [mpls ip-ttl-propagate disable](#) command.

For more information on mLDP configuration, see the *Multicast MLDP Profile 14 support on an Edge Router* section in *Multicast Configuration Guide for Cisco NCS 5500 Series Routers*.

Configuration for QoS IP DSCP Preservation for mLDP

Configuration Example

To configure QoS IP DSCP preservation for mLDP:

1. Match the DSCP value for a class map. (See [match dscp](#) for command details.)

```
Router(config)#class-map
Router(config)#class-map match-any exp_policy
Router(config-cmap)#match dscp cs2
Router(config-cmap)#end-class-map
Router(config)#
```

2. Set the MPLS EXP value for the same map. (See [set mpls experimental](#) for command details.)

```
Router(config)#policy-map exp_policy
Router(config-pmap)#class exp_policy
Router(config-pmap-c)#set mpls experimental imposition 7
Router(config-pmap-c)#class class-default
Router(config-pmap-c)#commit
```

3. Apply the policy on the ingress interface.

```
Router(config)#interface TenGigE0/0/0/0
Router(config-if)#service-policy input exp_policy
Router(config-if)#ipv4 address 192.0.2.1 255.255.255.0
Router(config-if)#commit
```

4. Next-hop reachability for ingress traffic is through mLDP Tunnels. To configure mLDP, see the *Multicast MLDP Profile 14 support on an Edge Router* section in *Multicast Configuration Guide for Cisco NCS 5500 Series Routers*.

Running Configuration

```
class-map
class-map match-any exp_policy
match dscp cs2
end-class-map
```

```

!
policy-map
policy-map exp_policy
  class exp_policy
    set mpls experimental imposition 7
  class class-default
!

interface TenGigE0/0/0/0
service-policy input exp_policy
ipv4 address 192.0.2.1 255.255.255.0
!
!

```

Verification

Run the command **show policy-map interface** and in the output, verify that in the ingress interface the traffic arrive with EXP value as **exp 4** and DSCP value as **cs6**. At egress, the DSCP value remains same as **cs6**.

```

RP/0/RP0/CPU0:ios#sh policy-map interface te0/0/0/0 input
Thu Jan  6 23:42:36.039 UTC

```

```

TenGigE0/0/0/0 input: exp_policy

```

```

Class exp_policy
  Classification statistics      (packets/bytes)      (rate - kbps)
  Matched                      :      86601328/11084969984      821764
  Transmitted                  :      86601328/11084969984      821764
  Total Dropped                :                0/0                0
Class class-default
  Classification statistics      (packets/bytes)      (rate - kbps)
  Matched                      :      2347/300416                0
  Transmitted                  :      2347/300416                0
  Total Dropped                :                0/0                0
Policy Bag Stats time: 1641512538846 [Local Time: 01/06/22 23:42:18.846]
RP/0/RP0/CPU0:ios#

```

Verify the DSCP value of cs6 (represented as 6C) in the IPv6 traffic:

```

IPv6 traffic
68 2C 7B 84 58 4B 00 8A 96 0D 40 44 88 47 05 DC C9 3E 6C 00 00 00 00 46 3B 3F
00 77 00 00 00 00 00 00 00 00 00 00 00 00 09 FF 33 00 00 00 00 00 00 00 00 00
00 08 06 10 71 CF 38 A2 D3 70 49 78 69 61 60 00 00 10 11 12 13 E0 EC 18 77 00 2A 1A 1B
1C 1D 1E 1F 20 21 22 23 24 25 26 27 28 29 2A 2B 2C 2D 2E 2F 30 31 32 33 34 35 36 37 38 39
3A 3B 3C 3D 3E 3F 40 41 42 43 FB 2A

```

Policy-map for conditional marking of incoming IPv4 and IPv6 traffic

The incoming packets are classified based on the ingress policy-map and the following actions are done.

- Set qos-group
- Discard class or drop precedence is set implicitly or as a result of a policing action.
- Packets that violate the configured policer are dropped in the ingress processing itself.

Running Configuration:

```

policy-map ingress
class af11

```

```

    police rate percent 10 peak-rate percent 20
    !
    set qos-group 1
    !
    class af22
    police rate percent 30 peak-rate percent 50
    !
    set qos-group 2
    !
    class af32
    set qos-group 3
    police rate percent 30 peak-rate percent 60
    !
    !
    class class-default
    !
    end-policy-map
    !

```

Policy-map for conditional marking of outgoing MPLS traffic

The IPv4 or IPv6 ingress packet undergoes MPLS encapsulation during the egress processing in the PE router which performs the label imposition. The MPLS experimental bits are marked on the basis of egress policy-map which performs the following actions:

- Match on qos-group or discard class or both
- Set the MPLS experimental bits based on the match criteria

Running Configuration:

```

policy-map egress
  class qos1_disc0 # This class matches on qos-group 1 and discard-class 0
    set mpls experimental imposition 1
  !
  class qos1_disc1 # This class matches on qos-group 1 and discard-class 1
    set mpls experimental imposition 5
  !
  class qos2_disc0 # This class matches on qos-group 2 and discard-class 0
    set mpls experimental imposition 2
  !
  class qos2_disc1 # This class matches on qos-group 2 and discard-class 1
    set mpls experimental imposition 6
  !
  class qos3_disc0 # This class matches on qos-group 3 and discard-class 0
    set mpls experimental imposition 3
  !
  class qos3_disc1 # This class matches on qos-group 3 and discard-class 1
    set mpls experimental imposition 7
  !
  class class-default
  !
end-policy-map
!

```


Conditional Marking of MPLS Experimental bits for L2VPN Traffic



Note This feature is not available on NC55-24DD and NC55-18DD-SE line cards for Cisco IOS XR Release 7.0.2.

Conditional marking of MPLS EXP bits is supported on Virtual Private Wire Service (VPWS), Virtual Private LAN Service (VPLS) and Ethernet Virtual Private Network (EVPN) traffic in the L2VPN network.

The conditional marking of MPLS experimental bits is achieved for Layer 2 Virtual Private Network (L2VPN) traffic by applying a combination of ingress and egress policy-maps on the Provider Edge (PE) router. In the ingress policy-map, the qos-group or discard-class is set either based on the result of the policing action or implicitly. The egress policy-map matches on qos-group or on a combination of qos-group and discard-class and sets the mpls experiment bits to the corresponding value.

Conditional marking can be used to mark the MPLS experimental bits differently for in-contract and out-of-contract packets. In-contract packets are the confirmed packets with the color green and discard-class set to 0. Out-of-contract packets are the packets which have exceeded the limit and have the color yellow and discard-class set to 1.

Conditional marking of MPLS experimental bits for L2VPN traffic is supported on both physical and bundle main interfaces as well as sub-interfaces.

Restrictions for Conditional Marking of MPLS Experimental bits on L2VPN

1. Conditional marking of MPLS EXP bits is not supported for broadcast, unknown unicast and multicast (BUM) traffic.
2. In the case of two PE routers connected back-to-back and the only label that the traffic between the routers have is the BGP label, then the explicit null label should be configured.
3. A maximum of two policy-maps which perform conditional marking of MPLS experimental bits can be configured per Network Processor Unit (NPU) of the Cisco NCS 5500 Series Routers. However, the same policy can be applied on multiple interfaces on the same NPU.
4. In the ingress policy-map if qos-group is being set for the incoming traffic packets, then setting of dscp and mpls experimental bits will not work.
5. Both the ingress and egress policy-maps must be applied in order to attain the expected behaviour. If either one of them is not applied then it may lead to undefined behaviour.
6. If the egress policy-map does not match on qos-group or discard-class and set the mpls experiment bits to the required value, then the mpls experimental bits will be set to a value of zero, by default.

Policy-map for conditional marking of incoming traffic

The incoming packets on the Power Edge router are classified based on the ingress policy-map and these actions are taken.

- Set qos-group
- Discard class or drop precedence is set implicitly or as a result of a policing action.
- Set traffic class

- Packets that violate the configured policer are dropped in the ingress processing itself.

Running Configuration:

```

class-map af11
  match cos 1
!

policy-map ingress
  class af11
    police rate percent 10 peak-rate percent 20
    !
    set qos-group 1
    set Traffic-class 3
    !
  class class-default
    !
  end-policy-map
!
```

Policy-map for conditional marking of outgoing MPLS traffic

The ingress packet undergoes MPLS encapsulation during the egress processing in the PE router which performs the label imposition. The MPLS experimental bits are marked on the basis of egress policy-map which performs the following actions:

- Match on qos-group or discard class or both
- Set the MPLS experimental bits based on the match criteria

Running Configuration:

```

class-map match-all qos-group2_0
  match qos-group 2
  match discard-class 0

policy-map egress-marking
  class qos-group2_0 # This class matches on qos-group 2 and discard-class 0
    set mpls experimental imposition 1
    !
  class class-default
    !
  end-policy-map
!

policy-map Egress-Queuing
  class Traffic-class3
    shape average 500 mbps
    !
  class class-default
    !
  end-policy-map
!
```

Conditional Marking of MPLS Experimental Bits for EVPN-VPWS Single-Homing Services

Table 16: Feature History Table

Feature Name	Release Information	Feature Description
Conditional Marking of MPLS Experimental Bits for EVPN-VPWS Single-Homing Services	Release 7.3.1	<p>This feature enables you to differentiate traffic in the MPLS forwarding domain and manage traffic from ingress PE to egress PE based on the MPLS EXP bit of the MPLS header.</p> <p>This feature is supported only for EVPN-VPWS single-homing services, and not supported for EVPN-VPWS multi-homing services.</p>

The conditional marking of MPLS experimental bits is achieved for EVPN-VPWS single-homing services by applying a combination of ingress and egress policy-maps on the provider edge (PE) router. In the ingress policy-map, the qos-group or discard-class is set either based on the result of the policing action or implicitly. The egress policy-map matches on qos-group or on a combination of qos-group and discard-class and sets the MPLS experiment bits to the corresponding value.

Conditional marking can be used to mark the MPLS experimental bits differently for in-contract and out-of-contract packets. In-contract packets are the confirmed packets with the color green and discard-class set to 0. Out-of-contract packets are the packets that have exceeded the limit and have the color yellow and discard-class set to 1.

Conditional marking of MPLS experimental bits for EVPN-VPWS single-homing services are supported on both physical and bundle main interfaces as well as sub-interfaces.

MPLS EXP Marking for EVPN Multi-Homed Services

Table 17: Feature History Table

Feature Name	Release Information	Feature Description
MPLS EXP Marking for EVPN Multi-Homed Services	Release 7.6.1	You can now differentiate traffic in the MPLS forwarding domain and manage traffic from ingress PE to egress PE based on the MPLS EXP bit of the MPLS header for EVPN multi-home scenarios.



Note This feature is not supported on routers that have the Cisco NC57 line cards installed and operate in the native mode or compatibility mode.

Configuration

- The ingress policing is applied on the UNI interface. It is with set qos-group and set traffic class.

- The marking policy is applied at the core facing NNI interface.
- MPLS EXP imposition is marked while packets egress from NNI Interface.

Running Configuration

```

interface TenGigE0/0/0/2.203 l2transport => This is UNI
encapsulation dot1q 203
service-policy input pol50-100

interface TenGigE0/0/0/10 =====> This is the core NNI
description *** CORE IF ***
cdp
service-policy input in_mpls
service-policy output eg_mark
ipv4 address 192.18.44.18 255.255.255.0
ipv6 address 2005:18:44::18/48
lldp
    enable
!
monitor-session test ethernet direction tx-only port-level
!
load-interval 30
!
l2vpn
xconnect group 203
    p2p 203
        interface TenGigE0/0/0/2.203
            neighbor evpn evi 1 service 203

policy-map pol50-100
class class-default
    set traffic-class 2
    set qos-group 4
    police rate 50 mbps peak-rate 100 mbps
!
!
end-policy-map
!

policy-map eg_mark
class qg4dc0
    set mpls experimental imposition 2
!
class qg4dc1
    set mpls experimental imposition 3
!
class class-default
!
end-policy-map
!
class-map match-all qg4dc0
match qos-group 4
match discard-class 0
end-class-map
!

class-map match-all qg4dc1
match qos-group 4
match discard-class 1
end-class-map

```

Verification

Verify that you have configured conditional marking of MPLS experimental bits for EVPN-VPWS single-homing services successfully.

```
Router#show qos int tenGigE 0/0/0/2.101 input
NOTE:- Configured values are displayed within parentheses
Interface TenGigE0/0/0/2.101 ifh 0x41da -- input policy
NPU Id: 0
Total number of classes: 1
Interface Bandwidth: 10000000 kbps
Policy Name: pol50-100
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Levell Class = class-default
New traffic class = 2
New qos group = 4

Policer Bucket ID = 0x18
Policer Stats Handle = 0x0
Policer committed rate = 49219 kbps (50 mbits/sec)
Policer peak rate = 98438 kbps (100 mbits/sec)
Policer conform burst = 62336 bytes (default)
Policer exceed burst = 187008 bytes (default)
-----
```

```
Router#show qos int tenGigE 0/0/0/10 output
Tue Sep 1 04:18:27.508 UTC
NOTE:- Configured values are displayed within parentheses
Interface TenGigE0/0/0/10 ifh 0xe0 -- output policy
NPU Id: 0
Total number of classes: 3
Interface Bandwidth: 10000000 kbps
Policy Name: eg_mark
VOQ Base: 0
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Levell Class = qg4dc0
New imposition exp = 2
Queue Max. BW. = no max (default)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 0 / (BWR not configured)

Levell Class = qg4dc1
New imposition exp = 3
Queue Max. BW. = no max (default)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 0 / (BWR not configured)

Levell Class = class-default
Queue Max. BW. = no max (default)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 0 / (BWR not configured)
-----
```

QPPB

QoS Policy Propagation via BGP (QPPB) is a mechanism that allows propagation of quality of service (QoS) policy and classification by the sending party that is based on the following:

- Access lists
- Community lists
- Autonomous system paths in the Border Gateway Protocol (BGP)

Thus, helps in classification that is based on the destination address instead of the source address.

QoS policies that differentiate between different types of traffic are defined for a single enterprise network. For instance, one enterprise may want to treat important web traffic, not-important web traffic, and all other data traffic as three different classes. And thereafter, use the different classes for the voice and video traffic.

Hence, QPPB is introduced to overcome the following problems:

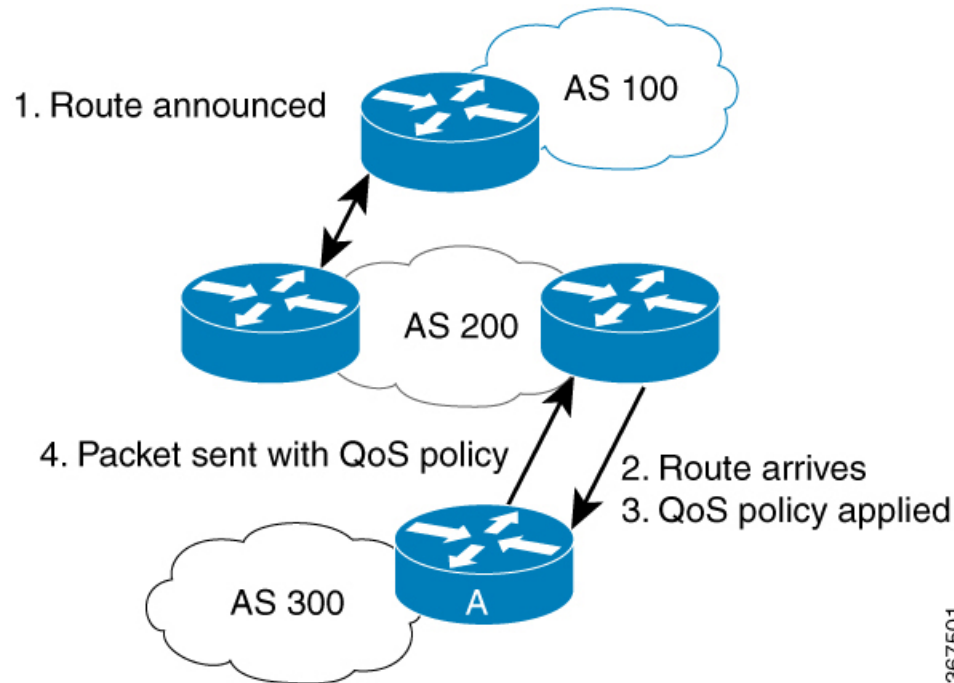
- The administrative challenges of classifying that is based on ACLs.
- The administrative problems of just listing the networks that need premium services.

QPPB allows marking of packets that are based on QoS group value associated with a Border Gateway Protocol (BGP) route.

Benefits of QPPB

- QPPB provides an IP prefix-based QoS capability.
- Traffic to IP addresses that have specific IP prefixes can be prioritized above other IP addresses.
- IP prefixes of interest are tagged through the control plane that uses common BGP route-map techniques, including the community attribute.
- Traffic to the tagged BGP prefixes is then classified and prioritized via the data forwarding plane by using the IOS-XR MQC (Modular QoS CLI) mechanisms, such as re-marking.
- QPPB provides the glue between the BGP control plane and the IP data forwarding plane in support of IP prefix-based QoS.
- BGP configuration within QPPB uses a table map to match specific prefixes learned through BGP neighbors, and then sets the router's local QoS Group variable maintained within the Forwarding Information Base (FIB) for those specific prefixes.
- The router supports a subset of full QPPB options - only IP destination prefix mode on input policy is supported.

Figure 3: Sample Scenario



Router A learns routes from AS 200 and AS 100. QoS policy is applied to any ingress interface of Router A to match the defined route maps with destination prefixes of incoming packets. Matching packets on Router A to AS 200 or AS 100 are sent with the appropriate QoS policy from Router A.

BGP maintains a scalable database of destination prefixes, QPPB, by using BGP table maps. BGP adds the ability to map a qos-group value to desired IP destinations. These qos-group values are used in QoS policies applied locally on ingress interfaces. Whenever a packet bound for such destinations is encountered, the qos-group value matching that destination route looks up with work inside the policy classmap, and marks that packet for any configured policy actions.

Configuration Workflow

Use the following configuration workflow for QPPB:

- Define route policy.
- Put Route policy at table-policy attach point under BGP.
- Define classmaps and ingress policy to use the qos-groups that are used in table-policy.
- Enable ipv4/ipv6 QPPB configuration under the desired interfaces.
- Configure the QPPB hardware profile, *hw-module profile qos ipv6 short*.
- If you use ipv6 QPPB, you must reload that linecard. If you use only ipv4 QPPB, linecard reload is not mandatory.

Define route policy

A routing policy instructs the router to inspect routes, filter them, and potentially modify their attributes as they are accepted from a peer, advertised to a peer, or redistributed from one routing protocol to another.

The routing policy language (RPL) provides a language to express routing policy. You must set up destination prefixes either to match inline values or one of a set of values in a prefix set.

Example:

```
prefix-set prefix-list-v4
  70.1.1.1,
  70.2.1.0/24,
  70.2.2.0/24 ge 28,
  70.2.3.0/24 le 28
end-set
prefix-set prefix-list-v6
  2001:300::2,
  2003:200::3
end-set

route-policy qppb1
  if destination in (60.60.0.2) then
    set qos-group 5
  elseif destination in prefix-list-v4 then
    set qos-group 4
  else
    set qos-group 1
  pass
endif
end-policy
```

For NC57 line cards, the set qos-group value under the route policy for QPPB support is 0-15. Even though the set qos-group value option displays supported values <0-31> as show below, the configuration fails for any value above 15 for NC57 line cards.

```
RP/0/RP0/CPU0:Fretta(config-rpl-if)#set qos-group ?
<0-31>      decimal number
```

Put Route policy at table-policy attach point under BGP

The table-policy attach point permits the route policy to perform actions on each route as they are installed into the RIB routing table. QPPB uses this attachment point to intercept all routes as they are received from peers. Ultimately the RIB will update the FIB in the hardware forwarding plane to store destination prefix routing entries, and in cases where table policy matches a destination prefix, the qos-group value is also stored with the destination prefix entry for use in the forwarding plane.

Example:

```
router bgp 900
  [vrf <name>]
  bgp router-id 22.22.22.22
  address-family ipv4 unicast
    table-policy qppb1
  address-family ipv6 unicast
    table-policy qppb2
  neighbor 30.2.2.1
    remote-as 500
    address-family ipv4 unicast
      route-policy pass in
      route-policy pass out
    address-family ipv6 unicast
```



```
route-policy pass in
route-policy pass out
```

Ingress interface QOS and ipv4/ipv6 bgp configuration

QPPB would be enabled per interface and individually for V4 and V6. An ingress policy would match on the qos groups marked by QPPB and take desired action.

If a packet is destined for a destination prefix on which BGP route policy has stored a qos-group, but it ingresses on an interface on which qppb is not enabled, it would not be remarked with qos-group.

Earlier, router supported matching on qos-group only in peering profile 'hw-module profile qos ingress-model peering location <>'. QPPB now permits classmaps to match qos-group in the default "non peering mode qos" as well. Also QPPB and hierarchical QOS policy profiles can work together if Hqos is used.

Example:

```
class-map match-any qos-group5
  match qos-group 5
end-class-map

class-map match-any qos-group4
  match qos-group 4
end-class-map

policy-map ingress-marker-pol
  class qos-group5
    set precedence 0
    set discard-class 0
    set traffic-class 1

  class qos-group4
    set precedence 1
    set discard-class 1
    set traffic-class 2
  class class-default

end-policy-map
```

Configuring QPPB on an Interface

1. RP/0/RP0/CPU0:router # configure

Enters interface configuration mode and associates one or more interfaces to the VRF.

2.

```
RP/0/RP0/CPU0:router(config)# interface
type interface-path-id
```

Enters interface configuration mode and associates one or more interfaces to the VRF.

3. ipv4 | ipv6 bgp policy propagation input qos-group destination

Example:

```
RP/0/RP0/CPU0:router(config-if)# ipv4 bgp policy propagation input qos-group destination
```

Enables QPPB on an interface

4. commit

Egress Interface Configuration

The traffic-class set on ingress has no existence outside the device. Also, traffic-class is not a part of any packet header but is associated internal context data on relevant packets. It can be used as a match criteria in an egress policy to set up various fields on the outgoing packet or shape flows.

Restrictions:


- No IP precedence marking.
- No policing on egress policy.

```
class-map match-any level1
  match traffic-class 1
end-class-map

class-map match-any level2
  match traffic-class 2
end-class-map

policy-map output-pol
  class level1
    bandwidth percent 50
  class level2
    bandwidth percent 20
    queue-limit 50 ms
end-policy-map

interface hun 0/5/0/0
  ipv4 address 30.1.1.1/24
  ipv6 address 2001:da8:b0a:12f0::1/64
  service-policy output output-pol
```



CHAPTER 4

Configuring Modular QoS Congestion Avoidance

This chapter covers the following topics:

- [Modular QoS Congestion Avoidance](#) , on page 81
- [Tail Drop and the FIFO Queue](#), on page 82
- [Random Early Detection and TCP](#), on page 84
- [Weighted Random Early Detection](#), on page 86
- [Explicit Congestion Notification \(ECN\)](#), on page 93

Modular QoS Congestion Avoidance

Congestion avoidance techniques monitor traffic flow to anticipate and avoid congestion at common network bottlenecks. Avoidance techniques are implemented before congestion occurs as compared with congestion management techniques that control congestion after it has occurred.



Note

From Cisco IOS XR Release 7.3.1 onwards, systems with Cisco NC57 line cards running in compatibility mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Starting with Cisco IOS XR Release 7.4.1 systems with Cisco NC57 line cards running in native mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS



Note For traffic requiring header decapsulation, the size of the header that is being removed is still included for the egress queuing actions. To offset this header size (required to achieve line rate for small frame sizes), configure an egress user policy with user overhead accounting on the egress interface. This policy can be a dummy policy configuration as well (allowing full traffic rate), if a policy isn't already in use or required on the egress interface.

You can enable user overhead accounting using the optional configuration of **accounting user-defined <overhead size in bytes>** while attaching the service policy on the egress interface.

Congestion avoidance is achieved through packet dropping. The router supports these QoS congestion avoidance techniques:

- [Tail Drop and the FIFO Queue, on page 82](#)
- [Random Early Detection and TCP, on page 84](#)
- [Weighted Random Early Detection, on page 86](#)

Tail Drop and the FIFO Queue

Tail drop is a congestion avoidance technique that drops packets when an output queue is full until congestion is eliminated. Tail drop treats all traffic flow equally and does not differentiate between classes of service. It manages the packets that are unclassified, placed into a first-in, first-out (FIFO) queue, and forwarded at a rate determined by the available underlying link bandwidth.

Configure Tail Drop

Packets satisfying the match criteria for a class accumulate in the queue reserved for the class until they are serviced. The **queue-limit** command is used to define the maximum threshold for a class. When the maximum threshold is reached, the enqueued packets to the class queue result in tail drop (packet drop).

Restrictions

- When configuring the **queue-limit** command, you must configure one of the following commands: **priority**, **shape average**, **bandwidth** or **bandwidth remaining**, except for the default class.

Configuration Example

You have to accomplish the following to complete the tail drop configuration:

1. Creating (or modifying) a policy map that can be attached to one or more interfaces to specify a service policy
2. Associating the traffic class with the traffic policy
3. Specifying the maximum limit the queue can hold for a class policy configured in a policy map.
4. Specifying priority to a class of traffic belonging to a policy map.

5. (Optional) Specifying the bandwidth allocated for a class belonging to a policy map or specifying how to allocate leftover bandwidth to various classes.
6. Attaching a policy map to an output interface to be used as the service policy for that interface.

```

Router# configure
Router(config)# class-map qos-1
Router(config-cmap)# match traffic-class 1
Router(config-cmap)# commit
Router(config-pmap)# exit

Router(config)# policy-map test-qlimit-1
Router(config-pmap)# class qos-1
Router(config-pmap-c)# queue-limit 100 us
Router(config-pmap-c)# priority level 7
Router(config-pmap-c)# exit
Router(config-pmap)# exit

Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy output test-qlimit-1
Router(config-if)# commit

```

Running Configuration

```

class-map qos-1
  match traffic-class 1
commit

policy-map test-qlimit-1
  class qos-1
    queue-limit 100 us
    priority level 7
  !
  class class-default
  !
end-policy-map
!

```

Verification

```
Router# show qos int hundredGigE 0/6/0/18 output
```

```

NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/6/0/18 ifh 0x3000220 -- output policy
NPU Id: 3
Total number of classes: 2
Interface Bandwidth: 100000000 kbps
VOQ Base: 11176
VOQ Stats Handle: 0x88550ea0
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Levell1 Class (HP7) = qos-1
Egressq Queue ID = 11177 (HP7 queue)
TailDrop Threshold = 1253376 bytes / 100 us (100 us)
WRED not configured for this class

Levell1 Class = class-default

```

```

Egressq Queue ID           = 11176 (Default LP queue)
Queue Max. BW.             = 101803495 kbps (default)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 (BWR not configured)
TailDrop Threshold         = 1253376 bytes / 10 ms (default)
WRED not configured for this class

```

Related Topics

- [Tail Drop and the FIFO Queue, on page 82](#)

Associated Commands

- [queue-limit](#)

Random Early Detection and TCP

The Random Early Detection (RED) congestion avoidance technique takes advantage of the congestion control mechanism of TCP. By randomly dropping packets prior to periods of high congestion, RED tells the packet source to decrease its transmission rate. Assuming the packet source is using TCP, it decreases its transmission rate until all packets reach their destination, indicating that the congestion is cleared. You can use RED as a way to cause TCP to slow transmission of packets. TCP not only pauses, but it also restarts quickly and adapts its transmission rate to the rate that the network can support.

RED distributes losses in time and maintains normally low queue depth while absorbing traffic bursts. When enabled on an interface, RED begins dropping packets when congestion occurs at a rate you select during configuration.

Configure Random Early Detection

The **random-detect** command with the **default** keyword must be used to enable random early detection (RED).

Guidelines

If you configure the **random-detect default** command on any class including class-default, you must configure one of the following commands: **shape average**, **bandwidth**, and **bandwidth remaining**.

Configuration Example

You have to accomplish the following to complete the random early detection configuration:

1. Creating (or modifying) a policy map that can be attached to one or more interfaces to specify a service policy
2. Associating the traffic class with the traffic policy
3. Enabling RED with default minimum and maximum thresholds.
4. (Optional) Specifying the bandwidth allocated for a class belonging to a policy map or specifying how to allocate leftover bandwidth to various classes.

5. (Optional) Shaping traffic to the specified bit rate or a percentage of the available bandwidth.
6. Attaching a policy map to an output interface to be used as the service policy for that interface.

```

Router# configure
Router(config)# class-map qos-1
Router(config-cmap)# match traffic-class 1
Router(config-cmap)# commit
Router(config-pmap)# exit

Router# configure
Router(config)# policy-map test-wred-2
Router(config-pmap)# class qos-1
Router(config-pmap-c)# random-detect default
Router(config-pmap-c)# shape average percent 10
Router(config-pmap-c)# end-policy-map
Router(config)# commit
Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy output test-wred-2
Router(config-if)# commit

```

Running Configuration

```

class-map qos-1
  match traffic-class 1
commit

policy-map test-wred-2
  class qos-1
    random-detect default
    shape average percent 10
  !
  class class-default
  !
  end-policy-map
!

interface HundredGigE 0/6/0/18
  service-policy output test-wred-2
!

```

Verification

```
Router# show qos int hundredGigE 0/6/0/18 output
```

```

NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/6/0/18 ifh 0x3000220 -- output policy
NPU Id:                                     3
Total number of classes:                   2
Interface Bandwidth:                       100000000 kbps
VOQ Base:                                  11176
VOQ Stats Handle:                          0x88550ea0
Accounting Type:                          Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class                               = qos-1
Egress Queue ID                            = 11177 (LP queue)
Queue Max. BW.                            = 10082461 kbps (10 %)

```

```

Queue Min. BW.                = 0 kbps (default)
Inverse Weight / Weight       = 1 (BWR not configured)
Guaranteed service rate       = 10000000 kbps
TailDrop Threshold            = 12517376 bytes / 10 ms (default)

Default RED profile
WRED Min. Threshold           = 12517376 bytes (10 ms)
WRED Max. Threshold           = 12517376 bytes (10 ms)

Level1 Class                   = class-default
Egressq Queue ID              = 11176 (Default LP queue)
Queue Max. BW.                = 101803495 kbps (default)
Queue Min. BW.                = 0 kbps (default)
Inverse Weight / Weight       = 1 (BWR not configured)
Guaranteed service rate       = 50000000 kbps
TailDrop Threshold            = 62652416 bytes / 10 ms (default)
WRED not configured for this class

```

Related Topics

- [Random Early Detection and TCP, on page 84](#)

Associated Commands

- [random-detect](#)

Weighted Random Early Detection

The Weighted Random Early Detection (WRED) drops packets selectively based on any specified criteria, like discard-class. WRED uses this matching criteria to determine how to treat different types of traffic.

You can configure WRED using the **random-detect** command and different discard-class values. The value can be range or a list of values that are valid for that field. You can also use minimum and maximum queue thresholds to determine the dropping point. Ensure that the WRED maximum threshold value is close to the queue limit. When the maximum threshold value is reached, packets start to get dropped.

When a packet arrives, the following actions occur:

- The average queue size is calculated.
- If the average queue size is less than the minimum queue threshold, the arriving packet is queued.
- If the average queue size is between the minimum queue threshold for that type of traffic and the maximum threshold for the interface, the packet is either dropped or queued, depending on the packet drop probability for that type of traffic.
- If the average queue size is greater than the maximum threshold, the packet is dropped.

Average Queue Size for WRED

The router automatically determines the parameters to use in the WRED calculations. The average queue size is based on the previous average and current size of the queue. The formula is:

$$\text{average} = (\text{old_average} * (1 - 2^{-x})) + (\text{current_queue_size} * 2^{-x})$$

where x is the exponential weight factor.

For high values of x , the previous average becomes more important. A large factor smooths out the peaks and lows in queue length. The average queue size is unlikely to change very quickly, avoiding a drastic change in size. The WRED process is slow to start dropping packets, but it may continue dropping packets for a time after the actual queue size has fallen below the minimum threshold. The slow-moving average accommodates temporary bursts in traffic.

**Note**

- The exponential weight factor, x , is fixed and is not user configurable.
- If the value of x gets too high, WRED does not react to congestion. Packets are sent or dropped as if WRED were not in effect.
- If the value of x gets too low, WRED overreacts to temporary traffic bursts and drops traffic unnecessarily.

For low values of x , the average queue size closely tracks the current queue size. The resulting average may fluctuate with changes in the traffic levels. In this case, the WRED process responds quickly to long queues. Once the queue falls below the minimum threshold, the process stops dropping packets.

Configure Weighted Random Early Detection

This configuration task is similar to that used for RED except that the **random-detect** command is not configured in RED.

Restrictions

- You cannot use the **random-detect** command in a class configured with the **priority** command, because WRED cannot be configured in a class that has been set for priority queueing (PQ).
- When configuring the **random-detect** command, you must configure one of the following commands: **shape average**, **bandwidth**, and **bandwidth remaining**.

Configuration Example

You have to accomplish the following to complete the random early detection configuration:

1. Creating (or modifying) a policy map that can be attached to one or more interfaces to specify a service policy
2. Associating the traffic class with the traffic policy
3. Enabling WRED by specifying the match criteria (discard-class).
4. (Optional) Specifying the bandwidth allocated for a class belonging to a policy map or specifying how to allocate leftover bandwidth to various classes.
5. (Optional) Shaping traffic to the specified bit rate or a percentage of the available bandwidth.
6. (Optional) Changing queue limit to fine-tune the amount of buffers available for each queue.
7. Attaching a policy map to an output interface to be used as the service policy for that interface.

Configure Weighted Random Early Detection

```

Router# configure
Router(config)# class-map qos-1
Router(config-cmap)# match traffic-class 1
Router(config-cmap)# commit
Router(config-pmap)# exit

Router# configure
Router(config)# policy-map test-wred-1
Router(config-pmap)# class qos-1
Router(config-pmap-c)# random-detect default
Router(config-pmap-c)# random-detect discard-class 0 10 ms 500 ms
Router(config-pmap-c)# shape average percent 10
Router(config-pmap-c)# commit

Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy output policy1
Router(config-if)# commit

```

Running Configuration

```

class-map qos-1
  match traffic-class 1
commit

policy-map test-wred-1
  class qos-1
    random-detect default
    random-detect discard-class 0 10 ms 500 ms
    shape average percent 10
  !
  class class-default
  !
end-policy-map
!

interface HundredGigE 0/6/0/18
  service-policy output test-wred-1
!

```

Verification

```
Router# show qos int hundredGigE 0/0/0/20 output
```

```

NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/0/0/20 ifh 0x38 -- output policy
NPU Id:                                0
Total number of classes:                2
Interface Bandwidth:                    100000000 kbps
Policy Name:                           test-wred-1
VOQ Base:                              1184
Accounting Type:                        Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class                           = qos-1
Egressq Queue ID                       = 1185 (LP queue)
Queue Max. BW.                         = 10000152 kbps (10 %)
Queue Min. BW.                         = 0 kbps (default)
Inverse Weight / Weight                 = 1 / (BWR not configured)
Guaranteed service rate                 = 10000000 kbps

```

```

Peak burst                               = 36864 bytes (default)
TailDrop Threshold                       = 1250000896 bytes / 1000 ms (default)

WRED profile for Discard_Class 0
WRED Min. Threshold                      = 12499968 bytes (10 ms)
WRED Max. Threshold                      = 624999936 bytes (500 ms)

Default RED profile
WRED Min. Threshold                      = 7499776 bytes (6 ms)
WRED Max. Threshold                      = 12499968 bytes (10 ms)

WRED ECN                                = Disabled

Levell1 Class                            = class-default
Egressq Queue ID                         = 1184 (Default LP queue)
Queue Max. BW.                           = no max (default)
Queue Min. BW.                           = 0 kbps (default)
Inverse Weight / Weight                  = 1 / (BWR not configured)
Guaranteed service rate                  = 50000000 kbps
Peak burst                               = 36864 bytes (default)
TailDrop Threshold                       = 62499840 bytes / 10 ms (default)
WRED not configured for this class

```

Related Topics

- [Weighted Random Early Detection, on page 86](#)
- [Configure Random Early Detection, on page 84](#)

Associated Commands

- [random-detect](#)

Configure WRED Counters by Class

Table 18: Feature History Table

Feature Name	Release Information	Feature Description
Configure WRED Counters by Class on Cisco NCS 5700 Series Base (Non-SE) Version Line Cards	Release 7.8.1	You can now view WRED statistics per class on Cisco NCS 5700 Series base (non-SE) version line cards operating in native mode. This functionality provides a more accurate and granular statistics profile for packet drops.

Feature Name	Release Information	Feature Description
Configure WRED Counters by Class	Release 7.4.1	<p>This feature enables the display of WRED statistics per class, thus providing a more accurate and granular statistics profile for packet drops. Such insight allows you to monitor, anticipate, and avoid congestion at common bottlenecks on your network.</p> <p>This functionality introduces the hw-module profile qos wred-stats-enable command and modifies the output of the show policy-map interface command.</p>

- **Statistics profile**—Up until now, running the **show policy-map interface** command displayed the drop statistics per class under tail drop. Which meant that no matter what drop mechanism you deployed ([Tail Drop and the FIFO Queue, on page 82](#), [Random Early Detection and TCP, on page 84](#), or [Weighted Random Early Detection, on page 86](#)), the statistics profile that displayed was a number aggregated under the tail drop entry. Given that WRED provides a far more granular control over packet drops than the other mechanisms, the generic statistics profile didn't provide you with the validation that you required for your WRED actions.
- **WRED counters by class**—This functionality introduces a new command, **hw-module profile qos wred-stats-enable**. Configuring this command enables the display of WRED and RED statistics per class, thus providing a more accurate and granular statistic for packet drops.
- **Why WRED counters by class**—Prior to this functionality, you did not have visibility into the WRED drop counters and had no way to validate WRED actions. With this improvement, the **show policy-map interface** command provides the much-needed visibility by clearly depicting the drop counters per class for WRED. This visibility gives you the confidence that WRED is indeed in action and is providing the requisite congestion avoidance.
- **Guidelines and limitations**
 - WRED max-drop threshold statistics are unavailable.
 - WRED statistics are collected separately for each drop precedence value.
- **WRED counters support cheat-sheet**—Here's a quick look at some key support areas and their details for WRED counters by class.

Support	Details
Line card	<p>Supported on Cisco NC57 line cards with external Ternary Content-Addressable Memory (TCAM) operating in native mode.</p> <p>Supported on Cisco NCS 5700 Series base (non-SE) version line cards operating in native mode.</p>

Support	Details
Statistics	WRED drop count is available for all the supported discard-values (from 0 through 2).

• **Configure WRED counters by class**—To configure WRED counters by class you must:

- Configure the **hw-module profile** for WRED statistics. Manually reload the chassis to activate the profile.
- Configure a class map.
- Create a service policy map.
- Attach the service policy to an interface.

```
/*Configure hw-module profile for wred stats*/
Router#conf
Router(config)#hw-module profile qos wred-stats-enable
Router(config)#commit
Router#reload

/*Configure a class map*/
Router(config)#class-map TC3
Router(config-cmap)#match traffic-class 3
Router(config-cmap)#commit
Router(config-cmap)#end-class-map

/*Create a service policy map*/
Router(config)#policy-map EGRESS_WRED
Router(config-pmap)#class TC3
Router(config-pmap-c)#shape average 1 gbps
Router(config-pmap-c)#random-detect discard-class 0 1 ms 2 ms
Router(config-pmap-c)#random-detect discard-class 1 3 ms 5 ms
Router(config-pmap-c)#random-detect discard-class 2 10 ms 20 ms
Router(config-pmap-c)#commit
Router(config-pmap)#end-policy-map

/*Attach the service policy to an interface*/
Router(config)#int hundredGigE 0/7/0/2
Router(config-if)#service-policy output EGRESS_WRED
Router(config-if)#commit
```

• **Running Configuration**

```
class-map match-any TC3
  match traffic-class 3
end-class-map
!
policy-map EGRESS_WRED
  class TC3
    shape average 1 gbps
    random-detect discard-class 0 1 ms 2 ms
    random-detect discard-class 1 3 ms 5 ms
    random-detect discard-class 2 10 ms 20 ms
  !
  class class-default
  !
```

```

        end-policy-map
    !
    interface HundredGigE0/7/0/2
        service-policy output EGRESS_WRED
        ipv4 address 192.168.0.1 255.255.255.0
    !

```

• Verification

```

Router#show qos int hundredGigE 0/7/0/2 output
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/7/0/2 ifh 0xe000088 -- output policy
NPU Id:                                0
Total number of classes:                2
Interface Bandwidth:                    100000000 kbps
Policy Name:                           EGRESS_WRED
SPI Id:                                0x0
VOQ Base:                               1160
PFC enabled:                            0
Accounting Type:                        Layer1 (Include Layer 1 encapsulation and above)
-----
Levell Class                           = TC3
Egressq Queue ID                       = 1163 (LP queue)
Queue Max. BW.                         = 1000000 kbps (1 gbits/sec)
Queue Min. BW.                         = 0 kbps (default)
Inverse Weight / Weight                 = 1 / (BWR not configured)
Guaranteed service rate                 = 1000000 kbps
Peak burst                             = 36864 bytes (default)
TailDrop Threshold                     = 4999168 bytes / 40 ms (default)

WRED profile for Discard_Class 0
WRED Min. Threshold                     = 124928 bytes (1 ms)
WRED Max. Threshold                     = 249856 bytes (2 ms)

WRED profile for Discard_Class 1
WRED Min. Threshold                     = 374784 bytes (3 ms)
WRED Max. Threshold                     = 624896 bytes (5 ms)

WRED profile for Discard_Class 2
WRED Min. Threshold                     = 1249792 bytes (10 ms)
WRED Max. Threshold                     = 2499840 bytes (20 ms)

Default RED profile
WRED Min. Threshold                     = 256 bytes (0 ms)
WRED Max. Threshold                     = 256 bytes (0 ms)

WRED ECN                               = Disabled

Levell Class                           = class-default
Egressq Queue ID                       = 1160 (Default LP queue)
Queue Max. BW.                         = no max (default)
Queue Min. BW.                         = 0 kbps (default)
Inverse Weight / Weight                 = 1 / (BWR not configured)
Guaranteed service rate                 = 50000000 kbps
Peak burst                             = 36864 bytes (default)
TailDrop Threshold                     = 62499840 bytes / 10 ms (default)
WRED not configured for this class

```

• show policy-map interface output for statistics verification

```

Router#show policy-map int hundredGigE 0/7/0/2 output

HundredGigE0/7/0/2 output: EGRESS_WRED

Class TC3

```

```

Classification statistics          (packets/bytes)      (rate - kbps)
Matched                          :          1151236/1151236000      0
Transmitted                      :          578807/578807000      0
Total Dropped                    :          572429/572429000      0
Queueing statistics
Queue ID                        : 1163
Taildropped(packets/bytes)      : 572429/572429000

WRED profile for Discard Class 0
RED Transmitted (packets/bytes)  : 0/0
RED random drops(packets/bytes) : 0/0
RED maxthreshold drops(packets/bytes) : N/A
RED ecn marked & transmitted(packets/bytes): N/A
WRED profile for Discard Class 1
RED Transmitted (packets/bytes)  : 0/0
RED random drops(packets/bytes) : 0/0
RED maxthreshold drops(packets/bytes) : N/A
RED ecn marked & transmitted(packets/bytes): N/A
WRED profile for Discard Class 2
RED Transmitted (packets/bytes)  : 578807/578807000
RED random drops(packets/bytes) : 572429/572429000
RED maxthreshold drops(packets/bytes) : N/A
RED ecn marked & transmitted(packets/bytes): N/A
RED profile (default)
RED Transmitted (packets/bytes)  : 0/0
RED random drops(packets/bytes) : 0/0
RED maxthreshold drops(packets/bytes) : N/A
RED ecn marked & transmitted(packets/bytes): N/A
Class class-default
Classification statistics          (packets/bytes)      (rate - kbps)
Matched                          :           0/0              0
Transmitted                      :           0/0              0
Total Dropped                    :           0/0              0
Queueing statistics
Queue ID                        : 1160
Taildropped(packets/bytes)      : 0/0

```

- **Clear WRED statistics on an interface**

Run the following command to clear WRED statistics on an interface.

```
Router#clear qos counters int hundredGigE 0/7/0/2 output
```

Associated Commands

[hw-module profile qos wred-stats-enable](#)

Explicit Congestion Notification (ECN)

Weighted Random Early Detection (WRED) is implemented at the core routers of a network. Edge routers assign IP precedences to packets, as the packets enter the network. With WRED, core routers then use these precedences to determine how to treat different types of traffic. WRED provides separate thresholds and weights for different IP precedences, enabling the network to provide different qualities of service, in regard to packet dropping, for different types of traffic. Standard traffic may be dropped more frequently than premium traffic during periods of congestion.

ECN is an extension to WRED. ECN marks packets instead of dropping them when the average queue length exceeds a specific threshold value. When configured, ECN helps routers and end hosts to understand that the network is congested and slow down sending packets. However If the number of packets in the queue is above

the maximum threshold, packets are dropped based on the drop probability. This is the identical treatment that a packet receives when WRED is enabled without ECN configured on the router.

RFC 3168, *The Addition of Explicit Congestion Notification (ECN) to IP*, states that with the addition of active queue management (for example, WRED) to the Internet infrastructure, routers are no longer limited to packet loss as an indication of congestion.



Note You cannot use this feature when you have set qos-group or mpls experimental along with a traffic class in the ingress policy.

Implementing ECN

Implementing ECN requires an ECN-specific field that has 2 bits—the ECN-capable Transport (ECT) bit and the CE (Congestion Experienced) bit—in the IP header. The ECT bit and the CE bit can be used to make four ECN field combinations of 00 to 11. The first number is the ECT bit and the second number is the CE bit.

Table 19: ECN Bit Setting

ECT Bit	CE Bit	Combination Indicates
0	0	Not-ECN-capable.
0	1	Endpoints of the transport protocol are ECN-capable.
1	0	Endpoints of the transport protocol are ECN-capable.
1	1	Congestion experienced.

The ECN field combination 00 indicates that a packet is not using ECN. The ECN field combinations 01 and 10—Called ECT(1) and ECT(0), respectively—are set by the data sender to indicate that the endpoints of the transport protocol are ECN-capable. Routers treat these two field combinations identically. Data senders can use either one or both of these two combinations. The ECN field combination 11 indicates congestion to the endpoints. Packets arriving a full queue of a router will be dropped.

Packet Handling When ECN Is Enabled

When the number of packets in the queue is below the minimum threshold, packets are transmitted. This happens whether ECN is enabled or not, and this treatment is identical to the treatment a packet receives when WRED only is being used on the network.

If the number of packets in the queue is above the maximum threshold:

- all incoming packets are dropped when you enable only WRED.
- all incoming packets are marked as CE when you configure ECN on the router.



Note When the number of packets reaches the queue limit, all packets are dropped. This is the identical treatment that a packet receives when you enable WRED without ECN configured on the router.

Three different scenarios arise if the number of packets in the queue is between the minimum threshold and the maximum threshold:

- If the ECN field on the packet indicates that the endpoints are ECN-capable (that is, the ECT bit is set to 1 and the CE bit is set to 0, or the ECT bit is set to 0 and the CE bit is set to 1)—and the WRED algorithm determines that the packet should have been dropped based on the drop probability—the ECT and CE bits for the packet are changed to 1, and the packet is transmitted. This happens because ECN is enabled and the packet gets marked instead of dropped.
- If the ECN field on the packet indicates that neither endpoint is ECN-capable (that is, the ECT bit is set to 0 and the CE bit is set to 0), packet is dropped once the queue limit is reached.
- If the ECN field on the packet indicates that the network is experiencing congestion (that is, both the ECT bit and the CE bit are set to 1), the packet is transmitted. No further marking is required.



Note When the incoming IP traffic with ECN bits set to 10 passes through the ingress qos-policy-map that has the class-map definition of `set DSCP/PREC <value>`, then the ECN bits in the IP header gets modified to 01. This is applicable to NC57 routers operating in the Native mode.

Limitations

- **The following limitation is not applicable to NC57 routers:**

SRv6 Disposition Node — For the incoming traffic that has the outer SRv6 IP header having the ECN bits as non-zero value and the inner IP header ECN bits as zero value, then the ECN bits of the inner IP header will be overwritten to 01.

- **The following limitation is specific to NC57 routers operating in Native mode.**



Note ECN remarking functionality is not supported with SRv6 tunnels. The inner header in the following scenario is IP and the outer header is SRv6.

SRv6 Imposition Node — The ECN bits in outer header traffic class is set to 0 when SRv6 encapsulation is set to traffic-class propagate along with Ingress MQC set to any DSCP/PREC value. This scenario does not impact the inner header ECN bits.

SRv6 Disposition Node — The ECN bits in the inner packet header is reset to 0 when the Ingress MQC is set to any DSCP/PREC value.

Configuration Example

```
Router# configure
Router(config)# policy-map policy1
Router(config-pmap)# class class1
Router(config-pmap-c)# bandwidth percent 50
Router(config-pmap-c)# random-detect 1000 packets 2000 packets
Router(config-pmap-c)# random-detect ecn
Router(config-pmap-c)# exit
Router(config-pmap)# exit
Router(config)# commit
```

Verification

Use the **show policy-map interface** to verify the configuration.

```
Router# show policy-map interface tenGigE 0/0/0/6 output
TenGigE0/0/0/6 output: pm-out-queue

Class cm-tc-1
  Classification statistics          (packets/bytes)      (rate - kbps)
    Matched                        :      85528554/87581239296      4830672
    Transmitted                    :      16240891/16630672384      966585
    Total Dropped                  :      69287663/70950566912      3864087
  Queueing statistics
    Queue ID                       :      1113
    Taildropped(packets/bytes)      :      69287663/70950566912

    WRED profile for
      RED Transmitted (packets/bytes)      : N/A
      RED random drops(packets/bytes)      : N/A
      RED maxthreshold drops(packets/bytes) : N/A
      RED ecn marked & transmitted(packets/bytes): N/A
Class class-default
  Classification statistics          (packets/bytes)      (rate - kbps)
    Matched                        :              0/0              0
    Transmitted                    :              0/0              0
    Total Dropped                  :              0/0              0
  Queueing statistics
    Queue ID                       :      1112
    Taildropped(packets/bytes)      :      0/0
```



Note No ECN-specific statistics are displayed in the show output for this command. ECN is enabled if all rows indicate **N/A**, as highlighted in the example.

Configure ECN Maximum Marking Probability

Table 20: Feature History Table

Feature Name	Release Information	Feature Description
Configure ECN Maximum Marking Probability	Release 7.3.3	<p>This feature allows you to configure percentages for ECN maximum marking probability after considering your network congestion and end-to-end application needs. With this flexibility, you can fine-tune the network's reactivity to congestion according to the tolerance levels of your end-to-end applications.</p> <p>By controlling the probability of packets being ECN marked at the ECN maximum threshold, you can optimize the congestion notification from the traffic destination to the traffic source. This optimization avoids aggressive throttling at the source (from too high ECN maximum marking probability) or traffic drops in transit due to significantly less throttling at the source (from too low ECN maximum marking probability).</p> <p>This functionality modifies the existing random-detect command.</p>

- **First, something about packet drop probability and Explicit Congestion Notification (ECN)**—The probability that a packet will be ECN-marked is based on the minimum threshold, maximum threshold, and maximum mark probability percent.

When the average queue size is above the minimum threshold, ECN starts marking packets to indicate congestion. When an ECN-marked packet reaches the end host, the host transmits confirmation of receipt of this packet to the transmitting host. The confirmation indicates to the transmitting host that congestion has begun and that the transmission or packet rate must be reduced.

Maximum mark probability is the percentage of packets marked when the average queue size is at the maximum threshold. The ECN mark probability increases linearly from 0 when queue length is less than minimum threshold to 10% when the queue length equals or exceeds the maximum threshold. For example, when the average queue size is halfway between minimum-threshold and maximum-threshold the mark probability is 5%, and half the packets transmitted will be ECN marked.

- **Configurable ECN Maximum Marking Probability**—So far, the ECN maximum marking probability was not user-configurable and was fixed at 10%. The arrangement of preset marking probabilities meant that the router couldn't adapt the congestion notification to network characteristics and application tolerance. With the flexibility to configure maximum ECN marking probability, you can configure the

maximum marking probability to a higher or lower percent, depending on the requirements of congestion handling in the network for a given traffic flow and application type. This also allows you to choose a consistent ECN maximum marking probability across all network elements instead of operating them at the factory default values. The following table maps the linear increase in mark probability percentage as the queue length conditions change from 0 to the tail drop threshold.

Queue length conditions	Mark probability increases linearly...
Queue length is less than the minimum threshold value	from 0
Average queue length is between minimum-threshold and maximum-threshold	corresponding to the average queue length: $(\text{maximum mark probability}) * (\text{average queue length} - \text{minimum threshold value}) / (\text{maximum threshold value} - \text{minimum threshold value})$
Queue length equals the maximum threshold	to the configured maximum probability percentage
maximum threshold value < queue length < tail drop threshold	to 100%

• Some FAQs

• What if I set the ECN marking probability percentage to 10%?

In such a case, there's no change in traffic behavior from what it was for previously ECN-marked packets. Tail drop comes into play when the queue length exceeds the tail drop threshold.

• How does traffic behave beyond my configured value of maximum ECN marking probability?

Let's say you've set the maximum ECN marking probability to 5%. Any more increase in the average queue length beyond maximum ECN threshold shifts the marking probability to 100%, and once the queue length exceeds the tail drop threshold, [Tail Drop and the FIFO Queue, on page 82](#) comes into effect.

• How do I enable this functionality?

Create a policy with **random-detect ecn** and **random-detect** *<min-threshold>* *<max-threshold>* **probability percent** *<value>* configured on an interface.

• How do I disable this functionality?

To disable this functionality, do not configure the **probability percent** value.

• Are the user-configured values programmed precisely in the hardware?

The user-configured values are rounded off to the nearest granularity during programming in the hardware. The round-off is minimal when the 'Pmax/(ECNmax-ECNmin)' ratio is a power of 2, where Pmax is the user-configured maximum marking probability in percent and ECNmin and ECNmax are the user-configured minimum and maximum ECN thresholds.

The **show qos interface** command output displays the actual value programmed in hardware and the user-configured value for reference.

• Important Guidelines

- **Line Cards that support this functionality:**
 - NC55-36X100G
 - NC55-18H18F
 - NC55-24X100G-SE
 - NC55-36X100G-S
 - NC55-24H12F-SE
 - NC55-36X100G-A-SE
 - NCS-55A1-36H-SE-S
 - NCS-55A1-36H-S
 - NCS-55A1-24H
 - NCS-55A1-48Q6H
- **The following interface types** support this functionality:
 - Physical interfaces
 - Bundle interfaces
 - Subinterfaces
 - Bundle subinterfaces
- This functionality is **supported for all interface speeds**.
- If you have a **policy map with one or multiple classes with maximum ECN marking probability enabled**, you can:
 - Apply the map to any of the supported interface types.
 - Remove the map from any of the supported interface types.
 - Modify the map while you're attaching it to multiple interfaces.
- The probability percentage option is supported only with **random-detect ecn** configured in the same class. Else, the policy is rejected when applied on an interface
- **Configuring ECN Maximum Marking Probability**—After you have configured [Explicit Congestion Notification \(ECN\)](#), on [page 93](#), the **random-detect** command displays an option to add the probability in percent.

Configuration Example

```
Router#configure
Router(config)#policy-map policy1
Router(config-pmap)#class class1
Router(config-pmap-c)#bandwidth percent 50
Router(config-pmap-c)#random-detect 1 mbytes 2 mbytes probability percent 5
Router(config-pmap-c)#random-detect ecn
Router(config-pmap-c)#commit
```

Verification

Run the **show qos interface** command to view the configured ECN Maximum Marking Probability value (displayed as **ECN Mark Probability**).

```
Router#sh qos interface FourHundredGigE 0/1/0/4 output
```

NOTE:- Configured values are displayed within parentheses

```
Interface FourHundredGigE0/1/0/4 ifh 0x800150 -- output policy
```

```
-----
.
.
Level2 Class                               =  LOSSLESSTCP_Egress
  Egressq Queue ID                         =  5475 (LP queue)
  Queue Max. BW.                           =  no max (default)
  Inverse Weight / Weight                  =  3 / (20)
  Guaranteed service rate                  =  78329670 kbps
  TailDrop Threshold                       =  390070272 bytes / 40 ms (default)

  Default RED profile
  WRED Min. Threshold                      =  9787392 bytes (1 ms)
  WRED Max. Threshold                      =  19580928 bytes (2 ms)
ECN Mark Probability                     =  4.99 (5)

  WRED ECN                                =  Enabled
```



CHAPTER 5

Priority Flow Control Overview

Table 21: Feature History Table

Feature Name	Release Information	Feature Description
Priority Flow Control	Release 7.3.2	Previously available in Release 6.6.3, this feature is a link-level flow control mechanism that enables you to selectively pause traffic based on its class of service (CoS).

Priority flow control (PFC; IEEE 802.1Qbb), which is also referred to as Class-based Flow Control (CBFC) or Per Priority Pause (PPP), is a mechanism that prevents frame loss that is due to transient congestion. PFC is similar to 802.3x Flow Control (pause frames) or link-level flow control (LLFC). However, instead of pausing all traffic on a link, PFC functions on a per class-of-service (CoS) basis.

During congestion, PFC sends a pause frame that indicates which CoS value must be paused. When the congestion is mitigated, the router stops sending the PFC frames to the upstream node.

For details about the Priority Flow Control feature, see the *Interface and Hardware Component Configuration Guide for Cisco NCS 5500 Series Routers*.

- [Restrictions, on page 101](#)
- [Configuring Priority Flow Control Thresholds, on page 103](#)
- [Priority Flow Control Watchdog Overview, on page 109](#)
- [Configure a Priority Flow Control Watchdog Interval, on page 109](#)
- [Monitoring and Logging Packet Drops on Lossless PFC-Enabled Queues, on page 112](#)

Restrictions

The following restrictions apply while configuring PFC thresholds, PFC watchdog and QoS policies in PFC profile.

- The PFC feature is only supported in the non-HQoS profile.
- The PFC feature is only supported on the following line card or fixed chassis PIDs of the NCS5500 Series:
 - NC55-36X100G

- NC55-18H18F
 - NC55-24X100G-SE
 - NC55-36X100G-S
 - NC55-24H12F-SE
 - NC55-36X100G-A-SE
 - NCS-55A1-36H-SE-S
 - NCS-55A1-36H-S
 - NCS-55A1-24H
 - NCS-55A1-48Q6H
- When PFC is enabled, only two parameter scheduling is supported. So, the egress actions could have either priority or weighted fair queue (WFQ) scheduling apart from a shaper action. For WFQ scheduling, either bandwidth remaining ratio (BRR) or bandwidth is supported and a mix of the two in the same policy-map is rejected. With PFC, BRR has a weight range of 1 to 256 as against 1 to 4096 for the non PFC, non-HQoS case.
 - The hw-module profile for enabling PFC per queue and defining PFC Tx thresholds are global configurations per line card location. So, only one set of pause, resume thresholds and headroom sizes per traffic-class can be configured for all ports on a given line card location. With the introduction of this new configuration model, the pause action under an egress policy class-map is deprecated.
 - While a line card reload is required for any addition or removal of traffic classes in the hw-module profile, you can update the pause, resume thresholds and buffer / headroom sizes on already configured traffic-classes without requiring a line card reload.
 - PFC watchdog (PFCWD) interval has a minimum granularity of 100ms. When you configure higher values, they are rounded up to the nearest multiple of 100ms.
 - There is no XML schema or yang model for the new hw-module profile. Only CLI configuration and show commands are available.
 - In the PFC profile, only monotonically increasing priority levels with traffic-class are supported. For example, TC7 should be p1, TC6 should be p2, and so on, until TC1 is p7 with class-default being lowest priority. Rate classes can be mixed anywhere.
 - When an egress queue is disabled as part of a watchdog queue shutdown action and once the pause frames stop, the queue drains the enqueued packets on the wire while it waits to be restored back.
 - If PFCWD is disabled and a traffic-class experiences sustained PFC storm beyond 7s, then the traffic on that queue continues to be dropped even when the PFC storm stops. Traffic needs to be stopped and restarted or all the incoming ports involved in sending traffic to this egress port needs to be flapped (shut and no shut) to manually recover the egress port.
 - The class-map actions of queue-limit and WRED thresholds without ECN enabled have no impact on PFC-enabled queues. Queue-limit is redundant because packet is buffered based on pause or x-off and headroom configured per source port.
 - Even when PFCWD shuts down an egress queue experiencing storm, if the ingress traffic was already causing a congestion due to the application of sub-rate shapers or BRR policy, then the ingress buffers build up and PFC Tx is still generated towards the sender. In other words, PFCWD only removes the

source of a storm. It does not suppress PFC Tx generated from sustained congestion due to user configuration.

- In the PFC profile, there is a 25% reduction in available buffer descriptors. Further, 20% of the remaining buffer space is reserved for headroom. This leads to only 80% of the 75% buffers being available in the shared pool for all ingress VoQs on that NP core.

Configuring Priority Flow Control Thresholds

You can set values for pause threshold (x-off), resume threshold (x-on) and headroom for a traffic class on all PFC enabled ports on a given line card location using the hw-module profile priority-flow-control command.

The existing queue-limit for that traffic class on an egress queuing policy on that line card will have no impact and the effective queue limit is pause threshold + headroom.



Note For optimal functionality in hardware, for a given traffic class, the resume threshold should not be more than 10% of the pause threshold and the headroom should be at least 100KB.

For details on how to configure PFC on an interface, see the *Interface and Hardware Component Configuration Guide for Cisco NCS 5500 Series Routers*.

Configuration Example

Apart from the line card level PFC parameters configured through the hw-module profile, PFC Receive functionality requires an egress policy to be applied and egress traffic to be mapped to unique queues per CoS value on all PFC enabled interfaces.

You can accomplish this with the following sequence of configuration.

1. Configure ingress classification policy to map incoming traffic to the right CoS / priority queue.
2. Configure egress queuing policy with the relevant shaping, priority and weighted fair queue scheduling actions
3. Attach the ingress policy on all interfaces where traffic can come in.
4. Attach the egress policy on all PFC enabled interfaces in the system.
5. Configure hw-module priority-flow-control profile and configure pause, resume and headroom for all PFC traffic-classes on all line card locations requiring PFC feature.

Reload the line card for all traffic-class addition and deletions. Parameters within an already configured traffic-class can be edited 'in place' without requiring a line card reload.

```
Hw-module configuration:
=====
RP/0/RP1/CPU0:NCS5504(config)#hw-module profile priority-flow-control location 0/0/CPU0
tRP/0/RP1/CPU0:NCS5504(config-pfc-loc)# traffic-class 3 pause-threshold 403200 bytes
resume-threshold 40320 bytes headroom 441600 bytes
RP/0/RP1/CPU0:NCS5504(config-pfc-loc)# traffic-class 4 pause-threshold 403200 bytes
resume-threshold 40320 bytes headroom 441600 bytes
RP/0/RP1/CPU0:NCS5504(config-pfc-loc)#
```

Class-map configuration:

```

=====
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_1
match prRP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 1
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_2
RP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 2
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_3
RP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 3
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_4
RP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 4
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_5
RP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 5
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_6
RP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 6
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any MAIN_IN_CMAP_7
RP/0/RP1/CPU0:NCS5504(config-cmap)# match precedence 7
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_1
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 1
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_2
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 2
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_3
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 3
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_4
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 4
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_5
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 5
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_6
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 6
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#class-map match-any PFC_OUT_CMAP_7
RP/0/RP1/CPU0:NCS5504(config-cmap)# match traffic-class 7
RP/0/RP1/CPU0:NCS5504(config-cmap)# end-class-map
RP/0/RP1/CPU0:NCS5504(config)#!
RP/0/RP1/CPU0:NCS5504(config)#

Ingress Policy-map configuration:
=====
RP/0/RP1/CPU0:NCS5504(config)#policy-map MAIN_OUT_TC_SUPPORTING_UUT
RP/0/RP1/CPU0:NCS5504(config-pmap)# class MAIN_IN_CMAP_1
RP/0/RP1/CPU0:NCS5504(config-pmap-c)# set traffic-class 1

```

```

RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class MAIN_IN_CMAP_2
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # set traffic-class 2
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class MAIN_IN_CMAP_3
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # set traffic-class 3
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class MAIN_IN_CMAP_4
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # set traffic-class 4
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class MAIN_IN_CMAP_5
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # set traffic-class 5
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class MAIN_IN_CMAP_6
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # set traffic-class 6
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class MAIN_IN_CMAP_7
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # set traffic-class 7
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) #
RP/0/RP1/CPU0:NCS5504(config-pmap-c) #

```

Egress policy-map configuration:

```

=====
RP/0/RP1/CPU0:NCS5504(config)#policy-map PFC_UUT_bwrr_microsoft
asRP/0/RP1/CPU0:NCS5504(config-pmap) # class PFC_OUT_CMAP_1
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 5
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # queue-limit 192 us
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class PFC_OUT_CMAP_2
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 5
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # queue-limit 192 us
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class PFC_OUT_CMAP_3
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 20
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # random-detect ecn
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # random-detect 224 kbytes 275 kbytes
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class PFC_OUT_CMAP_4
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 20
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # random-detect ecn
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # random-detect 224 kbytes 275 kbytes
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class PFC_OUT_CMAP_5
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 20
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # queue-limit 192 us
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class PFC_OUT_CMAP_6
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 1
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # queue-limit 192 us
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class PFC_OUT_CMAP_7
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # shape average percent 10
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # priority level 1
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # queue-limit 192 us
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # class class-default
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # bandwidth remaining ratio 20
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # queue-limit 192 us
RP/0/RP1/CPU0:NCS5504(config-pmap-c) # !

```

Applying policy to ingress and egress interface:

```
=====
```

```

RP/0/RP1/CPU0:NCS5504#show running-config interface hundredGigE 0/0/0/0
interface HundredGigE0/0/0/0
 service-policy input MAIN_OUT_TC_SUPPORTING_UUT
 ipv4 address 100.1.9.1 255.255.255.0
 ipv6 address 100:1:9::1/96
 priority-flow-control mode on

interface HundredGigE0/1/0/19
 service-policy output PFC_ECN_UUT_bwrr_microsoft
 ipv4 address 10.1.9.1 255.255.255.0
 ipv6 address 10:1:9::1/96
 priority-flow-control mode on
!

```

Running Configuration

```

RP/0/RP1/CPU0:NCS5504# show running-config hw-module profile priority-flow-control location
0/0/CPU0
Tue Oct 13 10:31:58.035 UTC
hw-module profile priority-flow-control location 0/0/CPU0
traffic-class 3 pause-threshold 403200 bytes resume-threshold 40320 bytes headroom 441600
bytes
traffic-class 4 pause-threshold 403200 bytes resume-threshold 40320 bytes headroom 441600
bytes
!

RP/0/RP1/CPU0:NCS5504#show running-config class-map
Tue Oct 13 10:32:02.400 UTC
class-map match-any MAIN_IN_CMAP_1
match precedence 1
end-class-map
!
class-map match-any MAIN_IN_CMAP_2
match precedence 2
end-class-map
!
class-map match-any MAIN_IN_CMAP_3
match precedence 3
end-class-map
!
class-map match-any MAIN_IN_CMAP_4
match precedence 4
end-class-map
!
class-map match-any MAIN_IN_CMAP_5
match precedence 5
end-class-map
!
class-map match-any MAIN_IN_CMAP_6
match precedence 6
end-class-map
!
class-map match-any MAIN_IN_CMAP_7
match precedence 7
end-class-map
!
class-map match-any PFC_OUT_CMAP_1
match traffic-class 1
end-class-map
!
class-map match-any PFC_OUT_CMAP_2
match traffic-class 2
end-class-map
!

```

```
class-map match-any PFC_OUT_CMAP_3
match traffic-class 3
end-class-map
!
class-map match-any PFC_OUT_CMAP_4
match traffic-class 4
end-class-map
!
class-map match-any PFC_OUT_CMAP_5
match traffic-class 5
end-class-map
!
class-map match-any PFC_OUT_CMAP_6
match traffic-class 6
end-class-map
!
class-map match-any PFC_OUT_CMAP_7
match traffic-class 7
end-class-map
!

RP/0/RP1/CPU0:NCS5504#show running-config policy-map PFC_UUT_bwrr_microsoft
Tue Oct 13 10:32:18.009 UTC
policy-map PFC_UUT_bwrr_microsoft
class PFC_OUT_CMAP_1
bandwidth remaining ratio 5
queue-limit 192 us
!
class PFC_OUT_CMAP_2
bandwidth remaining ratio 5
queue-limit 192 us
!
class PFC_OUT_CMAP_3
bandwidth remaining ratio 20
random-detect ecn
random-detect 224 kbytes 275 kbytes
!
class PFC_OUT_CMAP_4
bandwidth remaining ratio 20
random-detect ecn
random-detect 224 kbytes 275 kbytes
!
class PFC_OUT_CMAP_5
bandwidth remaining ratio 20
queue-limit 192 us
!
class PFC_OUT_CMAP_6
bandwidth remaining ratio 1
queue-limit 192 us
!
class PFC_OUT_CMAP_7
shape average percent 10
priority level 1
queue-limit 192 us
!
class class-default
bandwidth remaining ratio 20
queue-limit 192 us
!
end-policy-map
!

RP/0/RP1/CPU0:NCS5504#show running-config policy-map MAIN_OUT_TC_SUPPORTING_UUT
Tue Oct 13 10:32:31.430 UTC
```

```

policy-map MAIN_OUT_TC_SUPPORTING_UUT
class MAIN_IN_CMAP_1
set traffic-class 1
!
class MAIN_IN_CMAP_2
set traffic-class 2
!
class MAIN_IN_CMAP_3
set traffic-class 3
!
class MAIN_IN_CMAP_4
set traffic-class 4
!
class MAIN_IN_CMAP_5
set traffic-class 5
!
class MAIN_IN_CMAP_6
set traffic-class 6
!
class MAIN_IN_CMAP_7
set traffic-class 7
!
class class-default
!
end-policy-map
!

```

Verification

Incoming PFC frames are displayed on the PFC Rx statistics that are shown in the command below. If the interface is enabled for PFC and has an egress queuing policy applied to it, then the queue corresponding to the incoming CoS value on the PFC frames is paused

When PFC pause thresholds are crossed, PFC Tx frames are sent out of the PFC enabled interfaces sourcing such traffic on that network processing core. This can be verified against the PFC Tx statistics that are shown on this command. After the congestion condition ceases and the buffer resource usage goes below the PFC resume thresholds, a PFC resume frame is sent and the PFC Tx frames are stopped.

```
RP/0/RP1/CPU0:NCS5504#show controllers hundredGigE 0/0/0/0 priority-flow-control statistics
```

```
Mon Oct 12 12:22:39.362 UTC
```

```
Priority flow control information for interface HundredGigE0/0/0/0:
```

```

Priority Flow Control:
  Total Rx PFC Frames: 0
  Total Tx PFC Frames: 1764273
  Rx Data Frames Dropped: 0
  CoS   Status   Rx Frames   Tx Frames
  ---   -
  0    off           0           0
  1    off           0           0
  2    off           0           0
  3    on            0      882032
  4    on            0      882241
  5    off           0           0
  6    off           0           0
  7    off           0           0

```

The status of the traffic-class and the configured parameters under it on the hw-module priority-flow-control profile can be checked using the following command.

```
RP/0/RP1/CPU0:NCS5504#show controllers npu priority-flow-control loc 0/0/CPU0
Mon Oct 12 14:35:17.531 UTC
```

```
Location:      0/0/CPU0
PFC:           Enabled
TC      Pause-threshold  Resume-Threshold  Headroom
-----
3        403200 bytes      40320 bytes        441600 bytes
4        403200 bytes      40320 bytes        441600 bytes
```

Priority Flow Control Watchdog Overview

PFC Watchdog is a mechanism to identify any PFC storms (queue-stuck condition) in the network, and to prevent the PFC from propagating on the network and running in a loop. You can configure a PFC watchdog interval to detect whether packets in a no-drop queue are being drained within a specified time period. When the time period is exceeded, all outgoing packets are dropped on interfaces that match the PFC queue that is not being drained.

This requires monitoring PFC receiving on each port and detecting ports seeing an unusual amount of sustained pause frames. Once detected, the watchdog module can enforce several actions on such ports, which include generating a syslog message for network management systems, shutting down the queue, and auto-restoring the queue (after the PFC storm stops).

Related Topics

- [Priority Flow Control Overview, on page 101](#)

Configure a Priority Flow Control Watchdog Interval

You can configure PFC Watchdog parameters (Watchdog interval, shutdown multiplier, auto-restore multiplier) at the global or interface levels. Note that:

- When global Watchdog mode is disabled or off, Watchdog is disabled on all interfaces. This condition is regardless of the interface level Watchdog mode settings.
- When global Watchdog mode is enabled or on, the interface level Watchdog mode configuration settings override the global Watchdog mode values.
- When you configure interface level Watchdog attributes such as interval, shutdown multiplier, and auto-restore multiplier, they override the global Watchdog attributes.



Note Configuring the PFC mode and its policies is a prerequisite for PFC Watchdog.



Note PFC Watchdog also monitors, detects, and generates a syslog message every 5 minutes if global pause frames or link level flow control frames are received on a PFC-enabled port.

Such frames are ignored and discarded on the PFC-enabled port.

Configuration Example

You can configure the Watchdog at the global or at the interface level.



Note Watchdog is enabled by default, with system default values of:

Watchdog interval = 100 ms

Shutdown multiplier = 1

Auto-restart multiplier = 10

```
RP/0/RP1/CPU0:NCS5504# show controllers hundredGigE 0/0/0/0 priority-flow-control
watchdog-config
Mon Oct 12 14:32:47.056 UTC
```

Priority flow control information for interface HundredGigE0/0/0/0:

```
Priority flow control watchdog configuration:
(D) : Default value
U : Unconfigured
```

Configuration Item	Global	Interface	Effective
PFC watchdog state	: U	U	Enabled(D)
Poll interval	: U	U	100(D)
Shutdown multiplier	: U	U	1(D)
Auto-restore multiplier	: U	U	10(D)

```
RP/0/RP1/CPU0:NCS5504# show controllers hundredGigE 0/0/0/0 priority-flow-control
watchdog-stat
watchdog-state watchdog-stats
RP/0/RP1/CPU0:NCS5504# show controllers hundredGigE 0/0/0/0 priority-flow-control
watchdog-state
Mon Oct 12 14:32:56.760 UTC
```

Priority flow control information for interface HundredGigE0/0/0/0:

```
Priority flow control watchdog state machine state:
D - Disabled
M - Monitoring
S - Waiting For Shutdown
R - Waiting to Restore
```

```
-----
PFC Watchdog      : Enabled
Watchdog SM state : Traffic Class
                   7 6 5 4 3 2 1 0
                   - - - D D - - -
```

```
RP/0/RP1/CPU0:NCS5504# show controllers hundredGigE 0/0/0/0 priority-flow-control
watchdog-stats
Mon Oct 12 14:33:09.321 UTC
```


Priority flow control information for interface HundredGigE0/0/0/0:

Priority flow control watchdog statistics:

SAR: Auto restore and shutdown

Traffic Class	:	0	1	2	3	4	5	6
7								
Watchdog Events	:	0	0	0	0	0	0	0
0								
Shutdown Events	:	0	0	0	0	0	0	0
0								
Auto Restore Events	:	0	0	0	0	0	0	0
0								
SAR Events	:	0	0	0	0	0	0	0
0								
SAR Instantaneous Events	:	0	0	0	0	0	0	0
0								
Total Dropped Packets	:	0	0	0	0	0	0	0
0								
Dropped Packets	:	0	0	0	0	0	0	0
0								

RP/0/RP1/CPU0:NCS5504#

RP/0/RP1/CPU0:NCS5504#show controllers npu priority-flow-control loc 0/0/CPU0
Mon Oct 12 14:35:17.531 UTC

Location: 0/0/CPU0

PFC: Enabled

TC	Pause-threshold	Resume-Threshold	Headroom
3	403200 bytes	40320 bytes	441600 bytes
4	403200 bytes	40320 bytes	441600 bytes

RP/0/RP1/CPU0:NCS5504#

RP/0/RP1/CPU0:NCS5504#clear controller hundredGigE 0/0/0/0 priority-flow-control
watchdog-stats ?

traffic-class Traffic class to be cleared

<cr>

RP/0/RP1/CPU0:NCS5504#clear controller hundredGigE 0/0/0/0 priority-flow-control
watchdog-stats

Mon Oct 12 14:36:12.407 UTC

RP/0/RP1/CPU0:NCS5504#clear controller hundredGigE 0/0/0/0 priority-flow-control
watchdog-stats traffic-class ?

<0-7> Traffic class

RP/0/RP1/CPU0:NCS5504#clear controller hundredGigE 0/0/0/0 priority-flow-control
watchdog-stats traffic-class 3

NOTE: Use the clear commands to clear the statistics displayed by the show commands.

To restore a queue manually that has been shut down by the PFC watchdog action, run the command **set controller <> priority-flow-control recover traffic-class [0-7]**. Completion of this operation resets all internal watchdog state machines and the queue is back to a monitoring state.

RP/0/RP1/CPU0:NCS5504#show controllers hundredGigE 0/1/0/19 priority-flow-control
watchdog-state

Priority flow control information for interface HundredGigE0/1/0/19:

Priority flow control watchdog state machine state:

D - Disabled

M - Monitoring

```

S - Waiting For Shutdown
R - Waiting to Restore
-----
PFC Watchdog : Enabled
Watchdog SM state : Traffic Class
7 6 5 4 3 2 1 0
- - - M R - - -
RP/0/RP1/CPU0:NCS5504#

RP/0/RP1/CPU0:NCS5504#set controller hundredGigE 0/1/0/19 priority-flow-control recover
traffic-class 3

RP/0/RP1/CPU0:NCS5504#show controllers hundredGigE 0/1/0/19 priority-flow-control
watchdog-state

Priority flow control information for interface HundredGigE0/1/0/19:

Priority flow control watchdog state machine state:
D - Disabled
M - Monitoring
S - Waiting For Shutdown
R - Waiting to Restore
-----
PFC Watchdog : Enabled
Watchdog SM state : Traffic Class
7 6 5 4 3 2 1 0
- - - M M - - -

```

Related Topics

- [Priority Flow Control Overview, on page 101](#)

Monitoring and Logging Packet Drops on Lossless PFC-Enabled Queues

Table 22: Feature History Table

Feature Name	Release Information	Feature Description
Monitoring and Logging Packet Drops on Lossless PFC-Enabled Queues	Release 7.3.2	<p>In case of packet drops on lossless PFC-enabled queues that also have the PFC Watchdog feature enabled, this functionality generates syslogs in 10-minute intervals.</p> <p>Such timely alerts help you troubleshoot quickly, isolate issues, and reroute traffic if necessary, with minimal impact on end-user services.</p>

A Brief Background

Lossless PFC-enabled queues avoid queue tail drops during temporary congestion by sending PFC pause frames back to the previous network element to transmit lesser packets. Depending on the duration of the congestion, the previous network element may send PFC pause frames further upstream, hence triggering subsequent network elements to pause all the way up to the sender. Suppose that the congestion lasts longer or is persistent. In that case, the buffers on some of these network elements may still overflow, causing some packet drops eventually. PFC watchdog isolates such cases of persistent congestion and excessive PFC pause frames. PFC watchdog (enabled by default on all PFC enabled queues) monitors neighboring network elements that send excessive PFC pause frames and shuts down such queues until the PFC storm subsides. This action isolates the issue to just those queues, stops new PFC pause frame generation, and frees up the device buffers for other traffic flows, thus preventing the PFC storm from propagating throughout the network and affecting all other traffic flows.

From IOS XR Release 7.3.2, the PFC watchdog also monitors and periodically reports packet drops that occur due to excessive pausing, congestion, or PFC watchdog-triggered queue shutdown. While the **show controller** commands continue to provide an on-demand display of the total number of packets dropped for a given PFC enabled queue, the new logging mechanism provides timely alerts for proactive user intervention when drops occur on lossless PFC-enabled queues.

Highlights of Logging and Monitoring Packet Drops

From Release 7.3.2 onwards, packet drops on lossless PFC-enabled queues trigger syslog messages that record the drop statistics, alerting you about the event. These syslog messages enable monitoring of such queues and provide timely updates and alerts that help you take proactive action. You can troubleshoot quickly, isolate issues, and reroute traffic if necessary, with minimal impact on end-user services.

Here are some important highlights of the monitoring and logging of packet drops:

- *Only PFC-enabled queues* that also have the PFC Watchdog enabled are monitored.
- The syslog generation is event-based, where the event trigger is a packet drop on PFC Watchdog-monitored lossless queue between two consecutive measurement polls at a 10-minute interval. If the drop counters increment between such polls within a 10-minute interval, then this monitoring and logging functionality generates a syslog.
- Syslogs are generated for the following monitored parameters:
 - **Rx Data Frames Dropped** under **show controllers priority-flow-control statistics**. This counter accounts for virtual output queue (VOQ) tail drop statistics from all network processors across all ingress line cards in the system for a given egress port. VOQ tail drops occur when the PFC-enabled lossless queues build up beyond the user-configured headroom.
 - **Dropped Packets** under **show controllers priority-flow-control watchdog statistics**. This parameter accounts for all queue drops on the PFC-enabled lossless queues in the network processor on the egress line card.
- The **show controllers** command continues to display the total number of packets dropped.
- A typical syslog output looks like this:

```
%PLATFORM-PFC_WATCHDOG-5-PACKETDROP : PFC watchdog detected 150 packet drops on lossless
priority 3 of interface HundredGigE0_7_0_30. Total drops 5678 packets.
```

Restrictions for Monitoring and Logging Packet Drops

The following restrictions apply to the monitoring and logging of packet drops for lossless PFC-enabled queues:

- The 10-minute polling interval is fixed, and you can't modify it.
- The monitoring and logging functionality is enabled by default on all PFC Watchdog-enabled queues, and you can't disable it.



CHAPTER 6

Configuring Modular QoS Congestion Management

This chapter covers the following topics:

- [Congestion Management Overview, on page 115](#)
- [Ingress Traffic Management Model, on page 116](#)
- [Class-based Weighted Fair Queueing, on page 118](#)
- [Low-Latency Queueing with Strict Priority Queueing, on page 120](#)
- [Overhead Accounting, on page 122](#)
- [Traffic Shaping, on page 125](#)
- [4K Pseudowire on Bundle with QoS Enhancement , on page 127](#)
- [Traffic Policing, on page 129](#)
- [References for Modular QoS Congestion Management, on page 145](#)

Congestion Management Overview

Congestion management features allow you to control congestion by determining the order in which a traffic flow (or packets) is sent out an interface based on priorities assigned to packets. Congestion management entails the creation of queues, assignment of packets to those queues based on the classification of the packet, and scheduling of the packets in a queue for transmission.



Note

From Cisco IOS XR Release 7.3.1 onwards, systems with Cisco NC57 line cards running in compatibility mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Starting with Cisco IOS XR Release 7.4.1 systems with Cisco NC57 line cards running in native mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Table 23: Feature History Table

Feature Name	Release Information	Feature Description
Set Peak Burst Size for Egress Shaping	Release 7.4.1	<p>This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.</p> <p>With this feature you can configure the burst size along with the existing egress shaper bandwidth, using the "shape average shaper bandwidth burst size burst unit" command.</p> <p>This feature is used to control higher bursts of traffic being transmitted to the devices that have lower queue length configured to receive traffic. For more information, see Configure Traffic Shaping, on page 125.</p>
Absolute Rates for Bundle Interfaces for Traffic Shapers and Traffic Policers	Release 7.4.1	<p>This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes. For more information, see Configure Traffic Shaping, on page 125 and Configure Traffic Policing (Two-Rate Three-Color), on page 135.</p>
Ingress Policing on BVI (Bridge Group Virtual Interface) and Low-Latency Queueing (LLQ)	Release 7.4.1	<p>This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes. For more information, see "Restrictions for QoS on BVI" under Restrictions , on page 57 and Low-Latency Queueing with Strict Priority Queueing, on page 120</p>

The types of traffic regulation mechanisms supported are:

- [Class-based Weighted Fair Queueing, on page 118](#)
- [Low-Latency Queueing with Strict Priority Queueing, on page 120](#)
- [Traffic Shaping, on page 125](#)
- [Traffic Policing, on page 129](#)


Ingress Traffic Management Model

The ingress traffic management model relies on packet queueing on the egress interface using Virtual Output Queueing (VOQ) on the ingress. In this model, buffering takes place at ingress. Here's how the VOQ process works.

Your routers support up to eight output queues per main interface or physical port. For every egress output queue, the VOQ model earmarks buffer space on every ingress pipeline. This buffer space is in the form of

dedicated VOQs. These queues are called virtual because the queues physically exist on the ingress interface only when the line card actually has packets enqueued to it. To support the modular model of packet distribution, each network processing unit (NPU) core at the ingress needs connectors to every egress main interface and subinterface. The ingress traffic management model thus requires a mesh of connectors to connect the ingress NPU cores to the egress interfaces, as shown in **The Ingress Traffic Management Model**.

Figure 4: The Ingress Traffic Management Model



In the figure, every ingress interface (LC 0 through LC n) port has eight VOQs for the single egress line card LC 10.

Here's how packet transmission takes place:

1. When a packet arrives at ingress port (say on LC 0), the forwarding lookup on ingress line card points to the egress interface. Based on the egress interface (say it is on LC10), the packet is enqueued to the VOQ of LC 10. The egress interface is always mapped to a physical port.
2. Once egress bandwidth is available, the LC 10 ports ready to receive the packets (based on the packet marking and distribution model) send grants to the ingress ports via the connectors. (The figure shows a separate line for the grant for the sake of visual representation. In reality, the same connector is used for requests, grants, and transmission between an NPU core at the ingress and the egress port on LC 10.)
3. The ingress ports respond to this permission by transmitting the packets via FC to the LC 10 ports. (The time it takes for the ingress ports to request for egress port access, the egress port to grant access, and the packet to travel across FC is the round-trip time.)

The VOQ model thus operates on the principle of storing excess packets in buffers at ingress until bandwidth becomes available. Based on the congestion that builds up and the configured threshold values, packets begin to drop at the ingress itself, instead of having to travel all the way to the egress interface and then getting dropped.

Class-based Weighted Fair Queueing

Class-based Weighted Fair Queueing (CBWFQ) allows definition of traffic classes based on customer match criteria. With CBWFQ you can define traffic classes and assign guaranteed amount of minimum bandwidth to them. CBWFQ also allows for a strict priority queue for delay-sensitive traffic.

Bandwidth Remaining

The CBWFQ algorithm derives the weight for each class from the bandwidth remaining value allocated to the class. The **bandwidth remaining** option specifies a weight for the class to the CBWFQ. After the priority-queue is serviced, the leftover bandwidth is distributed as per bandwidth remaining ratio (BWRR) or percentage. If you do not configure this command for any class, the default value of the BWRR is considered as 1 (one). In the case of **bandwidth remaining percent**, the remaining bandwidth is equally distributed among other classes, to make it 100 percentage (100%).

Restrictions

- The **bandwidth remaining** command is supported only for egress policies.

Configure Minimum Bandwidth and Bandwidth Remaining

Guidelines

- The **priority** command cannot be configured along with **bandwidth**, **bandwidth remaining** commands, but can be configured with **shaping**, **queue-limit** and **random-detect** commands in the same class. Also see [Restrictions for Configuring H-QoS, on page 154](#).

Configuration Example

You have to accomplish the following to complete the minimum bandwidth and bandwidth remaining configuration:

1. Creating or modifying a policy-map that can be attached to one or more interfaces
2. Specifying the traffic class whose policy has to be created or changed
3. Allocating the minimum bandwidth and leftover bandwidth for the class
4. Attaching the policy-map to an output interface

```
Router# configure
Router(config)# policy-map test-bw-bw-rem
Router(config-pmap)# class qos-6
Router(config-pmap-c)# bandwidth percent 60
Router(config-pmap-c)# bandwidth remaining percent 60
```



```

Router(config-pmap-c# exit
Router(config-pmap)# exit
Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy output test-bw-bw-rem
Router(config-if)# commit

```

Running Configuration

```

policy-map test-bw-bw-rem
  class qos-6
    bandwidth percent 60
    bandwidth remaining percent 60
  !
  class qos-5
    bandwidth percent 20
    bandwidth remaining percent 40
  !
  class class-default
  !
end-policy-map
!

interface HundredGigE0/6/0/18
  service-policy input 100g-s1-1
  service-policy output test-bw-bw-rem
!

```

Verification

Router# **show qos interface HundredGigE 0/6/0/18 output**

```

NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/6/0/18 ifh 0x3000220 -- output policy
NPU Id:                                     3
Total number of classes:                   3
Interface Bandwidth:                       100000000 kbps
VOQ Base:                                  11176
VOQ Stats Handle:                          0x88550ea0
Accounting Type:                           Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class                               = qos-6
Egressq Queue ID                           = 11182 (LP queue)
Queue Max. BW.                             = 100824615 kbps (default)
Queue Min. BW.                             = 60494769 kbps (60 %)
Inverse Weight / Weight                    = 2 (60%)
Guaranteed service rate                    = 71881188 kbps
TailDrop Threshold                         = 90177536 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class                               = qos-5
Egressq Queue ID                           = 11181 (LP queue)
Queue Max. BW.                             = 100824615 kbps (default)
Queue Min. BW.                             = 20164923 kbps (20 %)
Inverse Weight / Weight                    = 3 (40%)
Guaranteed service rate                    = 27920792 kbps
TailDrop Threshold                         = 35127296 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class                               = class-default

```

```

Egressq Queue ID           = 11176 (Default LP queue)
Queue Max. BW.             = 101803495 kbps (default)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 120 (BWR not configured)
Guaranteed service rate    = 198019 kbps
TailDrop Threshold         = 247808 bytes / 10 ms (default)
WRED not configured for this class

```

Related Topics

- [Bandwidth Remaining, on page 118](#)

Associated Commands

- [bandwidth](#)
- [bandwidth remaining](#)

Low-Latency Queuing with Strict Priority Queuing

The Low-Latency Queuing (LLQ) feature brings strict priority queuing (PQ) to the CBWFQ scheduling mechanism. Priority queuing (PQ) in strict priority mode ensures that one type of traffic is sent, possibly at the expense of all others. For PQ, a low-priority queue can be detrimentally affected, and, in the worst case, never allowed to send its packets if a limited amount of bandwidth is available or the transmission rate of critical traffic is high.

Configuring Low Latency Queuing with Strict Priority queuing

Configuring low latency queuing (LLQ) with strict priority queuing (PQ) allows delay-sensitive data such as voice to be de-queued and sent before the packets in other queues are de-queued.

Guidelines

- Only priority level 1 to 7 is supported, with 1 being the highest priority and 7 being the lowest. However, the default CoSQ 0 has the lowest priority among all.
- Priority level 1 to 7 is supported for non-H-QoS profiles, with 1 being the highest priority and 7 being the lowest. For H-QoS profiles, priority level 1 to 4 is supported. For all profiles, however, the class default is CoSQ 0 and has the lowest priority among all.
- Egress policing is not supported. Hence, in the case of strict priority queuing, there are chances that the other queues do not get serviced.
- You can configure **shape average** and **queue-limit** commands along with **priority**.
- A Priority Queue (PQ) can oversubscribe bandwidth when other queues do not utilize the entire port bandwidth. However, oversubscription of traffic is supported only with a single priority level.

Configuration Example

You have to accomplish the following to complete the LLQ with strict priority queuing:

1. Creating or modifying a policy-map that can be attached to one or more interfaces
2. Specifying the traffic class whose policy has to be created or changed.
3. Specifying priority to the traffic class
4. (Optional) Shaping the traffic to a specific bit rate
5. Attaching the policy-map to an output interface

```
Router# configure
Router(config)#class-map qos-1
Router(config-cmap)#match traffic-class 1
Router(config-cmap)#commit

Router(config)#class-map qos-2
Router(config-cmap)#match traffic-class 2
Router(config-cmap)#commit

Router(config)# policy-map test-priority-1
Router(config-pmap)# class qos1
Router(config-pmap-c)# priority level 7
Router(config-pmap-c)# shape average percent 2
Router(config-pmap-c)# class qos-2
Router(config-pmap-c)# priority level 6
Router(config-pmap-c)# shape average percent 1
Router(config-pmap-c)# commit
Router(config-pmap-c)# exit
Router(config-pmap)# exit

Router(config)# interface HundredGigE 0/0/0/20
Router(config-if)# service-policy output test-priority-1
Router(config-if)# commit
```

Running Configuration

```
policy-map test-priority-1
  class qos-1
    priority level 7
    shape average percent 2
  !
  class qos-2
    priority level 6
    shape average percent 1
  !
  class class-default
  !
end-policy-map
!

interface HundredGigE0/0/0/20
  service-policy output test-priority-1
!
```

Verification

```
Router# show qos int hundredGigE 0/0/0/20 output
```

```

NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/0/0/20 ifh 0x38 -- output policy
NPU Id:                                0
Total number of classes:                3
Interface Bandwidth:                    100000000 kbps
Policy Name:                            test-priority-1
VOQ Base:                               1184
Accounting Type:                        Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class (HP7)                      = qos-1
Egressq Queue ID                        = 1185 (HP7 queue)
Queue Max. BW.                          = 2000000 kbps (2 %)
Guaranteed service rate                  = 2000000 kbps
Peak burst                              = 36864 bytes (default)
TailDrop Threshold                      = 2499840 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class (HP6)                      = qos-2
Egressq Queue ID                        = 1186 (HP6 queue)
Queue Max. BW.                          = 1000000 kbps (1 %)
Guaranteed service rate                  = 1000000 kbps
Peak burst                              = 36864 bytes (default)
TailDrop Threshold                      = 1249792 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class                            = class-default
Egressq Queue ID                        = 1184 (Default LP queue)
Queue Max. BW.                          = no max (default)
Queue Min. BW.                          = 0 kbps (default)
Inverse Weight / Weight                  = 1 / (BWR not configured)
Guaranteed service rate                  = 97000000 kbps
Peak burst                              = 36864 bytes (default)
TailDrop Threshold                      = 121249792 bytes / 10 ms (default)
WRED not configured for this class

```

Related Topics

- [Congestion Management Overview, on page 115](#)
- [Configure Traffic Shaping, on page 125](#)
- [Configure Minimum Bandwidth and Bandwidth Remaining, on page 118](#)

Associated Commands

- [priority](#)

Overhead Accounting

Traffic shapers and policers use packet traffic descriptors to ensure adherence to the service level agreement in QoS. However, when traffic flows from one hop to another in a network, headers added or removed at interim hops affect the packet bytes being accounted for by QoS at each hop. When your end-user network measures the packet bytes to ensure they receive the payload as agreed, these additional header bytes cause a discrepancy.

QoS overhead accounting provides the flexibility to operators to decide which header bytes can be excluded by the traffic shaper and policer and which can be included, depending on the end user's requirements and device capabilities, to meet the committed payload in units of bytes.

For example, if the QoS commitment includes the additional header bytes, the overhead accounting feature allows your router to account for this overhead and reduces the traffic policing and shaping rates accordingly. This is also called a **positive accounting overhead**.

If however, the committed rate doesn't include the additional bytes, overhead accounting allows your router to adjust the core stream traffic such that the traffic policing and shaping rates are increased. This is also called a **negative accounting overhead**.

To summarize, QoS overhead accounting enables the router to account for packet overhead when shaping and policing traffic to a specific rate. This accounting ensures that the router runs QoS features on the actual bandwidth that the subscriber traffic consumes.

Any interface that supports QoS policies supports overhead accounting.



Note You can enable user overhead accounting using the optional configuration of **accounting user-defined <overhead size in bytes>** while attaching the service policy on the egress interface.

Guidelines and Restrictions

- Overhead accounting for ingress shaping is not supported.

The following restrictions apply for routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.

- More than one compensation value can be programmed, provided you configure egress policy maps on different egress ports.
- You must configure a unique compensation value for a main interface and all sub-interfaces belonging to that main interface. You can't program different compensation values on different sub-interfaces sharing a common main interface.
- You can configure different compensation values on different sub-interfaces if they belong to other main interfaces.
- Compensation value programmed on egress queues (but not on VoQs) will remain active until the last egress policy map (with header compensation) is removed from main or sub-interfaces. This may impact traffic flow on main and sub-interfaces even though no compensation is set for them.

The following restrictions apply for routers that have line cards other than Cisco NC57 line cards.

- You can't program more than one compensation value per NPU or router, even if they're on different egress ports.
- You can configure the same egress compensation for different egress ports.
- NPUs can have different compensation values configured on different line cards in a modular system.

Configuring for Overhead Accounting

To configure overhead accounting, you must:

1. Create a policy map and configure QoS actions for that map.
2. Configure overhead accounting and attach the map to an interface.

```
/* create QoS policy */
Router#configure terminal
Router(config)#policy-map policer
Router(config-pmap)#class class-default
Router(config-pmap-c)#police rate percent 10
Router(config-pmap-c-police)#commit

/* configure account overhead value while attaching the QoS policy to interface */
Router(config)#int hundredGigE 0/0/0/2
Router(config-if)#service-policy input policer account user-defined 12
Router(config-if)#commit
Router(config-if)#root
Router(config)#end
```

Running Configuration

```
Router#sh run int hundredGigE 0/0/0/2
interface HundredGigE0/0/0/2
service-policy input policer account user-defined 12
!
```

The following example shows how to **configure a negative overhead accounting value**:

```
Router#conf
Router(config)#int hundredGigE 0/0/0/2
Router(config-if)#service-policy input policer account user-defined -12
Router(config-if)#commit
```

To **modify an overhead accounting value**, you must:

1. Remove the existing QoS policy and re-add it.
2. Configure the new overhead accounting value.

```
Router#conf
Router(config)#int hundredGigE 0/0/0/2
Router(config-if)#no service-policy input policer
Router(config-if)#service-policy input policer account user-defined -20
Router(config-if)#commit
Router#sh run int hundredGigE 0/0/0/2
interface HundredGigE0/0/0/2
service-policy input policer account user-defined -20
!
```

Positive Accounting Use Case

If QoS commitment includes Preamble, Frame Delimiter & Interframe Gap and has the following configuration:

```
service-policy input <foo> account user-defined +20
```

For QoS purposes, your router treats this packet as a packet of size = Actual Packet size + 20. Hence, the effective policing and shaping is *reduced* to match the downstream interface.

Negative Accounting Use Case

If QoS commitment to your router does not include VLAN header information, and has the following configuration:

```
service-policy input <foo> account user-defined -4
```

For QoS purposes, your router treats this packet as a packet of size = Actual Packet size – 4. Hence, the effective policing and shaping is *increased* to match the downstream interface.

Associated Commands

service-policy (overhead accounting)

Traffic Shaping

Traffic shaping allows you to control the traffic flow exiting an interface to match its transmission to the speed of the remote target interface and ensure that the traffic conforms to policies contracted for it. Traffic adhering to a particular profile can be shaped to meet downstream requirements, hence eliminating bottlenecks in topologies with data-rate mismatches.



Note If you apply a policy map that has configured traffic policing and traffic shaping on the basis of a percentage of bandwidth available on the interface and you change the speed of the interface, you must delete that policy map and reattach it to the interface. Else, QoS programming for the earlier speed remains in effect and does not change with change in the port speed.



Note Traffic shaping is supported only in egress direction.

Configure Traffic Shaping

The traffic shaping performed on outgoing interfaces is done at the Layer 1 level and includes the Layer 1 header in the rate calculation.

Guidelines

- Only egress traffic shaping is supported.
- It is mandatory to configure all the eight traffic-class classes (including class-default) for the egress policies.
- You can configure **shape average** command along with **priority** command.

Configuration Example

You have to accomplish the following to complete the traffic shaping configuration:

1. Creating or modifying a policy-map that can be attached to one or more interfaces
2. Specifying the traffic class whose policy has to be created or changed
3. Shaping the traffic to a specific bit rate and set peak burst size
4. Attaching the policy-map to an output interface

```

Router# configure
Router(config)#class-map c5
Router(config-cmap)#match traffic-class 5
Router(config-cmap)#commit

Router(config)# policy-map egress_policy1
Router(config-pmap)# class c5
Router(config-pmap-c)# shape average percent 50 1000
Router(config-pmap-c# exit
Router(config-pmap)# exit
Router(config)# interface HundredGigE 0/1/0/0
Router(config-if)# service-policy output egress_policy1
Router(config-if)# commit

```

Running Configuration

```

class-map c5
  match traffic-class 5
commit

policy-map egress_policy1
  class c5
    shape average percent 50 1000
  !
  class class-default
  !
end-policy-map
!

interface HundredGigE0/6/0/18
  service-policy output egress_policy1
!

```

Verification

```

Router# ios#show qos interface tenGigE 0/0/0/0 output

Wed Jul 10 14:18:37.783 UTC
NOTE:- Configured values are displayed within parentheses
Interface TenGigE0/0/0/0 ifh 0x120 -- output policy
NPU Id: 0
Total number of classes: 1
Interface Bandwidth: 10000000 kbps
Policy Name: test
VOQ Base: 1024
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = class-default
Egressq Queue ID = 1024 (Default LP queue)
Queue Max. BW. = 5031499 kbps (50 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 5000000 kbps
Peak burst = 2240 bytes (1000 bytes)
TailDrop Threshold = 6258688 bytes / 10 ms (default)

```


Important Notes

From Cisco IOS XR Release 6.6.25 onwards, a shaper on a bundle interface also allows absolute rates apart from the already supported units of percentage, per-thousand and per-million.

From Cisco IOS XR Release 7.4.1 onwards, this is also supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.

Related Topics

- [Congestion Management Overview, on page 115](#)

Associated Commands

- [shape average](#)

4K Pseudowire on Bundle with QoS Enhancement

Table 24: Feature History Table

Feature Name	Release Information	Feature Description
4K Pseudowire on Bundle with QoS Enhancement	Release 7.3.1	<p>With this feature you can configure a desired traffic policy, to which your network complies, by using the bandwidth management technique of two-level traffic shaping. You can also increase the Link Aggregation Group (LAG) sub-interface scale or pseudowires up to 4K.</p> <p>This increased scale value enables you to increase the number of devices connected to your router, resulting in benefits such as increased bandwidth and cost-effective operations.</p>

Your router supports two modes of egress traffic shaping:

- The **default non-Hierarchical QoS (H-QoS) mode**, where the egress traffic shaping is configured only on main interfaces and no hierarchical policies are supported. In other words, egress traffic shaping on subinterfaces isn't supported in this mode.
- The **H-QoS mode**, where egress traffic shaping is also supported on subinterfaces. This mode also supports hierarchical policies on the main and subinterfaces.

Prior to Cisco IOS XR Release 7.3.1, the H-QoS mode restricted the scale of subinterfaces you could configure. For example, the maximum Link Aggregation Group (LAG) subinterface scale or pseudowires is 1K in the H-QoS mode.

The enhancement described in this section, available from Cisco IOS XR Release 7.3.1 is applicable for the default non-H-QoS mode for egress QoS. It involves configuring two-level traffic shaper policy on the main interface, while enabling you to increase the LAG subinterface scale or pseudowires to as much as 4K. What this means is that you can also scale up the number of access devices you want to connect through LAG to your router, thus enabling:

- Increased reliability and availability
- Better use of physical resources
- Increased bandwidth
- Cost-effective operations

Restrictions and Guidelines

The following restrictions and guidelines apply while configuring two-level traffic shaper policy on the main interface in (default) non-H-QoS mode:

- The hierarchical egress policy support is only for main interfaces.
- Subinterface behavior remains the same in non-H-QoS mode. No egress QoS support is available for subinterfaces.
- This enhancement is applicable only for egress QoS, and there are no changes in ingress QoS behavior.
- There's no change to the current non-H-QoS flat policy behavior.
- The minimum shaper rate varies between different ASICs.

Configure Two-Level Shaper Policy on Main Interface

To configure two-level shaper policy on main interface, you must:

1. Enter global configuration mode and create a two-level policy map.
2. Attach this policy map to the main interface.

```
/* Enter the global configuration mode and create the two-level policy map */
Router#configure
Router(config)# policy-map two-level-pm
Router(config-pmap)#class class-default
Router(config-pmap-c)#shape average percent 20
Router(config-pmap-c)#service-policy child
Router(config-pmap-c)#root
Router(config)#policy-map child
Router(config-pmap)#class class-default
Router (config-pmap-c)#shape average percent 5
Router(config-pmap-c)#commit

/* Apply policy-map under interface */
Router(config)#interface hundredGigE 0/0/0/3
Router(config-intf)#service-policy output policy
Router(config-intf)#commit
```

Running Configuration

```
Router#show running-config policy-map two-level-pm
policy-map two-level-pm
  class class-default
    service-policy child
    shape average percent 20
  !
end-policy-map
!

Router#show running-config interface hundredGigE 0/0/0/3
interface HundredGigE0/0/0/3
service-policy output two-level-pm
!
```

Verification

Verify that the maximum bandwidth (shaping rate) for the parent policy (Level 1) is greater than the rate for the child policy at Level 2.

```
Router#show qos interface hundredGigE 0/0/0/3 output
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/0/0/3 ifh 0x68 -- output policy
NPU Id: 0
Total number of classes: 2
Interface Bandwidth: 100000000 kbps
Policy Name: two-level-pm
SPI Id: 0x0
VOQ Base: 1048
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = class-default
Queue Max. BW. = 20004211 kbps (20 %)
Queue Min. BW. = 20000305 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)

Level2 Class = class-default
Egressq Queue ID = 1048 (Default LP queue)
Queue Max. BW. = 1000000 kbps (5 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 1000000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 1249792 bytes / 10 ms (default)
WRED not configured for this class
```

Traffic Policing

Traffic policing allows you to control the maximum rate of traffic sent or received on an interface and to partition a network into multiple priority levels or class of service (CoS). Traffic policing manages the maximum rate of traffic through a token bucket algorithm. The token bucket algorithm uses user-configured values to determine the maximum rate of traffic allowed on an interface at a given moment in time. The token bucket algorithm is affected by all traffic entering or leaving the interface (depending on where the traffic policy with traffic policing is configured) and is useful in managing network bandwidth in cases where several large packets are sent in the same traffic stream. By default, the configured bandwidth value takes into account the Layer 2 encapsulation that is applied to traffic leaving the interface.

Traffic policing also provides a certain amount of bandwidth management by allowing you to set the burst size (Bc) for the committed information rate (CIR). See, [Committed Bursts and Excess Bursts, on page 130](#).

The router supports the following traffic policing mode(s):

- Single-Rate Two-Color (SR2C) in color-blind mode. See [Single-Rate Policer, on page 131](#).
- Single-Rate Three-Color (SR3C) in color-blind mode.
- Two-Rate Three-Color (2R3C) in color-blind mode. See [Two-Rate Policer, on page 135](#).



Note In NC57-18DD-SE, QoS enhanced stats is enabled by default.

Restrictions

- If you apply a policy map that has configured traffic policing and traffic shaping on the basis of a percentage of bandwidth available on the interface and you change the speed of the interface, you must delete that policy map and reattach it to the interface. Else, QoS programming for the earlier speed remains in effect and does not change with change in the port speed.
- Traffic policing is supported only in ingress direction, and only color-blind mode is supported.
- The policing rate accuracy may vary up to +/-2% from the configured policer value.
- Ensure that you don't configure a policer and match criteria for **discard-class** in the same class. Even though the configuration is allowed, the policer doesn't work and allows all traffic without dropping packets.
- Policer marking is not supported.
- Policers are configured in the interface at the core level and "show qos int <>" value is displayed at the NPU level.

For policers configured in a bundle interface where bundle members are from the same NPU but different cores (NPU cores), each member sends the traffic up to the core level policer configuration, but "show qos int <>" displays the NPU level policer output.

- Example:

For bundle interface with two 10GE members (same NPU, but one interface from core0, one interface from core1) 2R3C policer applied on bundle interface (1G confirm rate, 1G exceed rate – total 2G policer rate) will be shown on the "show qos int <>" output):

Interface in core0 – 500 Mbps confirm rate, 500 Mbps exceed rate

Interface in core1 – 500 Mbps confirm rate, 500 Mbps exceed rate

For traffic in one out of two interfaces, the policed rate will be 1Gbps. For traffic on two interfaces, policed rate will be 2Gbps.

Committed Bursts and Excess Bursts

Unlike a traffic shaper, a traffic policer does not buffer excess packets and transmit them later. Instead, the policer executes a "send or do not send" policy without buffering. Policing uses normal or committed burst

(bc) values and excess burst values (be) to ensure that the router reaches the configured committed information rate (CIR). Policing decides if a packet conforms or exceeds the CIR based on the burst values you configure. Burst parameters are based on a generic buffering rule for routers, which recommends that you configure buffering to be equal to the round-trip time bit-rate to accommodate the outstanding TCP windows of all connections in times of congestion. During periods of congestion, proper configuration of the excess burst parameter enables the policer to drop packets less aggressively.


For more details, see [Committed Bursts, on page 145](#) and [Excess Bursts, on page 145](#).

Single-Rate Policer

Single-Rate Two-Color Policer

A single-rate two-color (SR2C) policer provides one token bucket with two actions for each packet: a conform action and an exceed action.

Figure 5: Workflow of Single-Rate Two-Color Policer



Based on the committed information rate (CIR) value, the token bucket is updated at every refresh time interval. The Tc token bucket can contain up to the Bc value, which can be a certain number of bytes or a period of time. If a packet of size B is greater than the Tc token bucket, then the packet exceeds the CIR value and a default action is performed. If a packet of size B is less than the Tc token bucket, then the packet conforms and a different default action is performed.

Single-Rate Three-Color Policer

A single-rate three-color (SR3C) policer provides one token bucket with three actions for each packet: a conform action, an exceed action and a violate action. The packet is marked based on the CIR value and the two associated burst size - committed burst size (CBS) and excess burst size (EBS). If a packet does not exceed the CBS, it is marked as conformed packet. The packet is marked as exceeded if it exceeds CBS, but not the EBS. If it exceeds the EBS as well, it is marked as violate packet.

Configure Traffic Policing (Single-Rate Two-Color)

Traffic policing is often configured on interfaces at the edge of a network to limit the rate of traffic entering or leaving the network. The default conform action for single-rate two color policer is to transmit the packet and the default exceed action is to drop the packet. Users cannot modify these default actions.

Configuration Example

You have to accomplish the following to complete the traffic policing configuration:

1. Creating or modifying a policy-map that can be attached to one or more interfaces
2. Specifying the traffic class whose policy has to be created or changed
3. (Optional) Specifying the marking action
4. Specifying the policy rate for the traffic
5. Attaching the policy-map to an input interface

```
Router# configure
Router(config)# policy-map test-police-1
Router(config-pmap)# class ipv6-6
Router(config-pmap-c)# set dscp cs2 (optional)
Router(config-pmap-c)# set qos-group 7 (optional)
Router(config-pmap-c)# police rate percent 20 burst 10000 bytes
Router(config-pmap-c-police)# exit
Router(config-pmap-c)# exit
Router(config-pmap)# exit
Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy input test-police-1
Router(config-if)# commit
```

Running Configuration

```
class-map match-any ipv6-6
match precedence 3
end-class-map
!

policy-map test-police-1
class ipv6-6
set dscp cs2
set qos-group 7
police rate percent 20 burst 10000 bytes
!
class class-default
!
end-policy-map
!

interface HundredGigE0/6/0/18
service-policy input test-police-1
service-policy output test-priority-1
!
```

Verification

Router# **show qos interface hundredGigE 0/6/0/18 input**

```
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/6/0/18 ifh 0x3000220 -- input policy
NPU Id:                               3
Total number of classes:               2
Interface Bandwidth:                   100000000 kbps
Accounting Type:                       Layer1 (Include Layer 1 encapsulation and above)
-----
Levell1 Class                         =   ipv6-6
New dscp                              =   16
New qos group                         =   7

Policer Bucket ID                     =   0x102a0
Policer Stats Handle                  =   0x8a8090c0
Policer committed rate                 =   19980000 kbps (20 %)
Policer conform burst                  =   9856 bytes (10000 bytes)

Levell1 Class                         =   class-default

Default Policer Bucket ID              =   0x102a1
Default Policer Stats Handle            =   0x8a808e78
Policer not configured for this class
```

Related Topics

- [Traffic Policing, on page 129](#)

Associated Commands

- [police rate](#)

Configure Traffic Policing (Single-Rate Three-Color)

The default conform action and exceed actions for single-rate three-color policer are to transmit the packet and the default violate action is to drop the packet. User cannot modify these default actions.

Configuration Example

You have to accomplish the following to complete the traffic policing configuration:

1. Creating or modifying a policy-map that can be attached to one or more interfaces
2. Specifying the traffic class whose policy has to be created or changed
3. (Optional) Specifying the marking action
4. Configuring the policy rate for the traffic along with the peak-burst values
5. Attaching the policy-map to an input interface

```
Router# configure
Router(config)# policy-map test-police-1R3C
Router(config-pmap)# class ipv4-5
Router(config-pmap-c)# set qos-group 2 (optional)
```

```

Router(config-pmap-c)# police rate percent 20 burst 100000 bytes peak-burst 190000 bytes
Router(config-pmap-c-police)# exit
Router(config-pmap-c)# exit
Router(config-pmap)# exit
Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy input test-police-1R3C
Router(config-if)# commit

```

Running Configuration

```

class-map match-any ipv4-5
  match precedence 3
end-class-map
!

policy-map test-police-1R3C
  class ipv4-5
    set qos-group 7
    police rate percent 20 burst 100000 bytes peak-burst 190000 bytes
  !
  !
  class class-default
  !
end-policy-map
!

interface HundredGigE0/6/0/18
  service-policy input test-police-1R3C
  service-policy output test-priority-1
!

```

Verification

```
Router# show qos interface hundredGigE 0/6/0/18 input
```

```

NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/6/0/18 ifh 0x3000220 -- input policy
NPU Id:                                     3
Total number of classes:                   2
Interface Bandwidth:                       100000000 kbps
Accounting Type:                           Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class                               =   ipv4-5
New qos group                              =     2

Policer Bucket ID                          =   0x102a1
Policer Stats Handle                       =   0x8a8090c0
Policer committed rate                     =   19980000 kbps (20 %)
Policer conform burst                      =   99584 bytes (100000 bytes)
Policer exceed burst                       =   188672 bytes (190000 bytes)

Level1 Class                               =   class-default

Default Policer Bucket ID                  =   0x102a1
Default Policer Stats Handle                =   0x8a808e78
Policer not configured for this class

```


Related Topics

- [Traffic Policing, on page 129](#)

Associated Commands

- [police rate](#)


Two-Rate Policer

The two-rate policer manages the maximum rate of traffic by using two token buckets: the committed token bucket and the peak token bucket. The dual-token bucket algorithm uses user-configured values to determine the maximum rate of traffic allowed on a queue at a given moment. In this way, the two-rate policer can meter traffic at two independent rates: the committed information rate (CIR) and the peak information rate (PIR).

The dual-token bucket algorithm provides users with three actions for each packet—a conform action, an exceed action, and an optional violate action. Traffic entering a queue with the two-rate policer configured is placed into one of these categories. The actions are pre-determined for each category. The default conform and exceed actions are to transmit the packet, and the default violate action is to drop the packet.

This figure shows how the two-rate policer marks a packet and assigns a corresponding action to the packet.

Figure 6: Marking Packets and Assigning Actions—Two-Rate Policer



Also, see [Two-Rate Policer Details, on page 146](#).

The router supports Two-Rate Three-Color (2R3C) policer.

Configure Traffic Policing (Two-Rate Three-Color)

The default conform and exceed actions for two-rate three-color (2R3C) policer are to transmit the packet and the default violate action is to drop the packet. Users cannot modify these default actions.

Configuration Example

You have to accomplish the following to complete the two-rate three-color traffic policing configuration:

1. Creating or modifying a policy-map that can be attached to one or more interfaces
2. Specifying the traffic class whose policy has to be created or changed
3. Specifying the packet marking
4. Configuring two rate traffic policing
5. Attaching the policy-map to an input interface

```
Router# configure
Router(config)# policy-map policy1
Router(config-pmap)# class ipv4-7
Router(config-pmap-c)# set qos-group 4
Router(config-pmap-c)# police rate percent 20 burst 100000 bytes peak-rate percent 50
peak-burst 200000 bytes
Router(config-pmap-c-police)# exit
Router(config-pmap-c)# exit
Router(config-pmap)# exit
Router(config)# interface HundredGigE 0/6/0/18
Router(config-if)# service-policy input policy1
Router(config-if)# commit
```

Running Configuration

```
policy-map policy1
  class ipv4-7
    set qos-group 4
    police rate percent 20 burst 100000 bytes peak-rate percent 50 peak-burst 200000 bytes
  !
!

interface HundredGigE 0/6/0/18
  service-policy input policy1
!
```

Verification

```
Router# show policy-map interface HundredGigE 0/6/0/18
```

```
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/6/0/18 ifh 0x3000220 -- input policy
NPU Id: 3
Total number of classes: 8
Interface Bandwidth: 100000000 kbps
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = ipv4-4
- - -
- - -
Level1 Class = ipv4-7
New qos group = 4

Policer Bucket ID = 0x102a3
```

```

Policer Stats Handle           = 0x8a8089e8
Policer committed rate        = 19980000 kbps (20 %)
Policer peak rate              = 49860000 kbps (50 %)
Policer conform burst          = 99584 bytes (100000 bytes)
Policer exceed burst           = 199168 bytes (200000 bytes)

Level1 Class                    = class-default

Policer Bucket ID              = 0x102a7
Policer Stats Handle           = 0x8a7c8510
Policer committed rate        = 29880000 kbps (30 %)
Policer conform burst          = 4194304 bytes (default)

```

Important Notes

- From Cisco IOS XR Release 6.6.25 onwards:
 - a committed information rate of 0 kbps is supported, which is the only exception to the minimum rate of 22 kbps for both commit and excess rates.
 - a policer on a bundle interface also allows absolute rates apart from the already supported units of percent, per-thousand and per-million. From Cisco IOS XR Release 7.4.1 onwards, this is also supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.
- A policer is programmed per NPU core on a bundle interface. So, all members on a bundle interface from the same core share the policer.

Related Topics

- [Two-Rate Policer, on page 135](#)

Associated Commands

- [police rate](#)

Packets-Per-Second-Based Policer

Table 25: Feature History Table

Feature Name	Release Information	Feature Description
Packets-Per-Second-Based Policer	Release 7.4.1	<p>Prior to this functionality, when configuring policers, the only available option for policer rates was bit-rate measured in units of bits per second (bps). With this release, you can configure policer rates in units of packets per second (pps) as well. pps-based policer is critical in fending off malicious attacks—when attackers target your specific resources with a vast amount of traffic that contain higher number of packets, but move at a slower packet rate. Protection from such attacks is possible because pps-based policers ensure that regardless of the packet size and rate, the policer only accepts a fixed number of packets per second.</p> <p>This functionality modifies the police rate command.</p>

- **Policer rates so far**—You used the **police rate** command to configure policers based on two parameters:
 - bit-rates (default unit: bits per second or bps)
 - Burst size (default unit: bytes)
- **packets-per-second (pps)-based policer**—With this additional functionality, you can use the **police rate** command to configure policers in units of packets per second (pps). The pps configuration option is available as **police rate <pps-value> pps**. When you configure the pps option, ensure that you configure burst size in packets. (See **Restrictions and guidelines**.) Thus, the parameters for pps-based policer are:
 - packets per second (pps)
 - burst size (packets)
- **Why pps-based-policer**—Networks face newer types of attacks, and these days malicious operators don't necessarily employ aggressive tactics that involve overwhelming your bandwidth with large amount of traffic to cause distributed denial of service (DDoS). Now, some attackers go the 'softer' route, where they send smaller packet sizes at slower traffic rates. During such malicious network activity, a bandwidth-based policer can still aggregate up to many packets to be processed if the packet size is small. Attackers tend to use this behavior to bypass bandwidth-based policers to exploit vulnerabilities or try to hit performance limitations by increasing the packet rates.

Packets-per-second-based policers ensure that regardless of the packet size and traffic rate, the policer only accepts a fixed number of packets per second.

pps-based-policer support cheat-sheet—Here’s a quick look at some key support areas and their details for pps-based policer.

Support	Details
Classification and marking support	Same as that for bps-based-policer
Units	Equivalent kbps values display for QoS programming and statistics.
H-QoS	Support for parent and child policers
Bursts	Support for confirm burst (bc) and exceed burst (be) values in units of packets. The default value is in multiple of 128 bytes equivalent to 10 milliseconds.
Minimum pps value	For better granularity, recommended minimum value is 100 pps.

- **Restriction and guidelines**

- This functionality is **applicable only for ingress**.
- When using a pps-based policer, **ensure that you configure the burst-size value in number of packets as well**. This is because a policer burst rate determines whether a specific number of packets out of contract would be subject to the next action (that is, exceed or violate).
- **Within a QoS policy, configure the parent and child policies policers to either bps or pps**. Else, the configuration displays an error when you try attaching the policy to an interface.
- **For single-level policy maps:** under the same policy map, you can configure one class map with bps-based policer and the other class map with a pps-based policer.
- **For two-level hierarchical policy maps:**
 - The parent and child-level policy maps must use the same unit-based policer. That is, both must have either pps-based or bps-based policers.
 - If you configure the child-level policy map with pps-based policer, ensure that the parent policy-map class default has a pps-based policer.

- **Configure pps-based policer**—To configure pps-based policer, you must:

1. Configure a class map.
2. Create a service policy for the map and configure the pps values.
3. Attach the service policy to an interface.

```
/*Configure a class map*/
```

```
Router(config)#class-map prec1
Router(config-cmap)#match precedence 1
Router(config-cmap)# exit
Router(config)# commit
```

```

/*Create a service policy map*/

Router(config)# policy-map policy1
Router(config-pmap)# class prec1
Router(config-pmap-c)#police rate 1000 pps burst 300 packets
Router(config-pmap-c-police)#exit
Router(config-pmap-c)#exit
Router(config-pmap)#exit
Router(config)# commit

/*Attach the service policy to an interface*/
Router#int hundredGigE 0/7/0/2
Router(config-if)#service-policy input policy1
Router(config-if)#exit
Router(config)#commit

```

Running Configuration

```

class-map match-any prec1
match precedence 1
end-class-map
!
policy-map policy1
class prec1
    police rate 1000 pps burst 300 packets
!
class class-default
!
end-policy-map
!

```

Verification

```

Router#show qos int hundredGigE 0/7/0/2 input
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/7/0/2 ifh 0xe000088 -- input policy
NPU Id:                                0
Total number of classes:                2
Interface Bandwidth:                   100000000 kbps
Policy Name:                           policy1
SPI Id:                                0x0
Accounting Type:                        Layer2 (Include Layer 2 encapsulation and above)
-----
Level1 Class                            =    prec1

Policer Bucket ID                       =    0x9
Policer Stats Handle                    =    0x0
Policer committed rate                 =    998 kbps (1000 packets/sec)
Policer conform burst                  =    37632 bytes (300 packets)

Level1 Class                            =    class-default

Default Policer Bucket ID               =    0x8
Default Policer Stats Handle             =    0x0
Policer not configured for this class

```

Associated Commands

[police rate](#)

Per-thousand and Per-million Units

Shaper and policer rates can be configured in units of per-thousand and per-million on bundle interfaces. This provides the ability to provision shape and police rates down to 100 kbps on bundle or link aggregation (LAG) interfaces even with 100 GE bundle members.

For example, consider a 100GE interface and simple policy.

```
Interface HundredGig0/0/0/0
Service-policy output TEST
Policy-map TEST
Class C
  Shape average per-thousand 5
End-policy
```

Per thousand represents 0.1% of the link bandwidth and per million represents 0.0001% of the link bandwidth.

Which means that for a 100G link, 5 parts per thousand is 0.5% of the link bandwidth. Hence, the shape average per thousand of 5 in the above example enforces a shaper of 500 Mbps.

Shared Policer

Table 26: Feature History Table

Feature Name	Release Information	Feature Description
Shared Policer	Release 7.4.1	This feature is now supported on routers that have Cisco NC57 line cards installed and operate in native and compatibility modes.

The classification of the incoming packet occurs only once. However, based on the different classification criteria, the shared policer feature allows sharing of the policer bucket amongst two or more classes in a QoS policy map. That is, the same token bucket is used for a traffic flow matching against any of the classes sharing the policer.

For example, let us say a policer of 10 Mbps is shared among two classes C1 and C2. This feature ensures that both C1 and C2 get traffic flow assigned on First Come First Serve (FCFS) basis. Also that, if C2 does not have any traffic, C1 uses all of the 10 Mbps for transmission.

This feature includes two components:

- Policer Bucket Shared
- Policer Bucket Referred

Policer Bucket Shared

The policer bucket shared feature defines and shares a policer bucket instance among multiple classes.

Here is a sample configuration that defines and shares policer bucket instance sp1 :

```
policy-map parent
  class long-distance
    police bucket shared sp1 rate 1 mbps
```

In this configuration, a policy-map for class long-distance traffic type is created to police at 1Mbps rate and the policer bucket is shared.

Policer Bucket Referred

The policer bucket referred feature refers a defined policer bucket instance. Shared policer is not supported across policy levels. This means for example, that parent and child policy cannot share a common bucket.

Here is a sample configuration that refers shared policer bucket instance sp1 :

```
policy-map voip-child
  class long-distance-voip
    police bucket referred sp1
```

In this configuration, a policy-map for class long-distance-voip traffic type is created and the shared policer bucket sp1 is referred.

Shared Policer Statistics

Currently, individual class statistics are not available as a default option for shared policer. You can access statistics in the following modes.

Aggregate Mode

In this mode the policer bucket is shared among two or more classes. However, statistics are not available for every individual class. You can view the aggregate statistics that combine the numbers for all the classes sharing the policer bucket.

Per-Class Mode

In this mode the policer bucket is shared among two or more classes, and you can also view individual class statistics. However, when this mode is active, the Policy-Based Tunnel Selection (PBTS) mechanism is disabled. To enable the per-class mode, you must configure the **hw-module profile qos shared-policer-per-class-stats** command.

Restrictions and Guidelines

The following restrictions and guidelines apply while configuring the shared policer feature.

- When shared policer is enabled in per-class mode, Policy-Based Tunnel Selection (PBTS) mechanism is disabled. In other words, shared policer-per-class-mode and PBTS are mutually exclusive features.
- Shared policer is not supported across policy levels. This means, for example, that parent and child policies cannot share a common policer bucket.
- Shared policer is not supported in ingress peering mode.
- Shared policer is supported within classes of the same policy. However, cross-policy bucket sharing is not supported.
- There are no limitations on the number of classes that can share policer.
- There are no changes in policer scale numbers in the aggregate and per-class modes.
- All the existing policer types (1R2C, 1R3C and 2R3C) are supported.
- You must reload the affected line card to enable the per-class-stats mode.

Configuring Shared Policer

To configure shared policer, you must:

1. Create a class map to be used for matching packets to the specified class.
2. Create a policy map to be used for matching packets to the specified class.
3. Specify a class name.
4. Define and share a policer bucket.
5. Specify a class name.
6. Refer a shared policer bucket.

```
RP/0/RSP0/CPU0:ios(config)#class-map match-any c1
RP/0/RSP0/CPU0:ios(config-cmap)#match precedence 1
RP/0/RSP0/CPU0:ios(config-cmap)#end-class-map
RP/0/RSP0/CPU0:ios(config)#class-map match-any c2
RP/0/RSP0/CPU0:ios(config-cmap)#match precedence 2
RP/0/RSP0/CPU0:ios(config-cmap)#end-class-map
RP/0/RSP0/CPU0:ios(config)#policy-map s-pol
RP/0/RSP0/CPU0:ios(config-pmap)#class c1
RP/0/RSP0/CPU0:ios(config-pmap-c)#police bucket shared 1 rate 10 mbps
RP/0/RSP0/CPU0:ios(config-pmap-c-police)#exit
RP/0/RSP0/CPU0:ios(config-pmap-c)#exit
RP/0/RSP0/CPU0:ios(config-pmap)#class c2
RP/0/RSP0/CPU0:ios(config-pmap-c)#police bucket referred 1
RP/0/RSP0/CPU0:ios(config-pmap-c-police)#class class-default
RP/0/RSP0/CPU0:ios(config-pmap-c)#exit
RP/0/RSP0/CPU0:ios(config-pmap)#exit
RP/0/RSP0/CPU0:ios(config)#interface HundredGigE 0/6/0/18
RP/0/RSP0/CPU0:ios(config-if)#service-policy input s-pol
RP/0/RSP0/CPU0:ios(config-if)#commit
```

Running Configuration

```
class-map match-any c1
  match precedence 1
end-class-map

class-map match-any c2
  match precedence 2
end-class-map

policy-map s-pol
  class c1
    police bucket shared 1 rate 10 mbps
  !
  !
  class c2
    police bucket referred 1
  !
  !
  class class-default
  !
end-policy-map
!

interface HundredGigE0/6/0/18
  service-policy input s-pol
!
```

Verification

In Aggregate Mode

```
RP/0/RP0/CPU0:ios#sh policy-map interface tenGigE 0/0/0/0 input
Fri Nov 15 12:55:56.817 UTC

TenGigE0/0/0/0 input: s-pol

Class c1
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                    : 1784530245/228419871360      8640780
    Transmitted                : 2067504/264640512          10011
    Total Dropped              : 1782462741/228155230848      8630769
  Policing statistics          (packets/bytes)      (rate - kbps)
    Policed(conform)           : 2067504/264640512          10011
    Policed(exceed)            : 1782462741/228155230848      8630769
    Policed(violate)           : 0/0                      0
    Policed and dropped        : 1782462741/228155230848

Class c2
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                    : 0/0                      0
    Transmitted                : 0/0                  0
    Total Dropped              : 0/0                      0
  Policing statistics          (packets/bytes)      (rate - kbps)
    Policed(conform)           : 0/0                      0
    Policed(exceed)            : 0/0                      0
    Policed(violate)           : 0/0                      0
    Policed and dropped        : 0/0

Class class-default
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                    : 0/0                      0
    Transmitted                : 0/0                      0
    Total Dropped              : 0/0                      0
Policy Bag Stats time: 1573822531986 [Local Time: 11/15/19 12:55:31.986]
```

In Per-Class Mode

```
RP/0/RP0/CPU0:ios#sh policy-map interface tenGigE 0/0/0/0 input
Fri Nov 15 15:18:18.319 UTC

TenGigE0/0/0/0 input: s-pol

Class c1
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                    : 1005369276/128687267328      4320337
    Transmitted                : 1163300/148902400          5013
    Total Dropped              : 1004205976/128538364928      4315324
  Policing statistics          (packets/bytes)      (rate - kbps)
    Policed(conform)           : 1163300/148902400          5013
    Policed(exceed)            : 1004205976/128538364928      4315324
    Policed(violate)           : 0/0                      0
    Policed and dropped        : 1004205976/128538364928

Class c2
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                    : 1005341342/128683691776      4320335
    Transmitted                : 1166269/149282432          4997
    Total Dropped              : 1004175073/128534409344      4315338
  Policing statistics          (packets/bytes)      (rate - kbps)
    Policed(conform)           : 1166269/149282432          4997
    Policed(exceed)            : 1004175073/128534409344      4315338
    Policed(violate)           : 0/0                      0
    Policed and dropped        : 1004175073/128534409344
```

```

Class class-default
  Classification statistics      (packets/bytes)      (rate - kbps)
    Matched                    :          49159/6292352          0
    Transmitted                 :          49159/6292352          0
    Total Dropped               :              0/0              0
Policy Bag Stats time: 1573831087338 [Local Time: 11/15/19 15:18:07.338]

```

Related Commands hw-module profile qos shared-policer-per-class-stats

References for Modular QoS Congestion Management

Committed Bursts

The committed burst (bc) parameter of the police command implements the first, conforming (green) token bucket that the router uses to meter traffic. The bc parameter sets the size of this token bucket. Initially, the token bucket is full and the token count is equal to the committed burst size (CBS). Thereafter, the meter updates the token counts the number of times per second indicated by the committed information rate (CIR).

The following describes how the meter uses the conforming token bucket to send packets:

- If sufficient tokens are in the conforming token bucket when a packet arrives, the meter marks the packet green and decrements the conforming token count by the number of bytes of the packet.
- If there are insufficient tokens available in the conforming token bucket, the meter allows the traffic flow to borrow the tokens needed to send the packet. The meter checks the exceeding token bucket for the number of bytes of the packet. If the exceeding token bucket has a sufficient number of tokens available, the meter marks the packet.

Green and decrements the conforming token count down to the minimum value of 0.

Yellow, borrows the remaining tokens needed from the exceeding token bucket, and decrements the exceeding token count by the number of tokens borrowed down to the minimum value of 0.

- If an insufficient number of tokens is available, the meter marks the packet red and does not decrement either of the conforming or exceeding token counts.



Note When the meter marks a packet with a specific color, there must be a sufficient number of tokens of that color to accommodate the entire packet. Therefore, the volume of green packets is never smaller than the committed information rate (CIR) and committed burst size (CBS). Tokens of a given color are always used on packets of that color.

Excess Bursts

The excess burst (be) parameter of the police command implements the second, exceeding (yellow) token bucket that the router uses to meter traffic. The exceeding token bucket is initially full and the token count is equal to the excess burst size (EBS). Thereafter, the meter updates the token counts the number of times per second indicated by the committed information rate (CIR).

The following describes how the meter uses the exceeding token bucket to send packets:

- When the first token bucket (the conforming bucket) meets the committed burst size (CBS), the meter allows the traffic flow to borrow the tokens needed from the exceeding token bucket. The meter marks the packet yellow and then decrements the exceeding token bucket by the number of bytes of the packet.
- If the exceeding token bucket does not have the required tokens to borrow, the meter marks the packet red and does not decrement the conforming or the exceeding token bucket. Instead, the meter performs the exceed-action configured in the police command (for example, the policer drops the packets).

Two-Rate Policer Details

The committed token bucket can hold bytes up to the size of the committed burst (bc) before overflowing. This token bucket holds the tokens that determine whether a packet conforms to or exceeds the CIR as the following describes:

- A traffic stream is conforming when the average number of bytes over time does not cause the committed token bucket to overflow. When this occurs, the token bucket algorithm marks the traffic stream green.
- A traffic stream is exceeding when it causes the committed token bucket to overflow into the peak token bucket. When this occurs, the token bucket algorithm marks the traffic stream yellow. The peak token bucket is filled as long as the traffic exceeds the police rate.

The peak token bucket can hold bytes up to the size of the peak burst (be) before overflowing. This token bucket holds the tokens that determine whether a packet violates the PIR. A traffic stream is violating when it causes the peak token bucket to overflow. When this occurs, the token bucket algorithm marks the traffic stream red.

For example, if a data stream with a rate of 250 kbps arrives at the two-rate policer, and the CIR is 100 kbps and the PIR is 200 kbps, the policer marks the packet in the following way:

- 100 kbps conforms to the rate
- 100 kbps exceeds the rate
- 50 kbps violates the rate

The router updates the tokens for both the committed and peak token buckets in the following way:

- The router updates the committed token bucket at the CIR value each time a packet arrives at the interface. The committed token bucket can contain up to the committed burst (bc) value.
- The router updates the peak token bucket at the PIR value each time a packet arrives at the interface. The peak token bucket can contain up to the peak burst (be) value.
- When an arriving packet conforms to the CIR, the router takes the conform action on the packet and decrements both the committed and peak token buckets by the number of bytes of the packet.
- When an arriving packet exceeds the CIR, the router takes the exceed action on the packet, decrements the committed token bucket by the number of bytes of the packet, and decrements the peak token bucket by the number of overflow bytes of the packet.
- When an arriving packet exceeds the PIR, the router takes the violate action on the packet, but does not decrement the peak token bucket.

See [Two-Rate Policer](#), on page 135.



CHAPTER 7

Configuring Modular QoS on Link Bundles

This chapter covers the following topics:

- [QoS on Link Bundles, on page 147](#)

QoS on Link Bundles

A bundle is a group of one or more ports that are aggregated together and treated as a single link. The router supports Ethernet interfaces and VLAN interfaces (bundle sub-interfaces) bundles. All QoS features currently supported on physical interfaces, are also supported on all link bundle interfaces. Applying QoS on bundle members is not supported.



Note From Cisco IOS XR Release 7.3.1 onwards, systems with Cisco NC57 line cards running in compatibility mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Starting with Cisco IOS XR Release 7.4.1 systems with Cisco NC57 line cards running in native mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Restrictions for Link Bundles

- Only Ethernet link bundling is supported.
- A bundle interface can only contain physical interface.
- All links within a single bundle must be configured either to run 802.3ad (LACP) or Etherchannel (non-LACP). Mixed links within a single bundle are not supported.
- MAC accounting is not supported on Ethernet link bundles.
- Maximum number of links supported in each link bundle is 64.

- The maximum number of link bundles supported is 128.

Load Balancing

Load balancing function is a forwarding mechanism to distribute traffic over multiple links based on Layer 3 routing information in the router. Per-destination load balancing is only supported on the router, where the router is allowed to distribute packets over one of the links in the bundle. When the per-destination load balancing is enabled, all packets for a certain source-destination pair go through the same link, though there are multiple links available. In other words, per-destination load balancing can ensure that packets for a certain source-destination pair could arrive in order.

Layer 3 Load Balancing on Link Bundles

Layer 3 load balancing for link bundles is done on Ethernet Flow Points (EFPs) and is based on the IPv4 source and destination addresses in the packet. When Layer 3 service-specific load balancing is configured, all egress bundles are load balanced based on the IPv4 source and destination addresses. When packets do not have IPv4 addresses, default load-balancing (based on the MAC SA/DA fields in the packet header) is used.

Configure QoS on Link Bundles

QoS is configured on link bundles in the same way that it is configured on individual interfaces.

Guidelines

- When a QoS policy is applied on a bundle in the egress direction, it's also applied at each member interface.
- When a QoS policy is applied on a bundle (ingress direction), it's replicated at each NPU core.
- If a QoS policy is not applied to a bundle interface, both the ingress and egress traffic use the default queue of the per link member port.
- The shape rate that is specified in the bundle policy-map is not an aggregate for all bundle members. The shape rate applied to the bundle depends on the load balancing of the links. For example, if a policy map with a shape rate of 10 Mbps is applied to a bundle with two member links, and if the traffic is always load-balanced to the same member link, then an overall rate of 10 Mbps applies to the bundle. However, if the traffic is load-balanced evenly between the two links, the overall shape rate for the bundle becomes 20 Mbps.
- If a member is deleted from a bundle, the total bundle statistics change because the statistics that belongs to the detached link is lost.
- The QoS policy that is applied on bundle is inherited to all its member links and the reference bandwidth that is used to calculate shaper/bandwidth is applied as per the physical member interface bandwidth, and not the bundle as a whole.

Configuration Example

You have to accomplish the following to complete the QoS configuration on link bundles:



Note The policy works only if it is applied on the ingress direction. The egress is supported on COS, DEI and MPLS exp marking. So the below policy may not work when it is applied on egress.

1. Creating a class-map
2. Creating a policy-map and specifying the respective class-map
3. Specifying the action type for the traffic

Refer [Attach a Traffic Policy to an Interface, on page 18](#) for details on step 1, 2 and 3.

4. Creating a link bundle
5. Applying traffic policy to the link bundle

```
/* Configure Ether-Bundle and apply traffic policy */
Router(config)# interface Bundle-Ether 12000
Router(config-if)# mtu 9100
Router(config-if)# service-policy input ingress
Router(config-if)# service-policy output egress
Router(config-if)# ipv4 address 100.12.0.0 255.255.255.254
Router(config-if)# bundle maximum-active links 64
Router(config-if)# commit
```

Running Configuration

This example shows how a traffic policy is applied on an Ethernet link bundle. The policy is applied to all interfaces that are members of the Ethernet link bundle.

```
/* Policy-map */

policy-map ingress
  class inet4-classifier-af1
    set qos-group 1
  !
  class inet4-classifier-af2
    set qos-group 2
  !
  class inet4-classifier-af3
    set qos-group 3
  !
  class inet4-classifier-af4
    set qos-group 4
  !
  class inet4-classifier-bel
    set qos-group 5
  !
  class inet4-classifier-ncl
    set qos-group 6
  !
  class class-default
  !
end-policy-map
!

/* Ether Bundle */
```

```

interface Bundle-Ether12000
mtu 9100
service-policy input ingress
service-policy output egress
ipv4 address 100.12.0.0 255.255.255.254
load-interval 30
flow ipv4 monitor FMM-V4 sampler SM ingress
flow ipv6 monitor FMM-V6 sampler SM ingress
flow mpls monitor FMM-MPLS sampler SM ingress
ipv4 access-group IPV4ACL_101 ingress
ipv6 access-group IPV6ACL_101 ingress
!

```

Verification

- Verify that the bundle status is UP.

```

router# show bundle bundle-ether 1200
Wed Dec 16 19:55:49.974 PST

```

```

Bundle-Ether12000
Status: Up
Local links <active/standby/configured>: 35 / 0 / 35
Local bandwidth <effective/available>: 35000000000 (35000000000) kbps
MAC address (source): ea3b.745f.c4b0 (Chassis pool)
Inter-chassis link: No
Minimum active links / bandwidth: 1 / 1 kbps
Maximum active links: 64
Wait while timer: 2000 ms
Load balancing: Default
LACP: Operational
Flap suppression timer: Off
Cisco extensions: Disabled
Non-revertive: Disabled
mLACP: Not configured
IPv4 BFD: Not configured

```

Port	Device	State	Port ID	B/W, kbps
Hu0/4/0/0	Local	Active	0x8000, 0x0009	100000000
Link is Active				
Hu0/4/0/1	Local	Active	0x8000, 0x000a	100000000
Link is Active				
- - -				
Hu0/4/0/35	Local	Active	0x8000, 0x002b	100000000
Link is Active				

- Verify the bundle statistics:

```

router# show policy-map interface bundle-ether 12000

```

```

Bundle-Ether12000 input: ingress

```

```

Class inet4-classifier-af1
Classification statistics
Matched : 4647401962/21236124455654 26403040
Transmitted : 4647401962/21236124455654 26403040
Total Dropped : 0/0 0
Class inet4-classifier-af2

```



```

Classification statistics          (packets/bytes)    (rate - kbps)
  Matched                        :      4502980177/20576584333939      25571493
  Transmitted                    :      4502980177/20576584333939      25571493
  Total Dropped                  :                      0/0              0
Class inet4-classifier-af3
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      4647404125/21236213667880      26389086
    Transmitted                  :      4647404125/21236213667880      26389086
    Total Dropped                :                      0/0              0
Class inet4-classifier-af4
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      9291188840/42456120548683      52771168
    Transmitted                  :      9291188840/42456120548683      52771168
    Total Dropped                :                      0/0              0
Class inet4-classifier-bel
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      4647413429/21235847852686      26393414
    Transmitted                  :      4647413429/21235847852686      26393414
    Total Dropped                :                      0/0              0
Class inet4-classifier-ncl
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      9294887621/42473100149807      52778258
    Transmitted                  :      9294887621/42473100149807      52778258
    Total Dropped                :                      0/0              0

Class class-default
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :                      0/0              0
    Transmitted                  :                      0/0              0
    Total Dropped                :                      0/0              0

Bundle-Ether12000 output: egress

Class c1
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      16665494532/75878118942463      8760591
    Transmitted                  :      16655834643/75834136022017      8760591
    Total Dropped                :           9659889/43982920446          0
  Queueing statistics
    Queue ID                     : None (Bundle)
    Taildropped(packets/bytes)    : 9659889/43982920446
Class c2
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      16665421959/75877849543188      8718687
    Transmitted                  :      16665421959/75877849543188      8718687
    Total Dropped                :                      0/0              0
  Queueing statistics
    Queue ID                     : None (Bundle)
    Taildropped(packets/bytes)    : 0/0
Class c3
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      16665247833/75877509455458      8703470
    Transmitted                  :      16665187414/75877234624197      8703470
    Total Dropped                :           60419/274831261              0
  Queueing statistics
    Queue ID                     : None (Bundle)
    Taildropped(packets/bytes)    : 60419/274831261
Class c4
  Classification statistics          (packets/bytes)    (rate - kbps)
    Matched                      :      33330896131/151755393012945      17470745
    Transmitted                  :      33330745421/151754709368565      17470745
    Total Dropped                :          150710/683644380              0
  Queueing statistics
    Queue ID                     : None (Bundle)

```

```

    Taildropped(packets/bytes)           : 150710/683644380
Class c5
  Classification statistics              (packets/bytes)    (rate - kbps)
    Matched                             : 16878910340/76849791869834    8833394
    Transmitted                          : 16878849464/76849514633309    8833394
    Total Dropped                        : 60876/277236525              0
  Queueing statistics
    Queue ID                             : None (Bundle)
    Taildropped(packets/bytes)           : 60876/277236525
Class c6
  Classification statistics              (packets/bytes)    (rate - kbps)
    Matched                             : 33330898844/151756094112925    17456785
    Transmitted                          : 33330752668/151755427708382    17456785
    Total Dropped                        : 146176/666404543              0
  Queueing statistics
    Queue ID                             : None (Bundle)
    Taildropped(packets/bytes)           : 146176/666404543
Class c7
  Classification statistics              (packets/bytes)    (rate - kbps)
    Matched                             : 244106/79922040                74
    Transmitted                          : 244106/79922040                74
    Total Dropped                        : 0/0                          0
  Queueing statistics
    Queue ID                             : None (Bundle)
    Taildropped(packets/bytes)           : 0/0
Class class-default
  Classification statistics              (packets/bytes)    (rate - kbps)
    Matched                             : 267075066180/1215993441123215  139917482
    Transmitted                          : 267075066180/1215993441123215  139917482
    Total Dropped                        : 0/0                          0
  Queueing statistics
    Queue ID                             : None (Bundle)
    Taildropped(packets/bytes)           : 0/0

```

Related Topics

- [QoS on Link Bundles, on page 147](#)

Associated Commands

- `bundle maximu-active links`
- `interface Bundle-Ether`



CHAPTER 8

Configuring Hierarchical Modular QoS

Hierarchical QoS (H-QoS) is a QoS model that enables you to specify QoS behavior at multiple levels of hierarchy. This chapter provides information about this feature and the different steps involved in configuring it.

Table 27: Feature History for Hierarchical Modular QoS

Release	Modification
Release 6.3.1	Initial Release

This chapter covers the following topics:

- [Overview of Hierarchical Modular QoS, on page 153](#)
- [Restrictions for Configuring H-QoS, on page 154](#)
- [Configuring Hierarchical Queuing, on page 155](#)
- [Conform Aware Hierarchical Policy Overview, on page 161](#)

Overview of Hierarchical Modular QoS

Hierarchical QoS (H-QoS) allows you to specify QoS behavior at multiple policy levels, which provides a high degree of granularity in traffic management.

H-QoS is applied on the router interface using nested traffic policies. The first level of traffic policy, the parent traffic policy, is used for controlling the traffic at the main interface or sub-interface level. The second level of traffic policy, the child traffic policy, is used for more control over a specific traffic stream or class. The child traffic policy, is a previously defined traffic policy, that is referenced within the parent traffic policy using the **service-policy** command.

Two-level H-QoS is supported on both ingress and egress directions on all line cards and on physical or bundle main interfaces and sub-interfaces.

Three-level Hierarchical QoS (H-QoS) enables enforcement of class/service, group/ Ethernet Flow Point (EFP), and port level SLAs. You can apply regular two-level egress H-QoS policies on the sub-interfaces to achieve class and EFP SLAs at child and parent levels. In addition, you can apply a port shaper policy on the main interface to achieve an aggregated port level SLA in a 1+2 H-QoS or three-level H-QoS model.

An important point to note is that before Release 6.6.25 (where the three-level H-QoS capability was introduced), when you applied class-default shaper on a main interface, it was enforced *only* on the traffic going through the main interface. With three-level H-QoS, a class default shaper that is applied on the main

interface is considered as a port shaper and enforced on *all* traffic going out of that physical port. The advantage of three-level H-QoS is that the parent shaper on the sub-interfaces is allowed to oversubscribe, thus enabling best effort sharing of the aggregate port shaper at the third level.



Note From Cisco IOS XR Release 7.3.1 onwards, systems with Cisco NC57 line cards running in compatibility mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

Starting with Cisco IOS XR Release 7.4.1 systems with Cisco NC57 line cards running in native mode support QoS over Layer 2 services for:

- Local switching [xconnect or bridging]
- L2 VPN – VPWS

H-QoS support for EVPN ELAN on Cisco Network Convergence System 5700 Series Routers

Table 28: Feature History Table

Feature Name	Release Information	Feature Description
H-QoS support for EVPN ELAN on Cisco Network Convergence System 5700 Series Routers	Release 7.7.1	<p>You can now apply Hierarchical Modular QoS (H-QoS) on multipoint EVPN service (EVPN ELAN) traffic. You can use H-QoS to specify QoS behavior at multiple levels of hierarchy and allocate specific bandwidth and services to specific users.</p> <p>This feature is now supported on NCS 5700 Series routers operating in native mode (refers to the mode when the router has only NCS 5700 line cards).</p> <p>To enable native mode, use hw-module profile npu native-mode-enable in the configuration mode and reload the router. You can verify this configuration in the show run in hw command output.</p>

Restrictions for Configuring H-QoS

The following restrictions are applicable while configuring H-QoS:

1. The parent traffic policy only supports the traffic class of type class-default.
2. The parent traffic policy only supports the class-action **shape** and no other queuing action can be configured in it.
3. While configuring on the router, it is mandatory that the priority class must have traffic shaper in the child traffic policy.

4. The sum of the bandwidth of the child policies must be less than the parent policy's traffic shaper.
5. For congestion avoidance and management, the traffic shaper in the parent traffic policy calculates the queue limit and drop priority.
6. H-QoS profile and ingress peering profile don't work simultaneously. Hence, features requiring a peering profile also do not work with the H-QoS profile enabled.
7. PBTS feature does not work when the H-QoS profile is enabled. This is due to TCAM limitations.
8. A maximum of 896 bundle sub-interfaces are only supported in the system, even if there are no QoS policies applied. This is due to an internal LAG_ID resource consumption in HQoS profile mode for bundle sub-interfaces with or without QoS policies being applied.
9. A maximum of 4 priority levels are only supported in HQoS profile mode unlike the default mode where 7-priority levels are supported. The restriction also applies to physical and bundle main interface policies where 7-level priorities were previously used in non-H-QoS profile mode.
10. Bandwidth and Bandwidth remaining configurations are not supported simultaneously within the same policy-map. If a class has bandwidth (CIR), other classes must also have only bandwidth configuration. If a class-map has bandwidth remaining percent/ratio (EIR), other classes should also have only the bandwidth remaining configuration. Shaping is applied on any class.
11. Priority classes must have rate limit configuration by using a Shaping configuration. The effective shaper value is taken as priority bandwidth reservation. Sum of priority bandwidth reservations across all sub-interfaces and main interfaces must not exceed the network interface (NIF) port speed. This is to avoid over-subscription of priority traffic across the network interface port.

Rates of non-priority classes and parent shaping can be over-subscribed.
12. The granularity of bandwidth or bandwidth remaining ratio (BRR) is 1:64 as compared to 1:4096 in non-hqos mode. So, there could be accuracy differences in bandwidth performance based on the values used.
13. Filtering for egress IPv4 and IPv6 multicast traffic is not supported if H-QoS is configured on the router.

The following restrictions are applicable while configuring three-level H-QoS:

- There is no support for bandwidth action at the EFP parent level. All EFP/sub-interface policies get a fair share of the port shaper.
- Three-level H-QoS does not apply to ingress policies or to egress marking policies.
- Executing **clear qos counters** on the main interface clears only the main interface policy statistics. Use the "all" option to clear all sub-interface statistics or alternately, clear the sub-interface policy statistics individually.
- Main interface policy statistics do not reflect the sub-interface packet / byte counters, although the port shaper is enforced on all logical ports for a given physical interface. The sub-interface policy-map statistics reflect the transmitted and dropped packet/byte count post-port shaper enforcement.

Configuring Hierarchical Queuing

Before you configure H-QoS, you must enable the H-QoS profile on the router. After enabling H-QoS profile, reload the router, as shown in the following configuration.

```

admin
hw-module location all reload
Router# configure
Router(config)# hw-module profile qos hqos-enable
Router(config)# commit
Router# admin
sysadmin-vm:0_RP0# hw-module location all reload

```

The steps that are involved in configuring hierarchical queuing are as follows:

1. Configure a class-map.
2. Configure a child traffic policy using the class-map that was configured in the previous step.
3. Configure a parent traffic policy and add the child traffic policy in it.

The parent traffic policy is the H-QoS traffic policy and it can be applied on physical or bundle main interfaces and sub-interfaces.

Configuration Example

Configuration of a class-map is as follows:

```

Router# configure
Router(config)# class-map match-any tc2
Router(config-cmap)# match traffic-class 1
Router(config-cmap)# end-class-map
Router(config)# commit

```

Configuration of a child traffic policy is as follows:

```

Router# configure
Router(config)# policy-map child
Router(config-pmap)# class tc2
Router(config-pmap-c)# shape average percent 20
Router(config-pmap-c)# exit
Router(config-pmap)# class class-default
Router(config-pmap-c)# shape average percent 1
Router(config-pmap)# end-policy-map
Router(config)# commit

```

Configuration of a parent traffic policy is as follows:

```

Router# configure
Router(config)# policy-map parent
Router(config-pmap)# class class-default
Router(config-pmap-c)# service-policy child
Router(config-pmap-c)# shape average percent 50
Router(config-pmap)# end-policy-map
Router(config)# commit

```

Running Configuration

```

/* Configuration of a Class-map */
class-map match-any tc2
  match traffic-class 1
end-class-map
!

```

```

/* Configuration of a Child Traffic Policy */
policy-map child
  class tc2
    shape average percent 20
  !
  class class-default
    shape average percent 1
  !
end-policy-map
!
/* Configuration of a Parent Traffic Policy */
policy-map parent
  class class-default
    service-policy child
    shape average percent 50
  !
end-policy-map
!

```

Applying the Parent Traffic Policy on a Main Interface

```

Router# configure
Router(config)# interface TenGigE 0/0/0/10
Router(config-int)# service-policy output parent
Router(config-int)# commit

```

Applying the Parent Traffic Policy on a Sub-interface

```

Router# configure
Router(config)# interface TenGigE 0/0/0/10.1
Router(config-int)# service-policy output parent
Router(config-int)# commit

```

Verification

Verify if the H-QoS traffic policy is applied correctly on the interface using the commands **show qos interface interface-name output**. In the following example, the **Level1 Class** gives information about the class-map that is associated with the parent traffic policy and the **Level2 Class** gives information about the class-maps that are associated with the child traffic policy.

```
RP/0/RP0/CPU0:ios#show qos interface ten0/0/0/10 output
```

```

NOTE:- Configured values are displayed within parentheses
Interface TenGigE0/0/0/10 ifh 0x1e0 -- output policy
NPU Id:                                0
Total number of classes:                3
Interface Bandwidth:                    10000000 kbps
VOQ Base:                               1136
Accounting Type:                        Layer1 (Include Layer 1 encapsulation and above)
-----

```

Level1 Class	=	class-default
Queue Max. BW.	=	no max (50 %)
Queue Min. BW.	=	0 kbps (default)
Inverse Weight / Weight	=	0 / (BWR not configured)
Level2 Class	=	tc2
Egressq Queue ID	=	1138 (LP queue)
Queue Max. BW.	=	1020015 kbps (20 %)
Queue Min. BW.	=	0 kbps (default)
Inverse Weight / Weight	=	1 / (BWR not configured)
Guaranteed service rate	=	1000000 kbps

```

TailDrop Threshold                = 1253376 bytes / 10 ms (default)
WRED not configured for this class
Level2 Class                      = class-default
Egressq Queue ID                 = 1136 (Default LP queue)
Queue Max. BW.                   = 50625 kbps (1 %)
Queue Min. BW.                   = 0 kbps (default)
Inverse Weight / Weight           = 1 / (BWR not configured)
Guaranteed service rate           = 50000 kbps
TailDrop Threshold               = 62720 bytes / 10 ms (default)
WRED not configured for this class

```

The statistics for the packets that have matched the different traffic classes of the parent and child traffic policies can be viewed using the command **show policy-map interface interface-name output**. Also, this command also shows the number of packets that are transmitted or dropped when the specified action is applied on the packets that have matched the respective traffic class.

```
Router# show policy-map interface ten0/0/0/10 output
```

```

TenGigE0/0/0/10 output: parent
Class class-default
  Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                          : 2313578823/296138089344    8494665
  Transmitted                      : 232805738/29799134464      854465
  Total Dropped                    : 2080773085/266338954880    7640200
Policy child Class tc2
  Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                          : 2313578823/296138089344    8494665
  Transmitted                      : 232805738/29799134464      854465
  Total Dropped                    : 2080773085/266338954880    7640200
Queueing statistics
  Queue ID                        : 1138

  Taildropped(packets/bytes)       : 2080773085/266338954880
Policy child Class class-default
  Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                          : 0/0                      0
  Transmitted                      : 0/0                      0
  Total Dropped                    : 0/0                      0
Queueing statistics
  Queue ID                        : 1136

  Taildropped(packets/bytes)       : 0/0

```

When using hierarchical policers, there is no independent set of hardware counters to store the parent policer statistics. Instead, parent policer statistics are manipulated in the software to be the sum of all child policers under the same policy-map.

This is shown in the following example where two streams of traffic, with CoS value of 1 and 2 are sent at a speed of 3.5 Gbps each.

```

/*Hierarchical Policy Map Configuration*/
=====
Router# show running-config policy-map Hingress
policy-map Hingress
  class class-default
    service-policy ingress
    police rate 5 gbps peak-rate 9 gbps
  !
end-policy-map
!
/*Ingress Policy Map Configuration*/
=====

```



```

Router#show running-config policy-map ingress
policy-map ingress
  class cos1
    set traffic-class 1
    police rate 5 gbps
  !
  !
  class cos2
    set traffic-class 2
    police rate 5 gbps
  !
  !
  class class-default
  !
end-policy-map
!
/*Policy Map applied at TenGigE0/0/0/6.100 Interface*/
=====
Router#show policy-map interface tenGigE 0/0/0/6.100 input

TenGigE0/0/0/6.100 input: Hingress

Class class-default
  Classification statistics          (packets/bytes)      (rate - kbps)
    Matched                        :      856717937/109659895936      6683676
    Transmitted                    :      856717937/109659895936      6683676
    Total Dropped                  :              0/0                  0
  Policing statistics              (packets/bytes)      (rate - kbps)
    Policed(conform)               :      856717937/109659895936      6683674
    Policed(exceed)                :              0/0                  0
    Policed(violate)               :              0/0                  0
    Policed and dropped            :              0/0
Policy ingress Class cos1
  Classification statistics          (packets/bytes)      (rate - kbps)
    Matched                        :      437826303/56041766784      3341838
    Transmitted                    :      437826303/56041766784      3341838
    Total Dropped                  :              0/0                  0
  Policing statistics              (packets/bytes)      (rate - kbps)
    Policed(conform)               :      437826303/56041766784      3341838
    Policed(exceed)                :              0/0                  0
    Policed(violate)               :              0/0                  0
    Policed and dropped            :              0/0
    Policed and dropped(parent policer) : 0/0
Policy ingress Class cos2
  Classification statistics          (packets/bytes)      (rate - kbps)
    Matched                        :      418891634/53618129152      3341838
    Transmitted                    :      418891634/53618129152      3341838
    Total Dropped                  :              0/0                  0
  Policing statistics              (packets/bytes)      (rate - kbps)
    Policed(conform)               :      418891634/53618129152      3341838
    Policed(exceed)                :              0/0                  0
    Policed(violate)               :              0/0                  0
    Policed and dropped            :              0/0
    Policed and dropped(parent policer) : 0/0
Policy ingress Class class-default
  Classification statistics          (packets/bytes)      (rate - kbps)
    Matched                        :              0/0                  0
    Transmitted                    :              0/0                  0
    Total Dropped                  :              0/0                  0
Policy Bag Stats time: 0

```

Policy Bag Stats time: 0

Configuration Example for Three-Level H-QoS

To configure three-level H-QoS:

1. Configure the port shaper or EFP group shaper.
2. Configure the two-level H-QoS policy with an EFP parent shaper and class or service level actions
3. Enable the port or EFP group shaper on the main interface to become the root policy.
4. Enable the two-level H-QoS policy on each EFP instance, thus providing a three-level hierarchy of service, EFP and EFP group or port SLAs.

A sample configuration for three-level H-QoS is as follows.

```

policy-map port_shaper
  class class-default
    shape average 6 gbps
  !
end-policy-map
!

policy-map efp_policy
  class class-default
    service-policy efp_policy_child
    shape average 4 gbps
  !
end-policy-map
!

policy-map efp_policy_child
  class tc1
    shape average 50 mbps
    priority level 1
  !
  class tc2
    bandwidth percent 50
  !
  class tc3
    bandwidth percent 30
  !
  class class-default
  !
end-policy-map
!

interface TenGigE0/5/0/4
  service-policy output port_shaper
!

interface TenGigE0/5/0/4.1
  service-policy output efp_policy
  encapsulation dot1q 11
!

interface TenGigE0/5/0/4.2
  service-policy output efp_policy
  encapsulation dot1q 12
!


```

Verification

Run the **show policy-map interface** command in XR EXEC mode to display the packet/byte count and rate post port shaper enforcement on each of the sub-interface/EFP policies.

Conform Aware Hierarchical Policy Overview

Hierarchical QoS (H-QoS), while allowing for granular and multi-level management of traffic, does not allow for conform traffic from a child-level policy to a parent-level policy to get priority. This means that in case of excess traffic, the parent policer drops conform traffic packets from the child level as well.



The conform-aware hierarchical policy feature enables the parent-level policy to prioritize conform traffic from child-level policy over exceed and violate traffic.

Here is how it works: the child-level policer initially marks its packets as red, yellow, or green. Packets are marked based on the committed information rate (CIR) value and the two associated burst sizes - committed burst size (CBS) and excess burst size (EBS). If a packet does not exceed the CBS, it is marked as conform packet (green). The packet is marked as exceeded if it exceeds CBS, but not the EBS (yellow). If it exceeds the EBS as well, it is marked as violate packet (red).

When the packets arrive at the parent level policer (which is color aware), the policer prioritizes the packets marked green over the packets marked yellow. After all the conform traffic (green) is transmitted and there are tokens available still, the yellow packets are transmitted next, instead of being marked as violate traffic (red).

To enable the conform-aware hierarchical policy feature run the command.

```
hw-module profile qos conform-aware-policer
```

Conform Aware Hierarchical Policy Restrictions

The parent-policy traffic police rate must be greater than or equal to the sum of child conform rates.

Configuring Conform Aware Hierarchy Policy

To enable and configure shared policer:

1. Run the **hw-module profile qos conform-aware-policer** command.

2. Reload the affected line card.
3. Configure class maps to be used for matching packets to the class specified.
4. Create a child policy map.
5. Configure traffic policing for the classes in the child policy map.
6. Attach the child policy-map to the parent's class-default class.
7. Configure traffic policing for the parent policy map.

```
RP/0/RP0/CPU0:ios(config)#hw-module profile qos conform-aware-policer
RP/0/RP0/CPU0:router# reload location 0/0/CPU0
RP/0/RSP0/CPU0:ios(config)#class-map match-any c1
RP/0/RSP0/CPU0:ios(config-cmap)#match precedence 1
RP/0/RSP0/CPU0:ios(config-cmap)#end-class-map
RP/0/RSP0/CPU0:ios(config)#class-map match-any c2
RP/0/RSP0/CPU0:ios(config-cmap)#match precedence 2
RP/0/RSP0/CPU0:ios(config-cmap)#end-class-map
RP/0/RSP0/CPU0:ios(config-pmap)# policy-map child
RP/0/RSP0/CPU0:ios(config-pmap-c)#class class1
RP/0/RSP0/CPU0:ios(config-pmap-c-police)#police rate 20 mbps peak-rate 50 mbps
RP/0/RSP0/CPU0:ios(config-pmap-c)#class class2
RP/0/RSP0/CPU0:ios(config-pmap-c-police)#police rate 30 mbps peak-rate 60 mbps
RP/0/RSP0/CPU0:ios(config-pmap)#policy-map parent
RP/0/RSP0/CPU0:ios(config-pmap-c)#class class-default
RP/0/RSP0/CPU0:ios(config-pmap-c)#service-policy child
RP/0/RSP0/CPU0:ios(config-pmap-c-police)#police rate 60 mbps
RP/0/RSP0/CPU0:ios(config-pmap-c-police)#exit
RP/0/RSP0/CPU0:ios(config-pmap-c)#exit
RP/0/RSP0/CPU0:ios(config-pmap)#exit
RP/0/RSP0/CPU0:ios(config-if)#commit
```

Running Configuration

```
class-map match-any class1
match precedence 1
end-class-map

class-map match-any class2
match precedence 2
end-class-map

policy-map child
class class1
    police rate 20 mbps peak-rate 50 mbps
    !
class class2
    police rate 30 mbps peak-rate 60 mbps
    !

policy-map parent
class class-default
    service-policy child
    police rate 60 mbps
    !
    !
```

In this example, the 20 mbps rate in child class1 and 30 mbps rate in child class2 are guaranteed as conform traffic and the parent will not drop them.

Verification

Run the **show policy-map interface** command in XR EXEC mode to confirm that the committed information rate (CIR) is prioritized over the peak information rate (PIR).

In the example below, **Policed (conform)** or CIR from each class is prioritized over **Policed (exceed)** or PIR.

```
RP/0/RP0/CPU0:ios(config-if)#do show running-config policy-map parent_conform
Tue Apr 28 11:19:09.785 UTC
policy-map parent_conform
class class-default
  service-policy child_conform
  police rate 80 mbps peak-rate 120 mbps
  !
!
end-policy-map
!
```

```
RP/0/RP0/CPU0:ios(config-if)#do show running-config policy-map child_conform
Tue Apr 28 11:19:14.219 UTC
policy-map child_conform
class match_prec1
  police rate 30 mbps peak-rate 35 mbps
  !
  set qos-group 3
!
class match_prec2
  police rate 20 mbps peak-rate 25 mbps
  !
  set qos-group 4
!
class class-default
  police rate 20 mbps peak-rate 30 mbps
  !
!
end-policy-map
!
```

```
RP/0/RP0/CPU0:ios(config-if)#do show policy-map interface tenGigE 0/0/0/11/2 input
Tue Apr 28 11:23:48.915 UTC

TenGigE0/0/0/11/2 input: parent_conform

Class class-default
  Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                          :      217797200/111512166400      2344847
  Transmitted                      :      8314388/4256966656      88089
  Total Dropped                    :      209482812/107255199744      2256758

  Policing statistics              (packets/bytes)      (rate - kbps)
  Policed(conform)                 :      6602174/3380313088      69926
  Policed(exceed)                  :      1712214/876653568      18165
  Policed(violate)                 :      209482812/107255199744      2256782
  Policed and dropped              :      209482812/107255199744

Policy child_conform Class match_prec1
  Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                          :      54449300/27878041600      586215
  Transmitted                      :      3246813/1662368256      34399
  Total Dropped                    :      51202487/26215673344      551816
  Policing statistics              (packets/bytes)      (rate - kbps)
  Policed(conform)                 :      2818471/1443057152      29851
  Policed(exceed)                  :      428342/219311104      4547
  Policed(violate)                 :      51202487/26215673344      551816
```

```

Policed and dropped :          51202487/26215673344
Policed and dropped(parent policer) : 0/0

Policy child_conform Class match_prec2
Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                        :          54449300/27878041600      586213
  Transmitted                    :          2319731/1187702272      24577
  Total Dropped                  :          52129569/26690339328      561636
Policing statistics                (packets/bytes)      (rate - kbps)
  Policed(conform)               :          1891851/968627712      20037
  Policed(exceed)                :          427880/219074560      4540
  Policed(violate)               :          52129569/26690339328      561636
  Policed and dropped            :          52129569/26690339328
  Policed and dropped(parent policer) : 0/0

Policy child_conform Class class-default
Classification statistics          (packets/bytes)      (rate - kbps)
  Matched                        :      108898600/55756083200      1172419
  Transmitted                    :      2747844/1406896128      29113
  Total Dropped                  :      106150756/54349187072      1143306
Policing statistics                (packets/bytes)      (rate - kbps)
  Policed(conform)               :      1891852/968628224      20036
  Policed(exceed)                :      855992/438267904      9076
  Policed(violate)               :      106150756/54349187072      1143306
  Policed and dropped            :      106150756/54349187072
  Policed and dropped(parent policer) : 0/0
Policy Bag Stats time: 1588072996538

```

Related Commands

hw-module profile qos conform-aware-policer



CHAPTER 9

Configure Egress Traffic Management

This chapter covers the following topics:

- [Egress Traffic Management, on page 166](#)
- [More About the Egress Traffic Management Model, on page 168](#)
- [QoS Policy Configuration Rules, on page 169](#)
- [Configure Egress Traffic Management , on page 170](#)
- [Low Rate Shaper, on page 186](#)

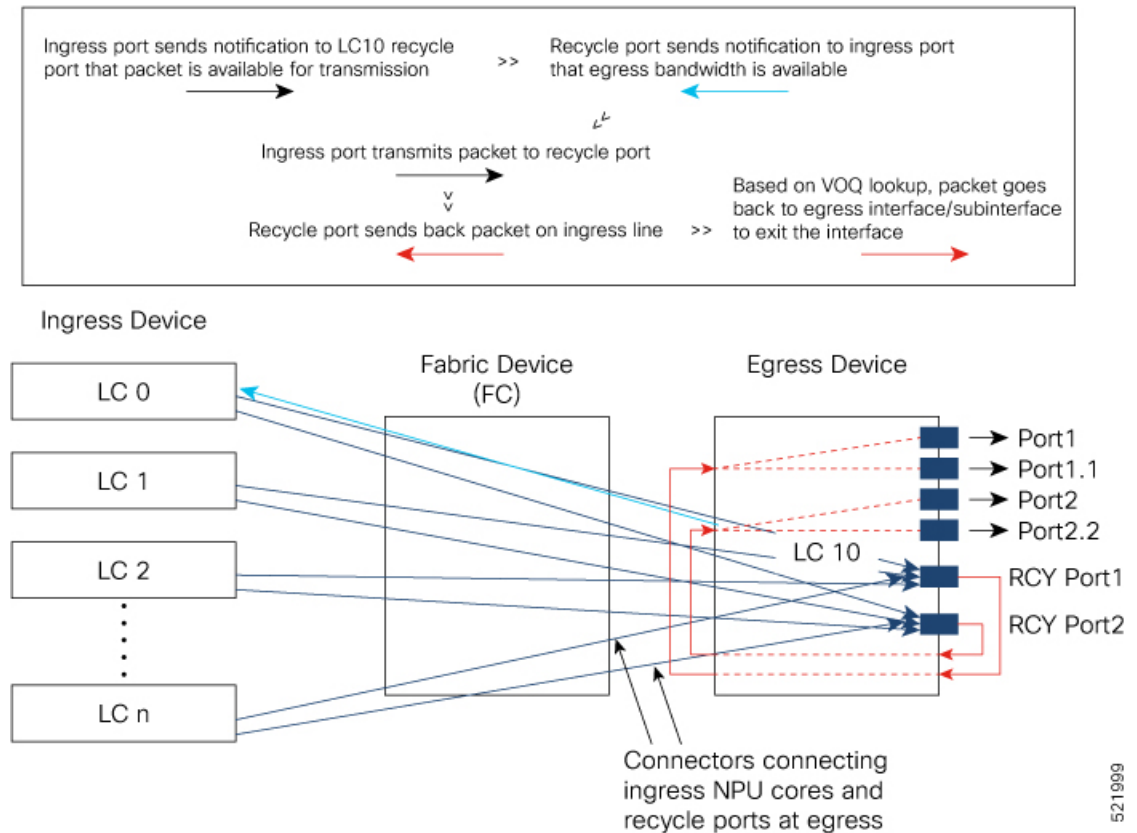
Egress Traffic Management

Table 29: Feature History Table

Feature Name	Release Information	Feature Description
Egress Traffic Management	Release 7.6.1	<p>The egress traffic management model employs a two-pass model per NPU core to get up to 50% of NPU core bandwidth as recycle bandwidth. VOQ lookup takes place in the second pass to get the outgoing interface or sub-interface VOQ. This lookup at the second pass enhances the egress queuing policy by making the feature-rich ingress pipeline available to the outgoing traffic.</p> <p>Also, this functionality is configurable at the port level and not at the global level. This approach limits the usage of VOQ connectors required on all ingress devices to only physical ports or main interfaces, thus enabling scaling of egress QoS policies.</p> <p>This functionality introduces the mode etm option for the controller optics command.</p>

The egress traffic management model modifies the existing [Ingress Traffic Management Model, on page 116](#) to ensure you can configure higher QoS scales. It does so by using a recycle model at egress NPU so that the packet buffering takes place at egress NPU. (See figure, **The Egress Traffic Management Model**.) Also, egress traffic management is configurable at the port-level and not at the global level, to maximize flexibility.

Figure 7: The Egress Traffic Management Model



In the figure, every ingress NPU port (LC 0 through LC n) has eight VOQs for the single egress NPU LC 10.

Here's how packet transmission takes place:

1. When a packet arrives at ingress port (say on LC 0), the forwarding lookup on ingress line card points to the egress interface that could be physical, subinterface, Link aggregation Group (LAG), or LAG subinterface. In the case of LAG, hashing takes place at the ingress line card to select the LAG member from where the packet will be sent out. Based on the egress interface (say it is on LC10), the packet is enqueued to the VOQ of LC 10. The egress interface is always mapped to a physical port.
2. However, unlike in the ingress traffic management model where the ingress NPU ports connect to the egress NPU ports (per traffic management hierarchy), the egress traffic management model connects ingress NPU ports to the recycle port that's mapped to the egress port.
3. Once egress bandwidth is available, the recycle ports (based on the packet marking and distribution model) send grants to the ingress ports via the connectors. (The figure shows a separate line for the grant for the sake of visual representation. In reality, the same connector is used for requests, grants, and transmission between an NPU core at the ingress and the corresponding recycle port.)
4. The ingress NPU responds to this permission by transmitting the packet via the fabric device (FC) to the egress NPU.
5. Once the packet reaches egress NPU, it gets recycled back to the ingress pipeline of the egress NPU. A VOQ lookup takes place at this stage that checks if there's a policy on the subinterface. If there's no policy on the subinterface, the main interface VOQ is taken. This behavior is the same as in the ingress traffic

management model, except that another VOQ lookup for subinterface and the main interface takes place in the second pass (recycle) on the egress NPU core.

6. Based on the lookup in the second pass, the packets go back into the egress to the specified main interface or subinterface. It's also in the second pass that any queueing action kicks in before the packets go out the egress interface.

More About the Egress Traffic Management Model

• Which routers and line cards support the egress traffic management model?

The egress traffic management model is supported on routers that have the following Cisco NC57 line cards that are installed and that operate in the native mode:

- NC57-18DD-SE
- NC57-36H-SE

This feature is also supported on the following Cisco NCS-57B1 chassis variants:

- NCS-57B1-5D24H-SE
- NCS-57B1-5DSE-SYS

• What benefits does the egress traffic management model offer?

- Configuring the egress traffic management functionality offers you **higher scales of egress QoS policies**. And this is because of the way the egress traffic management model is designed, reducing the size of the mesh of connectors to main and subinterface ingress ports.
- Plus, because of this model (where a second lookup and queueing action takes place in the recycled pass), the **egress queueing policy is enhanced**. An egress traffic management queueing policy contains the same match criteria that ingress policies typically contain. (See the following table for details.) In addition, match on traffic-class and remarking traffic-class is also supported in egress with the **hw-module profile qos ipv6 short-etm** command. (For details, see [Configure Egress Traffic Management](#), on page 170.) This means that the egress traffic management model offers you with even more flexibility in planning your egress policies, while supporting higher scale numbers.

Flow Type	Egress Classification Type Supported
IPv4/IPv6	<ul style="list-style-type: none"> • Prec / DSCP • SIP / DIP / SP / DP / Protocol / TCP-CTL via ACLs • COS / DEI (for VLAN) • Fragmentation
MPLS	<ul style="list-style-type: none"> • IPv4 ToS / IPv6 TC (for underlying for IPv4/IPv6) • COS / DEI (for VLAN) • MPLS Exp

Flow Type	Egress Classification Type Supported
Layer 2	<ul style="list-style-type: none"> • COS / DEI (for VLAN) • COS Inner / DEI Inner (for VLAN) • IPv4 ToS • IPv6 TC • MPLS Exp

- There's another important benefit that the egress traffic management model offers. You need higher QoS scale and features only for access- or edge-facing interfaces, and not on core-facing interfaces. The egress traffic management model provides the **flexibility to configure recycle mode only on access facing ports** while continue using core facing ports in single-pass mode.
- Configuring the egress traffic management functionality enables the low rate shaper feature. For details, see [Low Rate Shaper, on page 186](#).
- **In the recycle pass, packets are sent back on the ingress pipeline. What's the impact?**—The impact is **reduced throughput**. Recycle traffic competes with ingress traffic on the ingress pipeline; this impacts the traffic throughput. Thus, higher scales come at the cost of reduced throughputs. In such contention cases, device throughput drops down to 50%.
- **How do I handle Link Aggregation Group (LAG) configurations?**—All LAG members must have the same mode. This means that you must **configure the egress traffic management mode on all LAG members**.
- **Which MDB profiles are supported for the egress traffic management model?**—All SE-based MDB profiles are supported for the egress traffic management model. (SE-based MDB profiles are the scaled versions that allow for configuring more QoS policies on systems using eTCAM.)
- **Important Note on LOCP Statistics**—When you attach a policy map to a subinterface in the egress traffic management model, the Locally Originated Control Protocol (LOCP) packets exiting that subinterface are accounted for in the subinterface policy counters.

Without the egress traffic management model, egress inject LOCP packets are not accounted for in sub-interface egress policy counters.

QoS Policy Configuration Rules

- The **set traffic-class** configuration for all traffic classes is mandatory at egress for egress traffic management ports, regardless of your ingress configuration. This **set traffic-class val** represents the VOQ offset that the matching packets select. The only exception to this rule is the class-default configuration.
- The class-default can have a **set traffic-class val** other than 0.
- Even if the ingress and egress traffic classes match, you'll still need to configure a **set traffic-class val** at egress. If the match and set values differ, then the traffic-class is remarked. For example, say a traffic-class is set to 1 in the ingress policy and the match on that traffic-class is 1 on the egress in the

egress traffic management policy. Then, say, at the second pass, the **set traffic-class val** is remarked to 5. The packets that match this class in the egress traffic management policy go to VOQ offset 5.

- You can set traffic-class for multiple traffic classes to the same value. In such a situation, traffic matching those classes head to the same VOQ. Therefore, the statistics for all these classes point to the same VOQ counters. It's important to note that when multiple traffic classes have the same value in the **set traffic-class val** action, then the queuing parameters (for example, priority or bandwidth remaining, queue depth, WRED) of these classes **MUST** be the same.
- In a hierarchical policy, you can configure the **set traffic-class val** only for the child traffic classes. The parent classes can't have the **set traffic-class val** action.

Configure Egress Traffic Management

Before You Begin

Some critical points to note before you begin:

- Before you proceed any further, ensure that you delete your existing **hw-module profile qos hqos-enable** configuration and restart your system. Only then can you enable egress traffic management at the port level. H-QoS profiles and egress traffic management configurations can't coexist.
- The egress traffic management mode is configured on a per-port mode. This ensures that the core-facing ports continue using the current ingress-based VOQ model, while the access ports use the egress traffic management model. Such an approach reduces the consumption of the recycle bandwidth. See the following table for details.

Table 30: Global H-QoS versus Per Port Configuration

Global H-QoS with hw-module profile qos hqos-enable		Per Port Egress Traffic Management (ETM) without hw-module profile qos hqos-enable			
All Ports		ETM ports with controller optics mode etm configured		Non-ETM ports	
Main interface	Subinterface	Main interface	Subinterface	Main interface	Subinterface
H-QoS and flat policies	H-QoS and flat policies	H-QoS and flat policies	H-QoS and flat policies	H-QoS and flat policies	No policies, inherit from main

- The configuration is implemented at the controller level. This means that when you commit the egress traffic management configuration, **all existing configurations are deleted. All existing features and subinterface configurations are deleted and you must reconfigure all of them.**

Configuration Example for Policy Map with Class Match on traffic-class

To configure an egress traffic management queueing policy and reconfigure policies in **etm** mode for traffic class:

1. Delete the existing **hw-module profile qos hqos-enable** configuration and restart your system.

2. Enable egress traffic management (**etm**) on an interface using the **controller optics** command in the optics controller configuration mode. For 10 GbE interfaces, configure **etm** on the breakout port. After this configuration is enabled, it deletes all subinterfaces under the main interface.
3. Configure the subinterfaces. For policy maps with **match-all** specified for **traffic-class**, run **hw-module profile qos ipv6 short-etm** to truncate IPv6 fields. This configuration frees TCAM space for packet header marking for VOQ lookup in the second pass. Reload the line card.
4. Create the class maps.
5. Specify the match criteria, which can be: match on traffic class, DSCP, CoS, or any combination of these criteria.
6. Set the traffic classes for each class in the policy map.



Note The **class-default** need not have a **set-traffic class** value.

7. Specify the traffic class actions.
8. Attach the policy map to the subinterface.

```
/*Delete existing h-qos configuration*/
Router#configure terminal
Router(config)#no hw-module profile qos hqos-enable
In order to activate this new qos profile, you must manually reload the chassis/all line
cards
Router(config)#commit
Router(config)#exit
Router#reload

/*Enable etm on an interface*/
Router#configure terminal
Router(config)#controller optics 0/5/0/16
Router(config-Optics)#mode etm
Router(config-Optics)#commit
Router(config-Optics)#exit

/*To enable etm for breakout ports, as for 10 GbE*/
Router#configure terminal
Router(config)#controller optics 0/5/0/27/1
Router(config-Optics)#mode etm
Router(config-Optics)#commit
Router(config-Optics)#exit

/*Match on traffic-class on ETM-enabled interfaces requires running the following hw-module
command to truncate IPv6 fields, thus freeing TCAM space for packet header marking for
VOQ lookup in the second pass.*/
Router(config)#hw-module profile qos ipv6 short-etm
In order to activate this profile, you must manually reload the chassis/all line cards
Router(config)#commit
Router(config)#exit
Router#reload

/*Configure etm policies on subinterfaces
Configure class map. Specify the match criteria*/
Router(config)#class-map match-any MAIN_OUT_CMAP_1
Router(config-cmap)#match traffic-class 1
Router(config-cmap)#end-class-map
```

```

Router(config-cmap) #exit
Router(config) #class-map match-any MAIN_OUT_CMAP_2
Router(config-cmap) #match traffic-class 2
Router(config-cmap) #end-class-map
Router(config-cmap) #exit
Router(config) #class-map match-any MAIN_OUT_CMAP_3
Router(config-cmap) #match traffic-class 3
Router(config-cmap) #end-class-map
Router(config-cmap) #exit
Router(config) #class-map match-any MAIN_OUT_CMAP_4
Router(config-cmap) #match traffic-class 4
Router(config-cmap) #end-class-map
Router(config-cmap) #exit
Router(config) #class-map match-any MAIN_OUT_CMAP_5
Router(config-cmap) #match traffic-class 5
Router(config-cmap) #end-class-map
Router(config-cmap) #exit
Router(config) #class-map match-any MAIN_OUT_CMAP_6
Router(config-cmap) # match traffic-class 6
Router(config-cmap) #end-class-map
Router(config-cmap) #exit
Router(config) #class-map match-any MAIN_OUT_CMAP_7
Router(config-cmap) #match traffic-class 7
Router(config-cmap) #end-class-map
Router(config-cmap) #exit

/*Create a policy map*/
Router(config) #policy-map MAIN_OUT_PMAP5_PHY_SUB

/*Associate the traffic class with the policy map*/
Router(config-pmap) #class MAIN_OUT_CMAP_1

/*Specify the traffic class actions ('bandwidth remaining ratio, in this example)*/
Router(config-pmap-c) #bandwidth remaining ratio 20
Router(config-pmap-c) #set traffic-class 1
Router(config-cmap-c) #commit
/*Repeat for other classes*/
Router(config-pmap-c) #class MAIN_OUT_CMAP_2
Router(config-pmap-c) #bandwidth remaining ratio 30
Router(config-pmap-c) #set traffic-class 2
Router(config-pmap-c) #exit
Router(config-pmap-c) #class MAIN_OUT_CMAP_3
Router(config-pmap-c) #bandwidth remaining ratio 40
Router(config-pmap-c) #set traffic-class 3
Router(config-pmap-c) #exit
Router(config-pmap-c) #class MAIN_OUT_CMAP_4
Router(config-pmap-c) #bandwidth remaining ratio 50
Router(config-pmap-c) #set traffic-class 4
Router(config-pmap-c) #exit
Router(config-pmap-c) #class MAIN_OUT_CMAP_5
Router(config-pmap-c) #bandwidth remaining ratio 60
Router(config-pmap-c) #set traffic-class 5
Router(config-pmap-c) #exit
Router(config-pmap-c) #class MAIN_OUT_CMAP_6
Router(config-pmap-c) #bandwidth remaining ratio 60
Router(config-pmap-c) #set traffic-class 6
Router(config-pmap-c) #exit
Router(config-pmap-c) #class MAIN_OUT_CMAP_7
Router(config-pmap-c) #set traffic-class 7
Router(config-pmap-c) #priority level 1
Router(config-pmap-c) #exit
Router(config-pmap-c) #class class-default

```

```

Router(config-pmap-c)#bandwidth remaining ratio 10
Router(config-pmap-c)#exit
Router(config-pmap-c)#end-policy-map
Router(config)#exit
Router(config)#policy-map PARENT_MAIN_OUT_PHY_SUB
Router(config-pmap)#class class-default
Router(config-pmap-c)#service-policy MAIN_OUT_PMAP5_PHY_SUB
Router(config-pmap-c)#shape average 55 mbps
Router(config-pmap-c)#exit
Router(config-pmap-c)#end-policy-map
Router(config)#commit

/*Attach the policy map to the subinterface.*/
Router#configure
Router(config)#interface HundredGigE 0/5/0/16.1
Router(config-subif)#service-policy output PARENT_MAIN_OUT_PHY_SUB
Router(config-subif)#commit

```

You have enabled **etm** mode and configured policies in this mode. You are now ready to experience existing QoS features with enhanced egress classification for queueing policy and scale egress queueing policies.

Running Configuration

```

controller Optics 0/5/0/16
    mode etm
!
class-map match-any MAIN_OUT_CMAP_1
    match traffic-class 1
end-class-map
!
class-map match-any MAIN_OUT_CMAP_2
    match traffic-class 2
end-class-map
!
class-map match-any MAIN_OUT_CMAP_3
    match traffic-class 3
end-class-map
!
class-map match-any MAIN_OUT_CMAP_4
    match traffic-class 4
end-class-map
!
class-map match-any MAIN_OUT_CMAP_5
    match traffic-class 5
end-class-map
!
class-map match-any MAIN_OUT_CMAP_6
    match traffic-class 6
end-class-map
!
class-map match-any MAIN_OUT_CMAP_7
    match traffic-class 7
end-class-map
!
policy-map MAIN_OUT_PMAP5_PHY_SUB
    class MAIN_OUT_CMAP_1
        bandwidth remaining ratio 20
        set traffic-class 1
    !
    class MAIN_OUT_CMAP_2
        bandwidth remaining ratio 30
        set traffic-class 2
    !

```

```

class MAIN_OUT_CMAP_3
  bandwidth remaining ratio 40
  set traffic-class 3
!
class MAIN_OUT_CMAP_4
  bandwidth remaining ratio 50
  set traffic-class 4
!
class MAIN_OUT_CMAP_5
  bandwidth remaining ratio 60
  set traffic-class 5
!
class MAIN_OUT_CMAP_6
  bandwidth remaining ratio 60
  set traffic-class 6
!
class MAIN_OUT_CMAP_7
  set traffic-class 7
  priority level 1
!
class class-default
  bandwidth remaining ratio 10
!
end-policy-map
!
policy-map PARENT_MAIN_OUT_PHY_SUB
class class-default
  service-policy MAIN_OUT_PMAP5_PHY_SUB
  shape average 55 mbps
!
end-policy-map

```

Verification

To check if interface is in **etm** mode:

```
Router#show controllers npu voq-usage interface all instance all 1$
```

```

-----
Node ID: 0/0/CPU0
Intf Intf NPU NPU PP Sys VOQ Flow VOQ Port
name handle # core Port Port base base port speed
(hex) type
-----
Hu0/0/0/0 8 0 0 1 24577 38400 6912 local 100G
Hu0/0/0/0 8 0 0 154 1 1536 22272 local 100G ----->2nd Entry points to Recycle Port Mapping

```

To verify the output for policy map with class match on traffic class:

```

Router#show qos interface HundredGigE 0/5/0/16.1 output
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/5/0/16.1 ifh 0xa008002 -- output policy
NPU Id: 1
Total number of classes: 9
Interface Bandwidth: 100000000 kbps
Policy Name: PARENT_MAIN_OUT_PHY_SUB
SPI Id: 0x0
VOQ Base: 13328
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = class-default
Queue Max. BW. = 70311 kbps (55 mbits/sec)
Queue Min. BW. = 58593 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)

```



```

Level2 Class                               = MAIN_OUT_CMAP_1
New traffic class                          = 1
Egressq Queue ID                           = 13329 (LP queue)
Queue Max. BW.                             = no max (default)
Queue Min. BW.                             = 0 kbps (default)
Inverse Weight / Weight                    = 32 / (20)
Peak burst                                 = 36864 bytes (default)
TailDrop Threshold                         = 1280 bytes / 19 ms (default)
WRED not configured for this class

Level2 Class                               = MAIN_OUT_CMAP_2
New traffic class                          = 2
Egressq Queue ID                           = 13330 (LP queue)
Queue Max. BW.                             = no max (default)
Queue Min. BW.                             = 0 kbps (default)
Inverse Weight / Weight                    = 21 / (30)
Peak burst                                 = 36864 bytes (default)
TailDrop Threshold                         = 1280 bytes / 19 ms (default)
WRED not configured for this class

Level2 Class                               = MAIN_OUT_CMAP_3
New traffic class                          = 3
Egressq Queue ID                           = 13331 (LP queue)
Queue Max. BW.                             = no max (default)
Queue Min. BW.                             = 0 kbps (default)
Inverse Weight / Weight                    = 16 / (40)
Peak burst                                 = 36864 bytes (default)
TailDrop Threshold                         = 1280 bytes / 19 ms (default)
WRED not configured for this class

Level2 Class                               = MAIN_OUT_CMAP_4
New traffic class                          = 4
Egressq Queue ID                           = 13332 (LP queue)
Queue Max. BW.                             = no max (default)
Queue Min. BW.                             = 0 kbps (default)
Inverse Weight / Weight                    = 12 / (50)
Peak burst                                 = 36864 bytes (default)
TailDrop Threshold                         = 1280 bytes / 19 ms (default)
WRED not configured for this class

Level2 Class                               = MAIN_OUT_CMAP_5
New traffic class                          = 5
Egressq Queue ID                           = 13333 (LP queue)
Queue Max. BW.                             = no max (default)
Queue Min. BW.                             = 0 kbps (default)
Inverse Weight / Weight                    = 10 / (60)
Peak burst                                 = 36864 bytes (default)
TailDrop Threshold                         = 1280 bytes / 19 ms (default)
WRED not configured for this class

Level2 Class                               = MAIN_OUT_CMAP_6
New traffic class                          = 6
Egressq Queue ID                           = 13334 (LP queue)
Queue Max. BW.                             = no max (default)
Queue Min. BW.                             = 0 kbps (default)
Inverse Weight / Weight                    = 10 / (60)
Peak burst                                 = 36864 bytes (default)
TailDrop Threshold                         = 1280 bytes / 19 ms (default)
WRED not configured for this class

Level2 Class (HP1)                         = MAIN_OUT_CMAP_7
New traffic class                          = 7
Egressq Queue ID                           = 13335 (HP1 queue)
Queue Max. BW.                             = no max (default)

```

```

Peak burst                               = 36864 bytes (default)
TailDrop Threshold                       = 68608 bytes / 10 ms (default)
WRED not configured for this class

Level2 Class                             = class-default
Egressq Queue ID                         = 13328 (Default LP queue)
Queue Max. BW.                           = no max (default)
Queue Min. BW.                           = 0 kbps (default)
Inverse Weight / Weight                   = 64 / (10)
Peak burst                               = 36864 bytes (default)
TailDrop Threshold                       = 1280 bytes / 19 ms (default)

```

Configuration Example for Policy Map with Class Match on DSCP

```

Router(config)#class-map match-any MATCH_DSCP_CONTROL
Router(config-cmap)#description "Customer Control Traffic"
Router(config-cmap)#match dscp 48-58
Router(config-cmap)#end-class-map
Router(config)#class-map match-any MATCH_DSCP_COS1_RT
Router(config-cmap)#description "Real Time Traffic"
Router(config-cmap)#match dscp 40 46
Router(config-cmap)#end-class-map
Router(config)#class-map match-any MATCH_DSCP_COS2_IC
Router(config-cmap)#description "Bursty High In Contract Traffic"
Router(config-cmap)#match dscp 24 26
Router(config-cmap)#end-class-map
Router(config)#class-map match-any MATCH_DSCP_COS2_OC
Router(config-cmap)#description "Bursty High Out of Contract Traffic"
Router(config-cmap)#match dscp 25 27 28 29 30 31
Router(config-cmap)#end-class-map
Router(config)#class-map match-any MATCH_DSCP_COS3_IC
Router(config-cmap)#description "Bursty Low In Contract Traffic"
Router(config-cmap)#match dscp 16 18
Router(config-cmap)#end-class-map
Router(config)#class-map match-any MATCH_DSCP_COS3_OC
Router(config-cmap)#description "Bursty Low Out of Contract Traffic"
Router(config-cmap)#match dscp 17 19 20 21 22 23
Router(config-cmap)#end-class-map
Router(config)#commit

Router(config)#policy-map M4_Ex_1G_Child_80:40:30:30_out_ss
Router(config-pmap)#class MATCH_DSCP_CONTROL
Router(config-pmap-c)#shape average percent 5
Router(config-pmap-c)#set traffic-class 7
Router(config-pmap-c)#priority level 2

Router(config-pmap-c)#class MATCH_DSCP_COS1_RT
Router(config-pmap-c)#shape average percent 80
Router(config-pmap-c)#set traffic-class 5
Router(config-pmap-c)#priority level 1
Router(config-pmap-c)#class MATCH_DSCP_COS2_IC
Router(config-pmap-c)#set traffic-class 2
Router(config-pmap-c)#bandwidth remaining percent 40
Router(config-pmap-c)#queue-limit 4096 packets
Router(config-pmap-c)#class MATCH_DSCP_COS3_IC
Router(config-pmap-c)#bandwidth remaining percent 30
Router(config-pmap-c)#set traffic-class 3
Router(config-pmap-c)#queue-limit 3072 packets
Router(config-pmap-c)#class class-default
Router(config-pmap-c)#bandwidth remaining percent 30
Router(config-pmap-c)#set traffic-class 0
Router(config-pmap-c)#queue-limit 3072 packets

```

```

Router(config-pmap-c) #end-policy-map
Router(config) #policy-map M4_Ex_1G_80:40:30:30_out_ss
Router(config-pmap) #class class-default
Router(config-pmap-c) #service-policy M4_Ex_1G_Child_80:40:30:30_out_ss
Router(config-pmap-c) #shape average 100 mbps
Router(config-pmap-c) #commit
Router(config-pmap-c) #exit
Router(config-pmap) #exit
Router(config) #interface HundredGigE 0/5/0/16.1
Router(config-subif) #service-policy output M4_Ex_1G_80:40:30:30_$
Router(config-subif) #commit

```

Running Configuration

```

class-map match-any MATCH_DSCP_CONTROL
  description "Customer Control Traffic"
  match dscp 48-58
end-class-map
!
class-map match-any MATCH_DSCP_COS1_RT
  description "Real Time Traffic"
  match dscp 40 46
end-class-map
!
class-map match-any MATCH_DSCP_COS2_IC
  description "Bursty High In Contract Traffic"
  match dscp 24 26
end-class-map
!
class-map match-any MATCH_DSCP_COS3_IC
  description "Bursty Low In Contract Traffic"
  match dscp 16 18
end-class-map
!
policy-map M4_Ex_1G_Child_80:40:30:30_out_ss
  class MATCH_DSCP_CONTROL
    shape average percent 5
    set traffic-class 7
    priority level 2
  !
  class MATCH_DSCP_COS1_RT
    shape average percent 80
    set traffic-class 5
    priority level 1
  !
  class MATCH_DSCP_COS2_IC
    set traffic-class 2
    bandwidth remaining percent 40
    queue-limit 4096 packets
  !
  class MATCH_DSCP_COS3_IC
    bandwidth remaining percent 30
    set traffic-class 3
    queue-limit 3072 packets
  !
  class class-default
    bandwidth remaining percent 30
    set traffic-class 0
    queue-limit 3072 packets
  !
end-policy-map
!
policy-map M4_Ex_1G_80:40:30:30_out_ss

```

```

class class-default
  service-policy M4_Ex_1G_Child_80:40:30:30_out_ss
  shape average 100 mbps
!
end-policy-map

```

Verify the Output for Policy Map with Class Match on DSCP

Router#show qos interface HundredGigE 0/5/0/16.1 output

NOTE:- Configured values are displayed within parentheses

Interface HundredGigE0/5/0/16.1 ifh 0xa008002 -- output policy

```

NPU Id: 1
Total number of classes: 6
Interface Bandwidth: 100000000 kbps
Policy Name: M4_Ex_1G_80:40:30:30_out_ss
SPI Id: 0x0
VOQ Base: 13328
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = class-default
Queue Max. BW. = 113282 kbps (100 mbits/sec)
Queue Min. BW. = 101564 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)

Level2 Class (HP2) = MATCH_DSCP_CONTROL
New traffic class = 7
Egressq Queue ID = 13335 (HP2 queue)
Queue Max. BW. = 7812 kbps (5 %)
Guaranteed service rate = 5000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 6144 bytes / 10 ms (default)
WRED not configured for this class

Level2 Class (HP1) = MATCH_DSCP_COS1_RT
New traffic class = 5
Egressq Queue ID = 13333 (HP1 queue)
Queue Max. BW. = 82031 kbps (80 %)
Guaranteed service rate = 80000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 99840 bytes / 10 ms (default)
WRED not configured for this class

Level2 Class = MATCH_DSCP_COS2_IC
New traffic class = 2
Egressq Queue ID = 13330 (LP queue)
Queue Max. BW. = no max (default)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 2 / (40%)
Guaranteed service rate = 6000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 1048576 bytes / 1398 ms (4096 packets)
WRED not configured for this class

Level2 Class = MATCH_DSCP_COS3_IC
New traffic class = 3
Egressq Queue ID = 13331 (LP queue)
Queue Max. BW. = no max (default)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 2 / (30%)
Guaranteed service rate = 4500 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 786432 bytes / 1398 ms (3072 packets)
WRED not configured for this class

Level2 Class = class-default

```

New traffic class	=	0
Egressq Queue ID	=	13328 (Default LP queue)
Queue Max. BW.	=	no max (default)
Queue Min. BW.	=	0 kbps (default)
Inverse Weight / Weight	=	2 / (30%)
Guaranteed service rate	=	4500 kbps
Peak burst	=	36864 bytes (default)
TailDrop Threshold	=	786432 bytes / 1398 ms (3072 packets)
WRED not configured for this class		

Configuration Example for Policy Map with Class Match on COS

```

Router(config)#class-map match-any cos2
  Router(config-cmap)#match cos 2
Router(config-cmap)#end-class-map
Router(config)#class-map match-all cos3
  Router(config-cmap)#match cos 3
Router(config-cmap)#end-class-map
Router(config)#class-map match-all cos4
  Router(config-cmap)#match cos 4
Router(config-cmap)#end-class-map
Router(config)#class-map match-all cos5
  Router(config-cmap)#match cos 5
Router(config-cmap)#end-class-map
Router(config)#class-map match-all cos6
  Router(config-cmap)#match cos 6
Router(config-cmap)#end-class-map
Router(config)#class-map match-all cos7
  Router(config-cmap)#match cos 7
Router(config-cmap)#end-class-map
Router(config)#class-map match-all cos1
  Router(config-cmap)#match cos 1
Router(config-cmap)#end-class-map
Router(config)#policy-map L2VPN-OUT
  Router(config-pmap)#class cos2
    Router(config-pmap-c)#set traffic-class 2
    Router(config-pmap-c)#shape average 45 mbps
  Router(config-pmap-c)#class cos3
    Router(config-pmap-c)#set traffic-class 3
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#class cos4
    Router(config-pmap-c)#set traffic-class 4
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#class cos5
    Router(config-pmap-c)#set traffic-class 5
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#class cos6
    Router(config-pmap-c)#set traffic-class 6
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#class cos7
    Router(config-pmap-c)#set traffic-class 7
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#class cos1
    Router(config-pmap-c)#set traffic-class 1
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#class class-default
    Router(config-pmap-c)#shape average 50 mbps
  Router(config-pmap-c)#end-policy-map
Router(config)#commit
Router(config)#exit

```

```
Router(config)#interface HundredGigE 0/5/0/16.1
Router(config-subif)#no service-policy output
Router(config-subif)#commit

Router(config-subif)#service-policy output L2VPN-OUT
Router(config-subif)#commit
Router(config-subif)#end
```

Running Configuration

```
class-map match-any cos2
  match cos 2
end-class-map
!
class-map match-all cos3
  match cos 3
end-class-map
!
class-map match-all cos4
  match cos 4
end-class-map
!
class-map match-all cos5
  match cos 5
end-class-map
!
class-map match-all cos6
  match cos 6
end-class-map
!
class-map match-all cos7
  match cos 7
end-class-map
!
class-map match-all cos1
  match cos 1
end-class-map
!
class-map match-any class-default
end-class-map
!
policy-map L2VPN-OUT
  class cos2
    set traffic-class 2
    shape average 45 mbps
  !
  class cos3
    set traffic-class 3
    shape average 50 mbps
  !
  class cos4
    set traffic-class 4
    shape average 50 mbps
  !
  class cos5
    set traffic-class 5
    shape average 50 mbps
  !
  class cos6
    set traffic-class 6
    shape average 50 mbps
  !
  class cos7
    set traffic-class 7
    shape average 50 mbps
```

```

!
class cos1
  set traffic-class 1
  shape average 50 mbps
!
class class-default
  shape average 50 mbps
!
end-policy-map

```

Verify the Output for Policy Map with Class Match on COS

Router#show qos interface HundredGigE 0/5/0/16.1 output

NOTE:- Configured values are displayed within parentheses

Interface HundredGigE0/5/0/16.1 ifh 0xa008002 -- output policy

```

NPU Id: 1
Total number of classes: 8
Interface Bandwidth: 100000000 kbps
Policy Name: L2VPN-OUT
SPI Id: 0x0
VOQ Base: 13328
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----

```

```

Level1 Class = cos2
New traffic class = 2
Egressq Queue ID = 13330 (LP queue)
Queue Max. BW. = 46875 kbps (45 mbits/sec)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 45000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 56064 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class = cos3
New traffic class = 3
Egressq Queue ID = 13331 (LP queue)
Queue Max. BW. = 50781 kbps (50 mbits/sec)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 50000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 62464 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class = cos4
New traffic class = 4
Egressq Queue ID = 13332 (LP queue)
Queue Max. BW. = 50781 kbps (50 mbits/sec)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 50000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 62464 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class = cos5
New traffic class = 5
Egressq Queue ID = 13333 (LP queue)
Queue Max. BW. = 50781 kbps (50 mbits/sec)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 50000 kbps
Peak burst = 36864 bytes (default)
TailDrop Threshold = 62464 bytes / 10 ms (default)

```

WRED not configured for this class

```

Level1 Class                = cos6
New traffic class           = 6
Egressq Queue ID           = 13334 (LP queue)
Queue Max. BW.             = 50781 kbps (50 mbits/sec)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 / (BWR not configured)
Guaranteed service rate    = 50000 kbps
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 62464 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class                = cos7
New traffic class           = 7
Egressq Queue ID           = 13335 (LP queue)
Queue Max. BW.             = 50781 kbps (50 mbits/sec)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 / (BWR not configured)
Guaranteed service rate    = 50000 kbps
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 62464 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class                = cos1
New traffic class           = 1
Egressq Queue ID           = 13329 (LP queue)
Queue Max. BW.             = 50781 kbps (50 mbits/sec)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 / (BWR not configured)
Guaranteed service rate    = 50000 kbps
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 62464 bytes / 10 ms (default)
WRED not configured for this class

```

```

Level1 Class                = class-default
Egressq Queue ID           = 13328 (Default LP queue)
Queue Max. BW.             = 50781 kbps (50 mbits/sec)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 / (BWR not configured)
Guaranteed service rate    = 50000 kbps
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 62464 bytes / 10 ms (default)
WRED not configured for this class

```

Configuration Example for Policy Map with Class Match on Traffic Class and DSCP

```

Router(config)#class-map match-any cm-dscp-cs1
Router(config-cmap)#match dscp cs1
Router(config-cmap)#class-map match-any cm-dscp-cs2
Router(config-cmap)#match dscp cs2
Router(config-cmap)#class-map match-any cm-dscp-af11
Router(config-cmap)#match dscp af11
Router(config-cmap)#class-map match-any cm-tc-5
Router(config-cmap)#match traffic-class 5
Router(config-cmap)#class-map match-all cm-tc-6-and-dscp-cs3
Router(config-cmap)#match traffic-class 6
Router(config-cmap)#match dscp cs3
Router(config-cmap)#policy-map pm-out-etm-queueing
Router(config-pmap)#class cm-dscp-cs1
Router(config-pmap-c)#set traffic-class 1
Router(config-pmap-c)#shape average percent 45

```



```

Router(config-pmap-c)#class cm-dscp-cs2
Router(config-pmap-c)#set traffic-class 1
Router(config-pmap-c)#shape average percent 45
Router(config-pmap-c)#class cm-dscp-af11
Router(config-pmap-c)#set traffic-class 2
Router(config-pmap-c)#shape average percent 10
Router(config-pmap-c)#class cm-tc-5
Router(config-pmap-c)#shape average percent 5
Router(config-pmap-c)#set traffic-class 3
Router(config-pmap-c)#class cm-tc-6-and-dscp-cs3
Router(config-pmap-c)#set traffic-class 7
Router(config-pmap-c)#priority
Router(config-pmap-c)#commit
Router(config-pmap-c)#exit
Router(config-pmap)#exit
Router(config)#interface HundredGigE 0/5/0/16.1
Router(config-subif)#service-policy output pm-out-etm-queueing
Router(config-subif)#commit
Router(config-subif)#end

```

Running Configuration

```

class-map match-any cm-dscp-cs1
match dscp cs1
end-class-map
!
class-map match-any cm-dscp-cs2
match dscp cs2
end-class-map
!
class-map match-any cm-dscp-af11
match dscp af11
end-class-map
!
class-map match-any cm-tc-5
match traffic-class 5
end-class-map
!
class-map match-all cm-tc-6-and-dscp-cs3
match traffic-class 6
match dscp af22 cs3
end-class-map
!
policy-map pm-out-etm-queueing
class cm-dscp-cs1
set traffic-class 1
shape average percent 45
!
class cm-dscp-cs2
set traffic-class 1
shape average percent 45
!
class cm-dscp-af11
set traffic-class 2
shape average percent 10
!
class cm-tc-5
shape average percent 5
set traffic-class 3
!
class cm-tc-6-and-dscp-cs3
set traffic-class 7
priority level 1
!

```

```

class class-default
!
end-policy-map

```

Notes on the Configuration Example

- Both the traffic classes **cm-dscp-cs1** and **cm-dscp-cs2** have the same action for set traffic-class in the EM mode, which is **set traffic-class 1**. This means that both these classes have matching traffic going to the same VOQ, and hence their statistics refer to the statistics for VOQ=1.
- The traffic class **cm-dscp-af11** has the action set traffic-class 2 at the recycle pass in the ETM mode. Hence, its matching traffic and statistics are from VOQ=2.
- The traffic class **cm-tc-6-and-dscp-cs3** is an example of remarking traffic-class. This class has traffic-class=6 configured at the ingress, and matches all traffic that has incoming traffic-class=6 AND incoming DSCP=CS3. However, at egress, the traffic-class is remarked to 7. Hence, matching traffic goes to VOQ=7, and the statistics for this class are taken from VOQ=7.

Verify the Output for Policy Map with Class Match on Traffic Class and DSCP

```

Router#show qos interface HundredGigE 0/5/0/16.1 output
NOTE:- Configured values are displayed within parentheses
Interface HundredGigE0/5/0/16.1 ifh 0xa008002 -- output policy
NPU Id: 1
Total number of classes: 6
Interface Bandwidth: 100000000 kbps
Policy Name: pm-out-etm-queueing
SPI Id: 0x0
VOQ Base: 13328
Accounting Type: Layer1 (Include Layer 1 encapsulation and above)
-----
Level1 Class = cm-dscp-cs1
New traffic class = 1
Egressq Queue ID = 13329 (LP queue)
Queue Max. BW. = 45003261 kbps (45 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Peak burst = 36864 bytes (default)
TailDrop Threshold = 1249792 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class = cm-dscp-cs2
New traffic class = 1
Egressq Queue ID = 13329 (LP queue)
Queue Max. BW. = 45003261 kbps (45 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Peak burst = 36864 bytes (default)
TailDrop Threshold = 1249792 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class = cm-dscp-af11
New traffic class = 2
Egressq Queue ID = 13330 (LP queue)
Queue Max. BW. = 10000152 kbps (10 %)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Peak burst = 36864 bytes (default)
TailDrop Threshold = 1249792 bytes / 10 ms (default)
WRED not configured for this class

Level1 Class = cm-tc-5
New traffic class = 3

```

```

Egressq Queue ID           = 13331 (LP queue)
Queue Max. BW.             = 5000076 kbps (5 %)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 / (BWR not configured)
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 1249792 bytes / 10 ms (default)
WRED not configured for this class

Levell Class (HP1)         = cm-tc-6-and-dscp-cs3
New traffic class          = 7
Egressq Queue ID          = 13335 (HP1 queue)
Queue Max. BW.             = no max (default)
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 124999936 bytes / 10 ms (default)
WRED not configured for this class

Levell Class               = class-default
Egressq Queue ID          = 13328 (Default LP queue)
Queue Max. BW.             = no max (default)
Queue Min. BW.             = 0 kbps (default)
Inverse Weight / Weight    = 1 / (BWR not configured)
Peak burst                 = 36864 bytes (default)
TailDrop Threshold         = 1249792 bytes / 10 ms (default)
WRED not configured for this class

```

Associated Commands

controller optics

Low Rate Shaper

Table 31: Feature History Table

Feature Name	Release Information	Feature Description
Low Rate Shaper	Release 7.6.1	<p>This functionality allows your hardware to program low shaper bandwidth with accuracy. Earlier, if one of your customers wanted, say, a bandwidth of 500 kbps, the Cisco NC 57 line card ASIC could only program a minimum bandwidth of 3.9 Mbps for a particular traffic class, resulting in unnecessary wastage of bandwidth. With low shaper rates now made possible, you can limit bandwidth for specific subscribers whose requirements don't call for excessive bandwidth. This functionality thus allows you to save on your bandwidth costs.</p> <p>This functionality is only available for and automatically enabled in the egress traffic management mode.</p>

- **What is the low rate shaper functionality all about?**—As the name suggests, the functionality is about providing accurate low shaper rate (less than 5 Mbps). But the powerful aspect of this feature is that the low shaper rates that it enables are far lesser what your existing Cisco NC 57 line card ASICs allow.

This functionality is available only when you configure your egress interface in the Egress Traffic Management (Egress TM) mode. It's automatically enabled in this mode.

- **What are the benefits of this functionality?**—This functionality is supported on Cisco NC 57 line cards. The forwarding ASIC in this line card is capable of supporting only 3.9 Mbps as the lowest rate for a particular traffic class, even if you configure a shaping bandwidth less than 3.9 Mbps. When you have a typical network set up that has hundreds of edge-facing routers and thousands of interfaces, this restriction meant that excess bandwidth was distributed even to those subscribers who didn't require it, and hence didn't pay for it.

With this functionality, you can save on your bandwidth costs by limiting the bandwidth for specific subscribers who require a lower spectrum of capacities. Hence, the low-rate shaper also provides you with more accuracy in shaping your hardware's bandwidth. The result is optimal bandwidth distribution and savings for your operations.

- **How does this functionality work? Any important points?**

- When you configure a shaper rate of any value below 5 MBps for the virtual output queue (VOQ) of an ETM-enabled egress interface, the system automatically shifts into low-rate shaper mode.

- Once in low-rate shaper mode, your system supports traffic shaping rates to a particular VOQ in a specific range. This range is from 122 kbps through 80 Gbps.
- To support low-shaping capability, you need to reserve low-rate connectors. The system has 64 connector regions and each region contains 1024 connectors.
- Low connectors get allocated statically during system initialization, and the allocation can't be changed dynamically. You can't use connectors that are reserved and configured as low connectors, as normal connectors, and the converse holds good as well.
- You have the flexibility to reserve the required connectors for low and normal regions before applying all QoS policies at egress.
- A connector switches from low to normal or the other way around if:
 - a low shaper rate (less than 5 Mbps) is detected on any of the traffic classes (VOQ) of an interface, where egress QoS policy map is applied. In this case, all the eight normal connectors are replaced by low connectors.
 - post-policy update or removal, no low shaper rate (less than 5 Mbps) is detected on an egress QoS Policy Map that is applied on VOQs attached with low connectors. In this case, low connectors are replaced by normal connectors.

• **How many egress policy maps can I configure with low rate shaper?**

By default, 2 connector regions per core are reserved for low connectors.

Therefore, per core, the number of low rate connectors are: number of low rate connectors per region x number of connector regions per core = 1024 x 2 = **2048**

Each interface VOQ set requires **8** connectors (or 8 traffic classes).

Therefore, the number of interface VOQ sets (or policy maps) that you can configure per core = number of low rate connectors/number of connectors per interface VOQ set = 2048/8 = **256**

• **Any limitations or restrictions?**

- A low connector can't support more than 80 Gbps traffic shaping rate. Which means that a VOQ attached to a low connector can't forward more than 80 Gbps of traffic.
- Attaching and detaching low connector sets to and from egress interfaces impacts traffic flows in unrelated VOQs. This could affect traffic flows when connectors are switched.
- This functionality is applicable only for local switching on an ETM-enabled egress interface. This is why low-rate shaper benefits do not accrue to ingress traffic from a remote line card.

• **How do I verify the low shaper rate action?**—Run the **show qos interface** command to see the low shaper statistics. For example:

```
router#show qos interface GigabitEthernet 0/2/0/0 output
Level2 Class = expl
New traffic class = 1
Egressq Queue ID = 1793 (LP queue)
Queue Max. BW. = 3051 kbps (3 mbits/sec)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 625 kbps
Peak burst = 32832 bytes (default)
TailDrop Threshold = 1280 bytes / 16 ms (default)
```

```
LOW SHAPER = Enabled

WRED not configured for this class

Level2 Class = exp2
New traffic class = 2
Egressq Queue ID = 1794 (LP queue)
Queue Max. BW. = 5004 kbps (5 mbits/sec)
Queue Min. BW. = 0 kbps (default)
Inverse Weight / Weight = 1 / (BWR not configured)
Guaranteed service rate = 625 kbps
Peak burst = 32832 bytes (default)
TailDrop Threshold = 1280 bytes / 16 ms (default)
```

LOW SHAPER = Enabled

```
WRED not configured for this class
```